

# Succinct

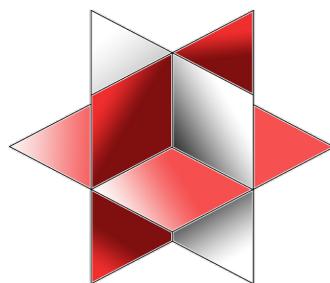
## Enabling Queries on Compressed Data



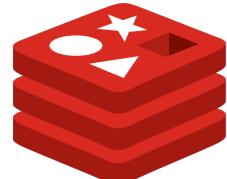
# Distributed data stores

---

Fundamental building blocks for many cloud services



elasticsearch



redis

MICA

RAMCloud

MemC3



# Records, primary key, secondary keys

---

Data stored as collection of records

USER	POSITION	BOSS	BOSS' Description
Anurag	Student	Rachit	Serious, Pushy
....			
Rachit	Postdoc	Ion	Curious, funny

# Records, primary key, secondary keys

---

Data stored as collection of records

Primary key

USER	POSITION	BOSS	BOSS' Description
Anurag	Student	Rachit	Serious, Pushy
....			
Rachit	Postdoc	Ion	Curious, funny

# Records, primary key, secondary keys

---

Data stored as collection of records

Secondary keys

USER	POSITION	BOSS	BOSS' Description
Anurag	Student	Rachit	Serious, Pushy
....			
Rachit	Postdoc	Ion	Curious, funny

# Queries in data stores

---

## Queries on primary keys

- GET, PUT, DELETE
- Very efficient [MICA, Redis, RAMCloud]

# Queries in data stores

---

## Queries on primary keys

- GET, PUT, DELETE
- Very efficient [MICA, Redis, RAMCloud]

## Queries on secondary keys

- E.g., search
- Executed via
  - Data scans [columnar stores]
  - Indexes [Cassandra, MongoDB]
- Either slow or large storage overhead

# Flat (unstructured) files

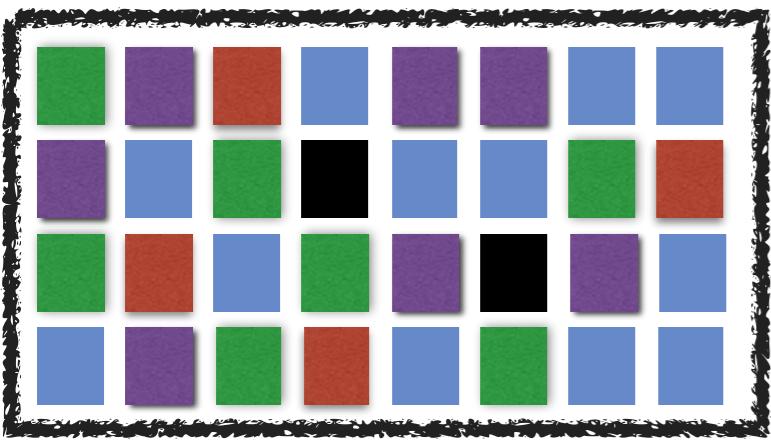
---

For now, focus on:

- Queries on flat (unstructured) files
- Very powerful primitive
  - Allows queries on semi-structured data
  - Discussed later

# Example: Search( ■ , file)

---

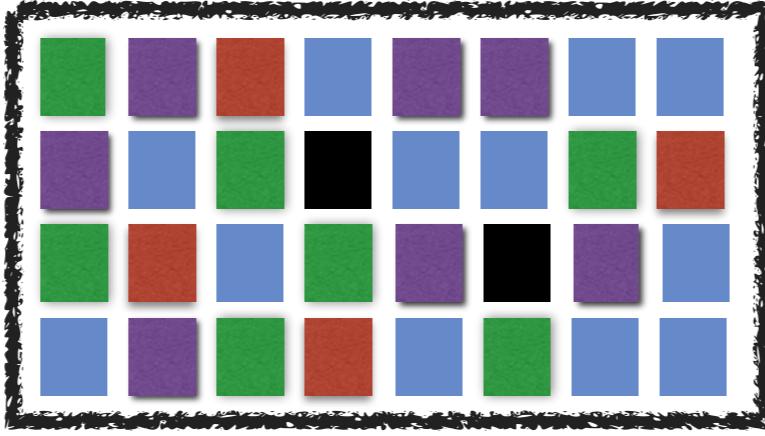
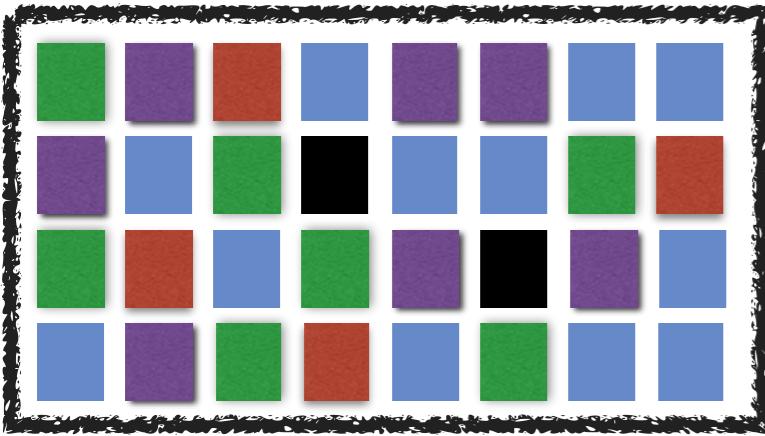


Search( ■ )

# Example: Search( ■ , file)

---

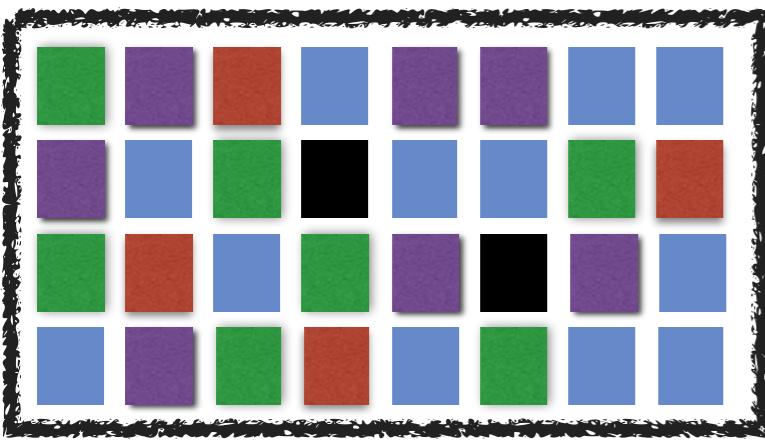
Data Scans



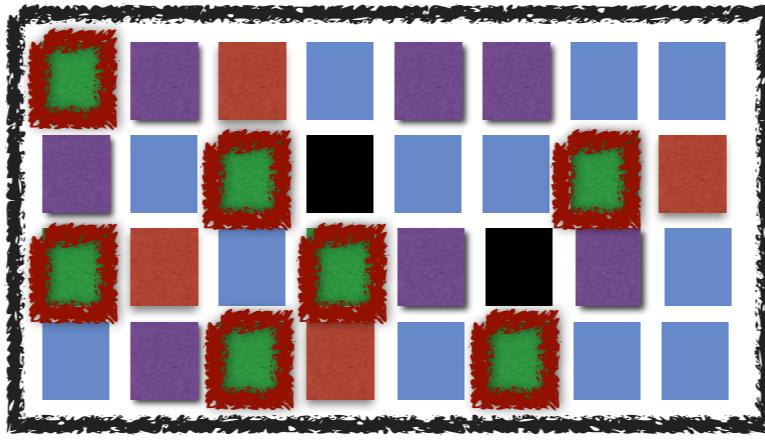
Search( ■ )

# Example: Search( ■ , file)

---



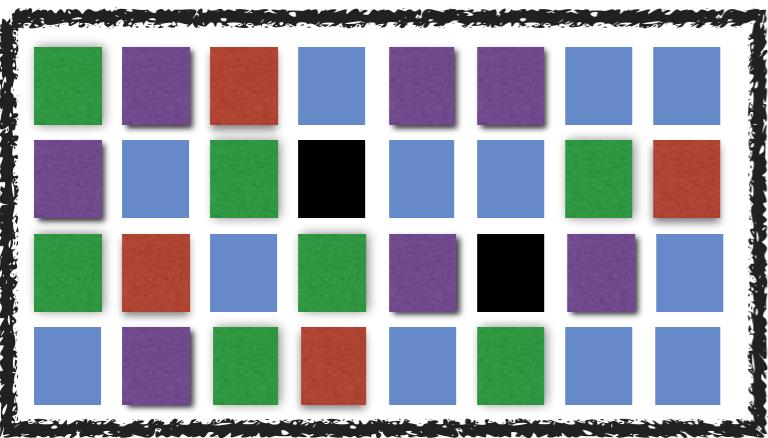
Search( ■ )



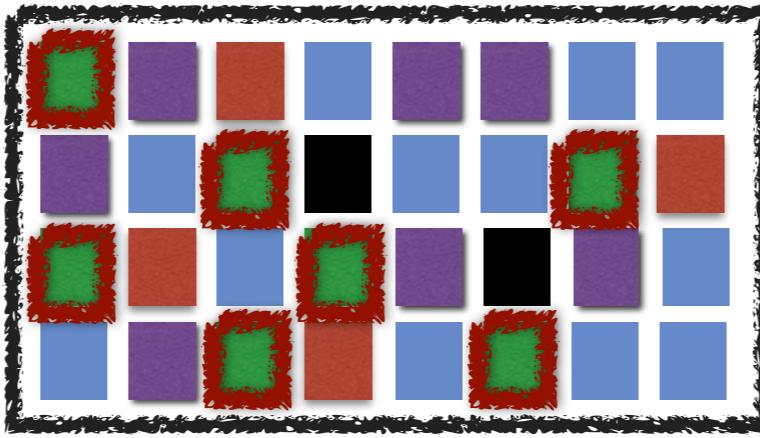
Data Scans

# Example: Search( , file)

---



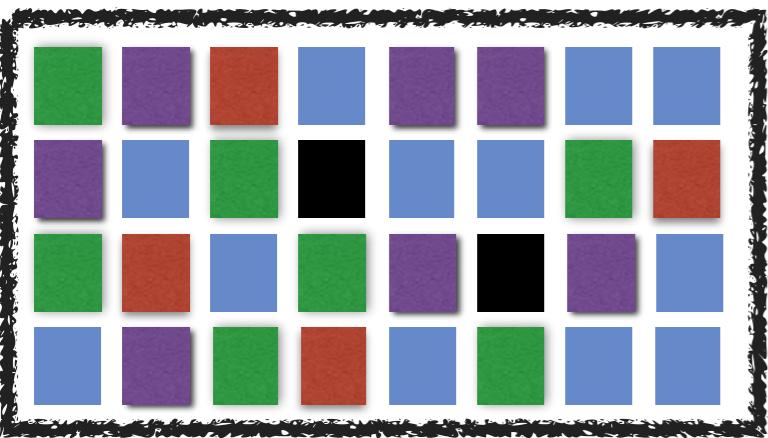
Search(  )



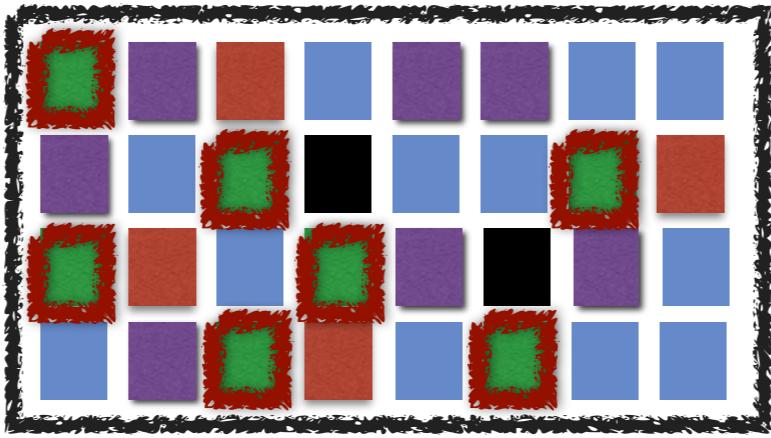
Data Scans

Low storage  
High Latency

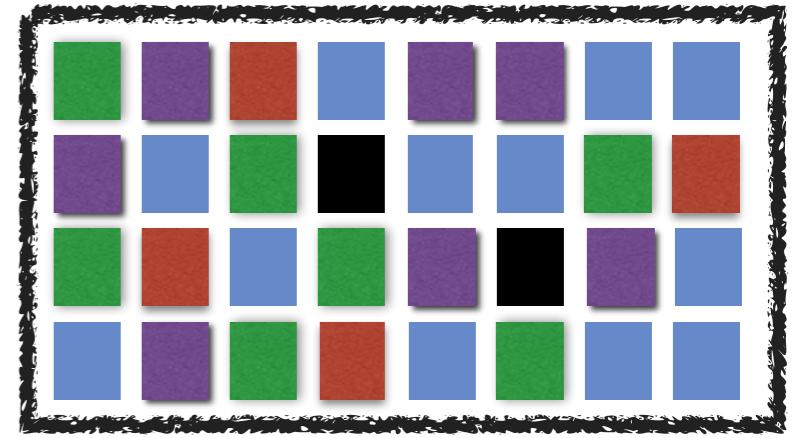
# Example: Search( ■ , file)



Search( ■ )



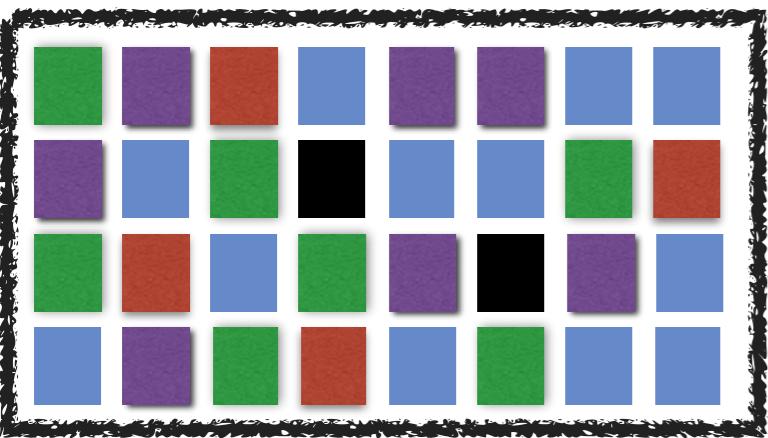
Data Scans



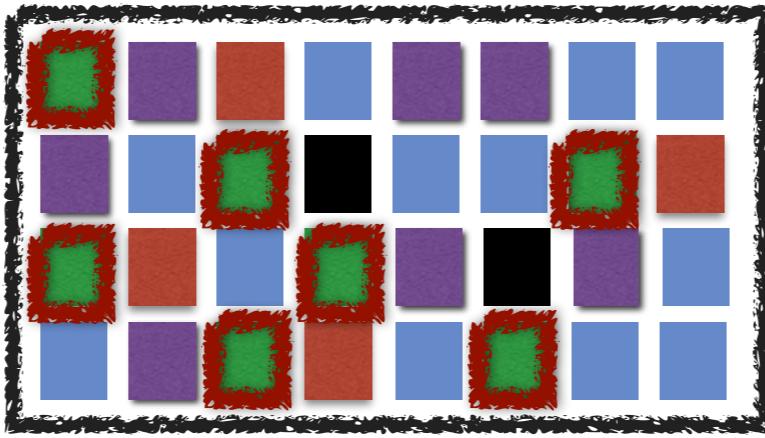
Indexes

Low storage  
High Latency

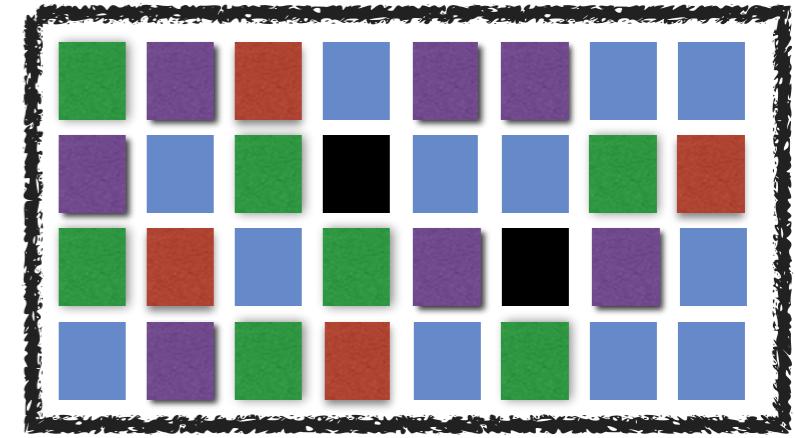
# Example: Search( ■ , file)



Search( ■ )



Data Scans

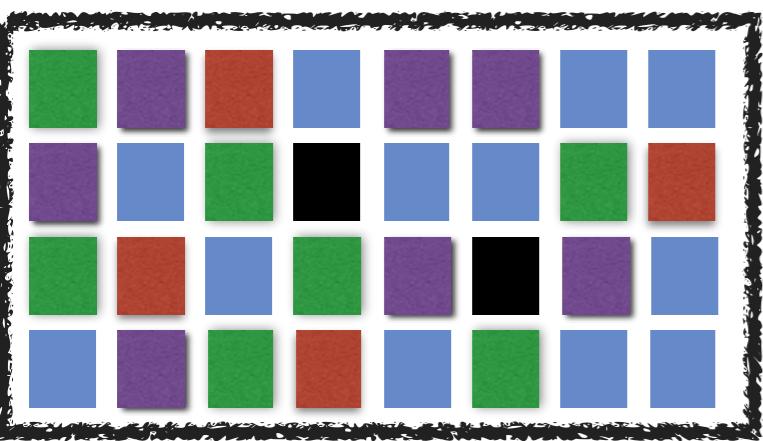


Indexes

Low storage  
High Latency

- 0, 10, 14, 16, 19, 26, 29
- 1, 4, 5, 8, 20, 22, 24
- 2, 15, 17, 27
- 3, 6, 7, 9, 12, 13, 18, 23 ..
- 11, 21

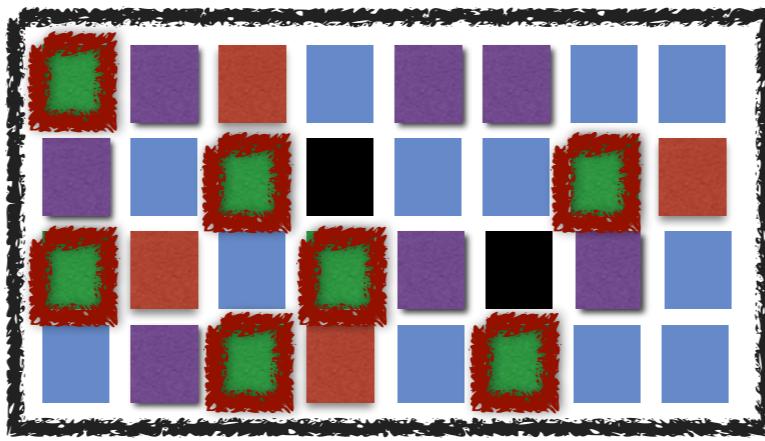
# Example: Search( ■ , file)



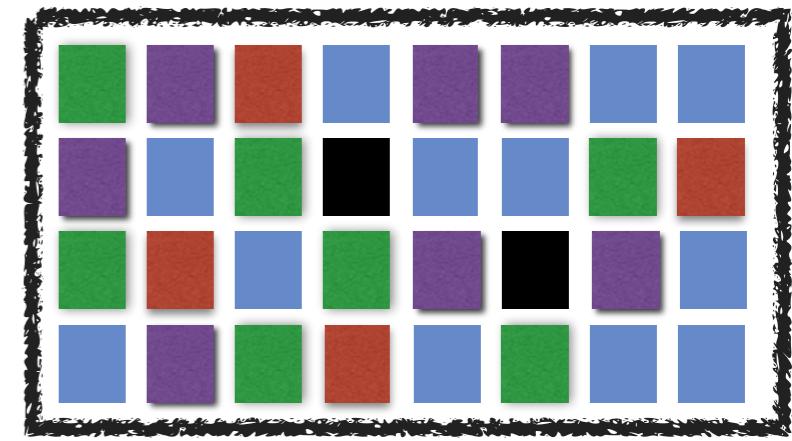
Search( ■ )

Low storage  
High Latency

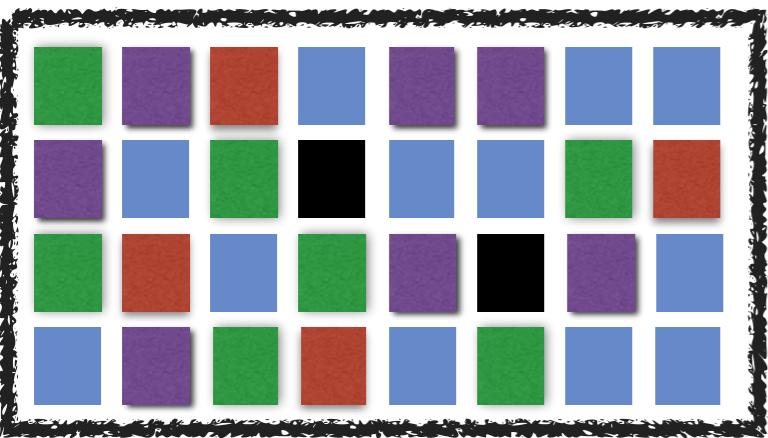
Data Scans



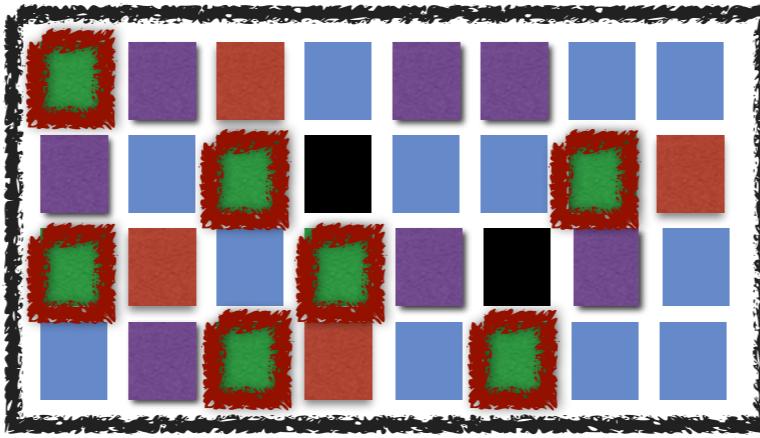
Indexes



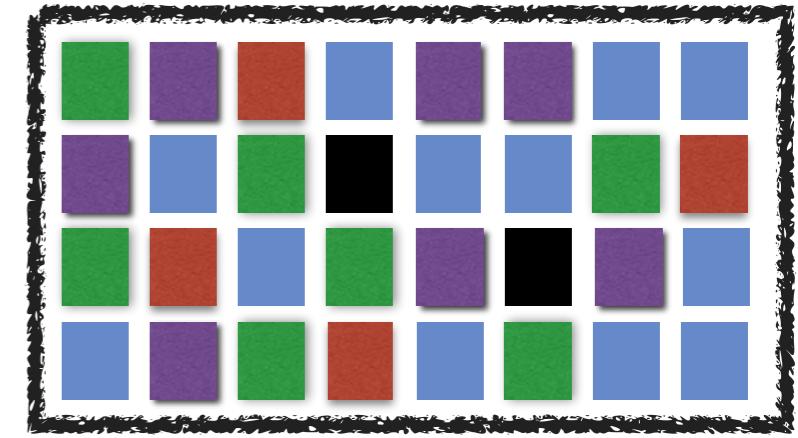
# Example: Search( ■ , file)



Search( ■ )



Data Scans



Indexes

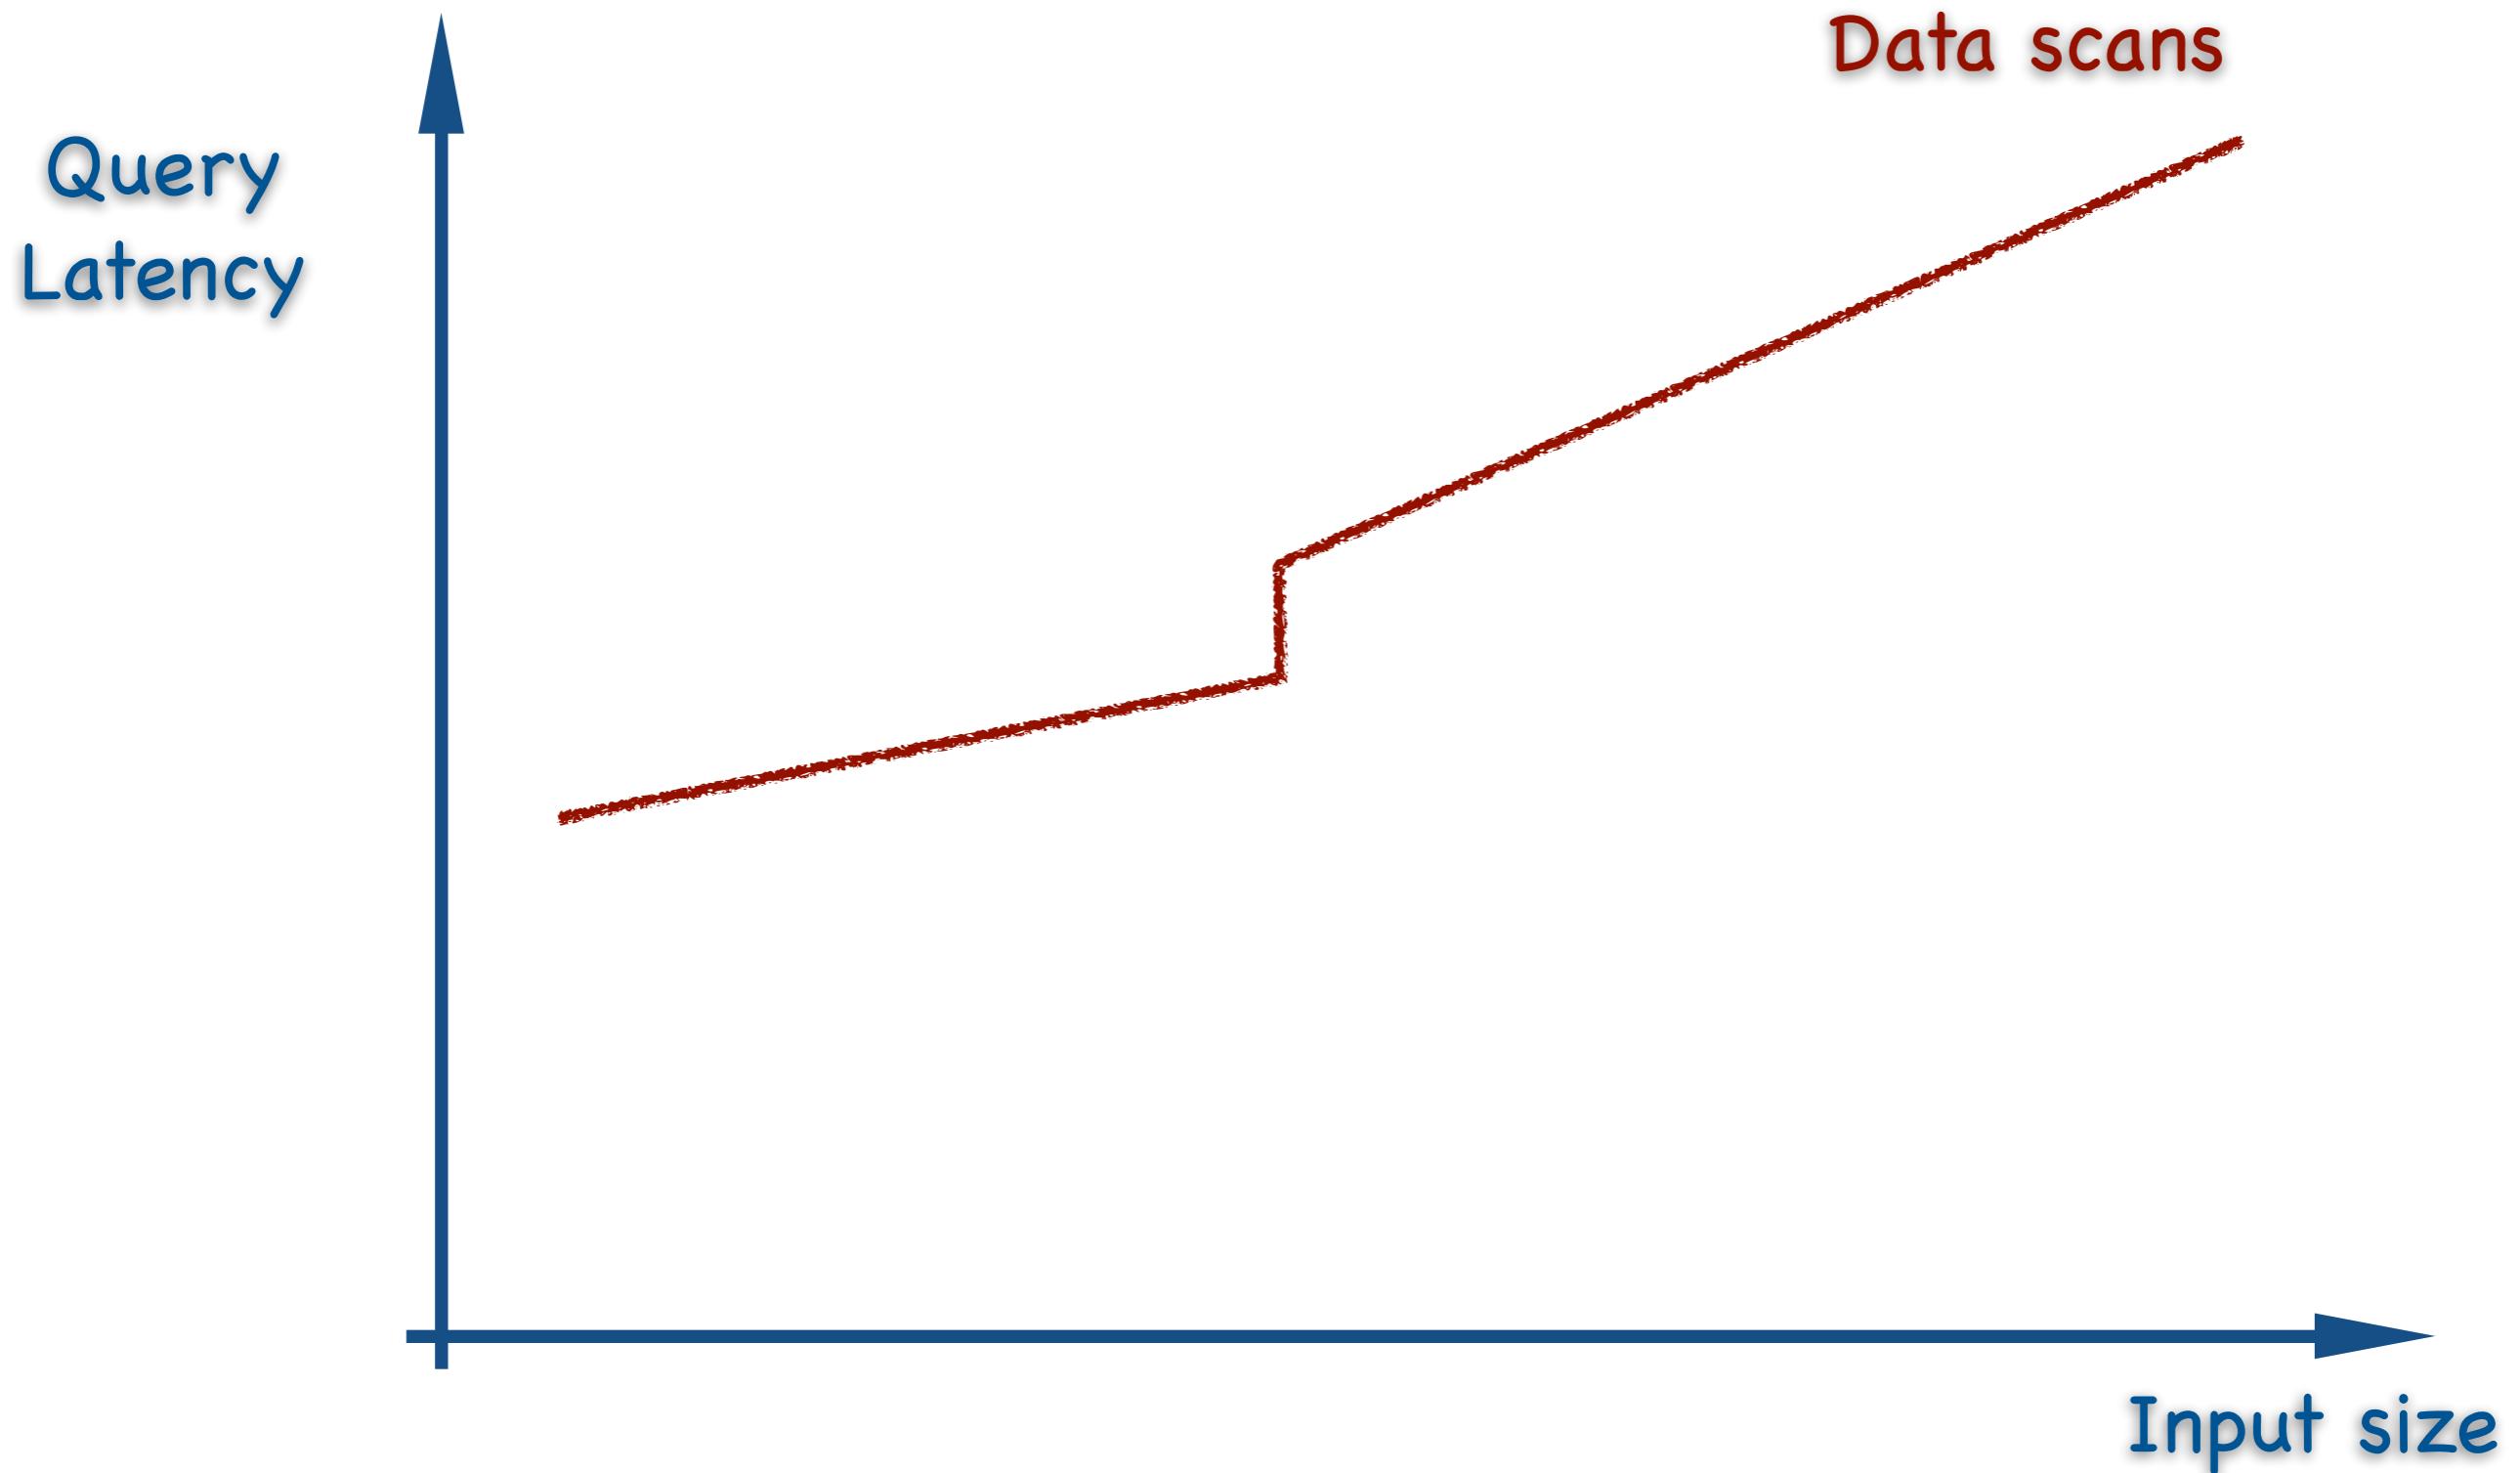
Low storage  
High Latency

■	0, 10, 14, 16, 19, 26, 29
■	1, 4, 5, 8, 20, 22, 24
■	2, 15, 17, 27
■	3, 6, 7, 9, 12, 13, 18, 23 ..
■	11, 21

High storage  
Low Latency

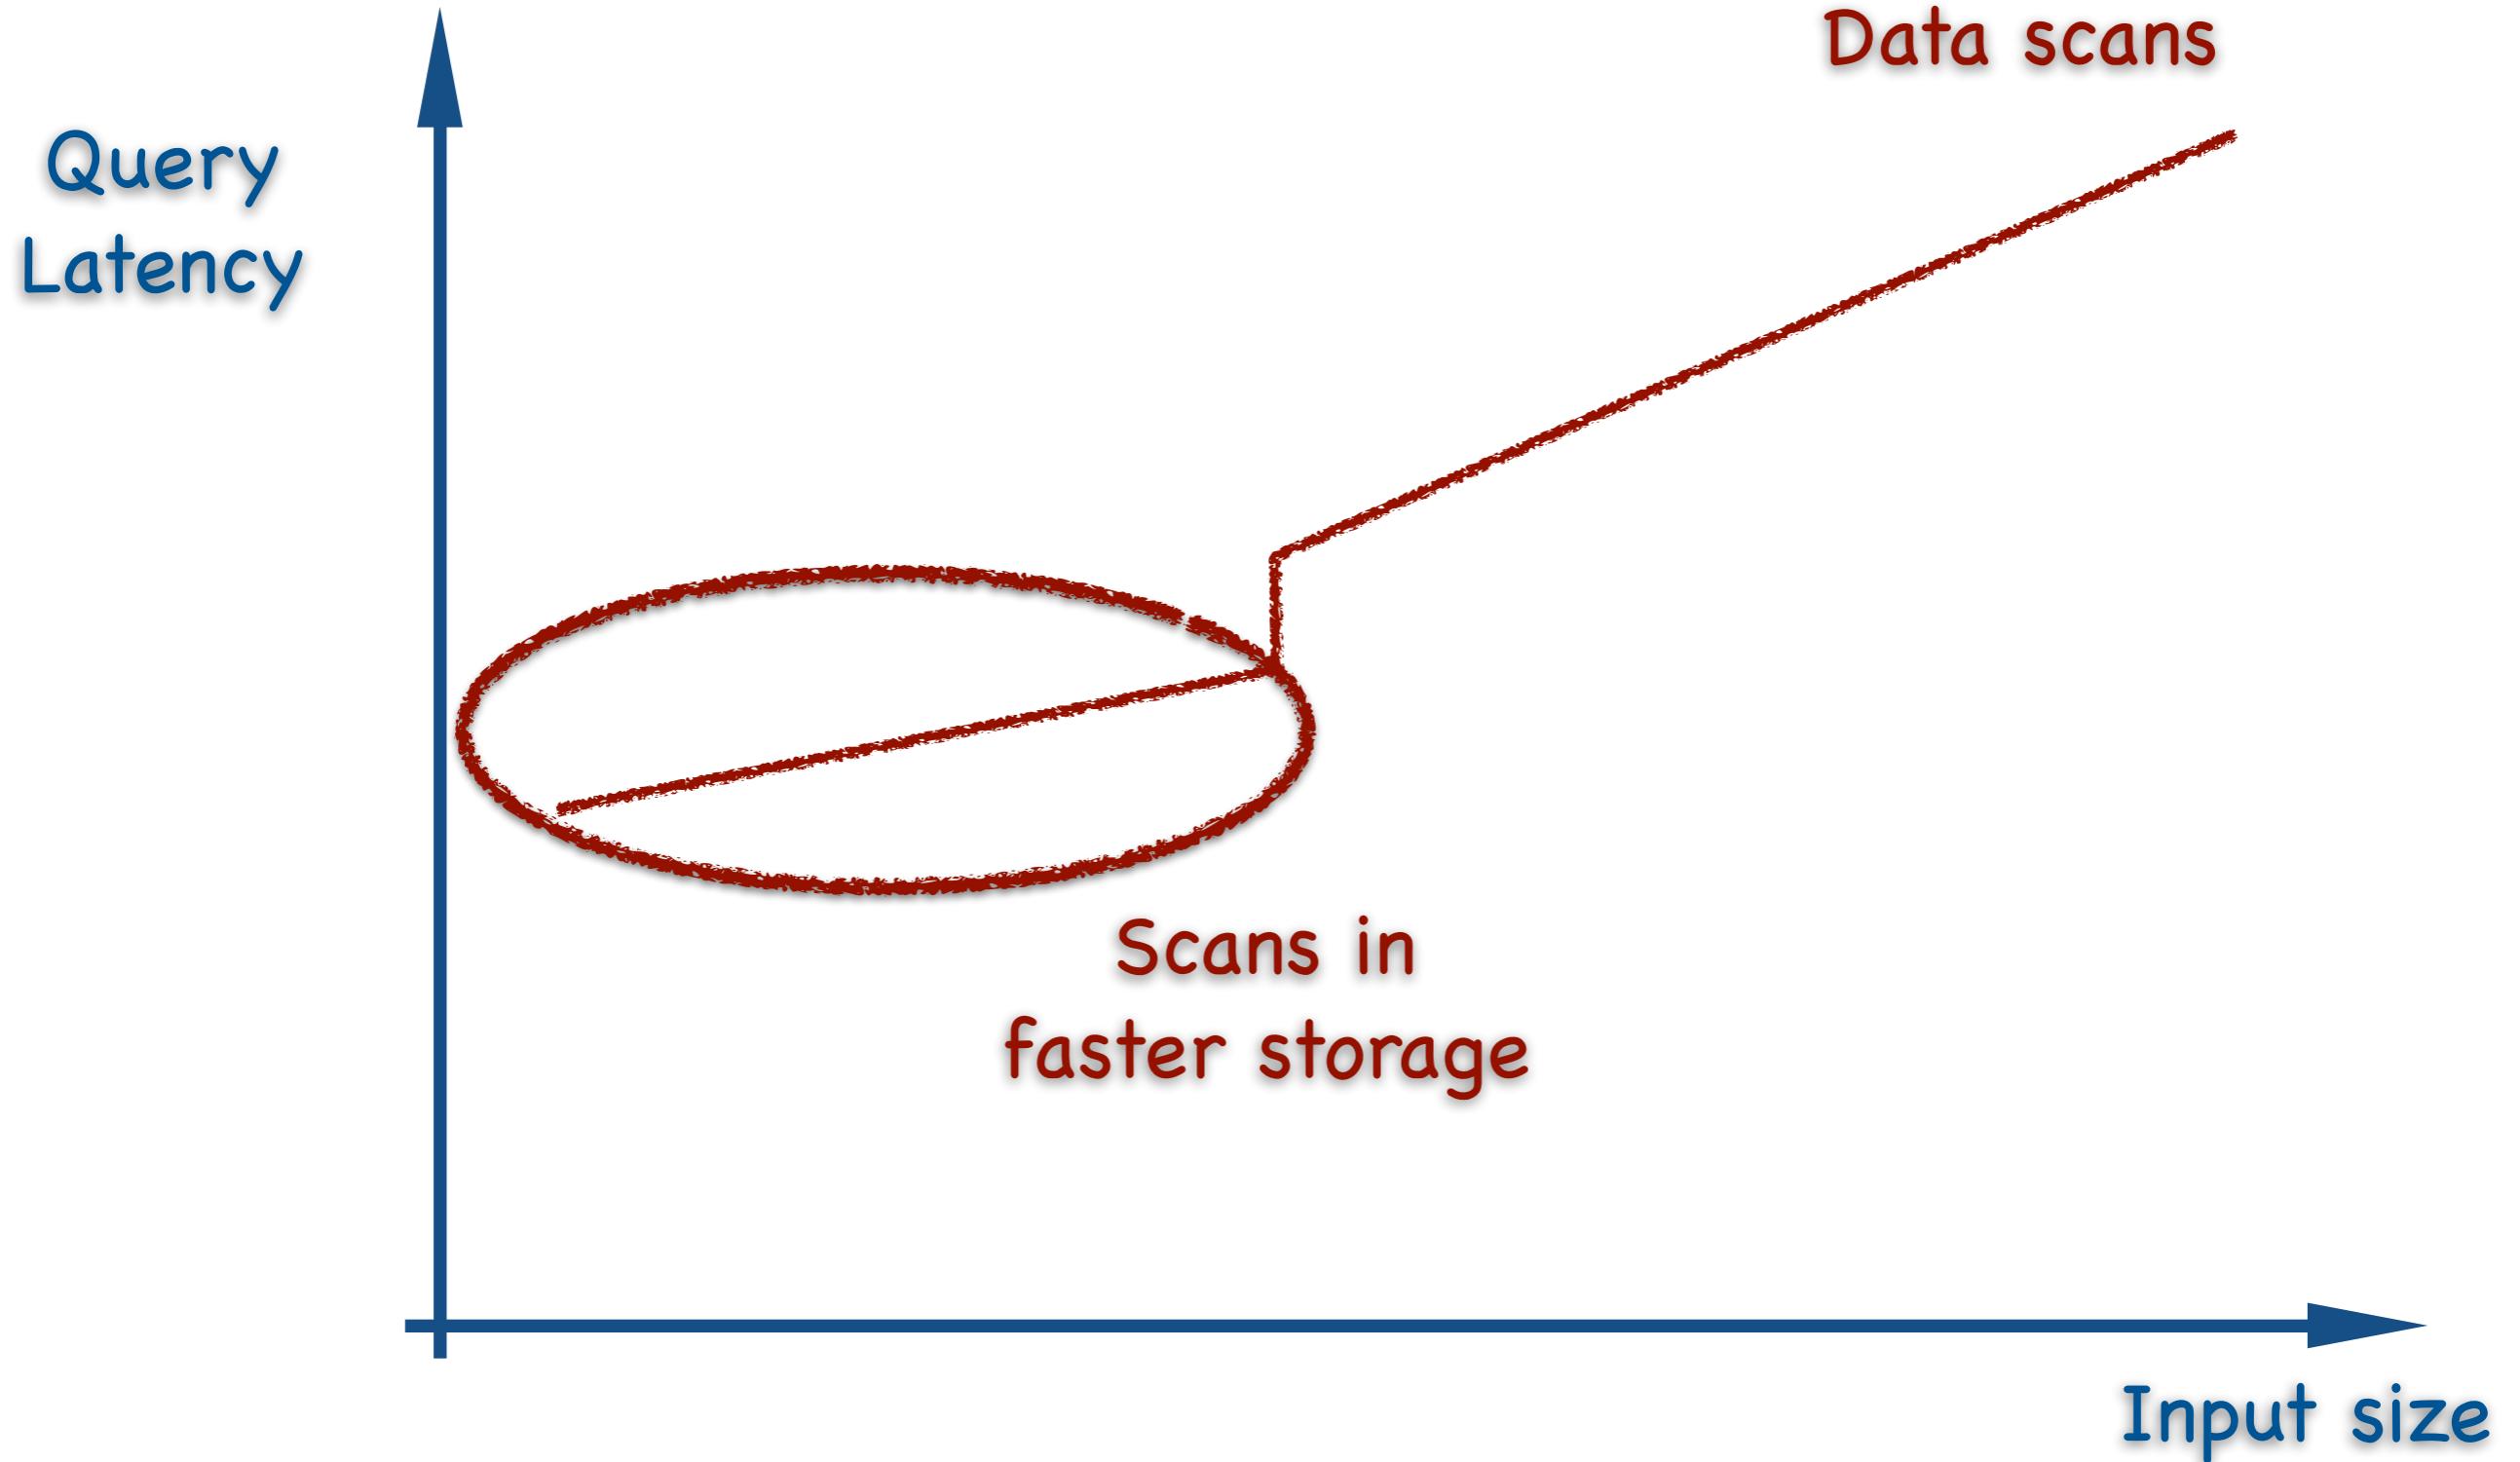
# Latency with increasing input size

---



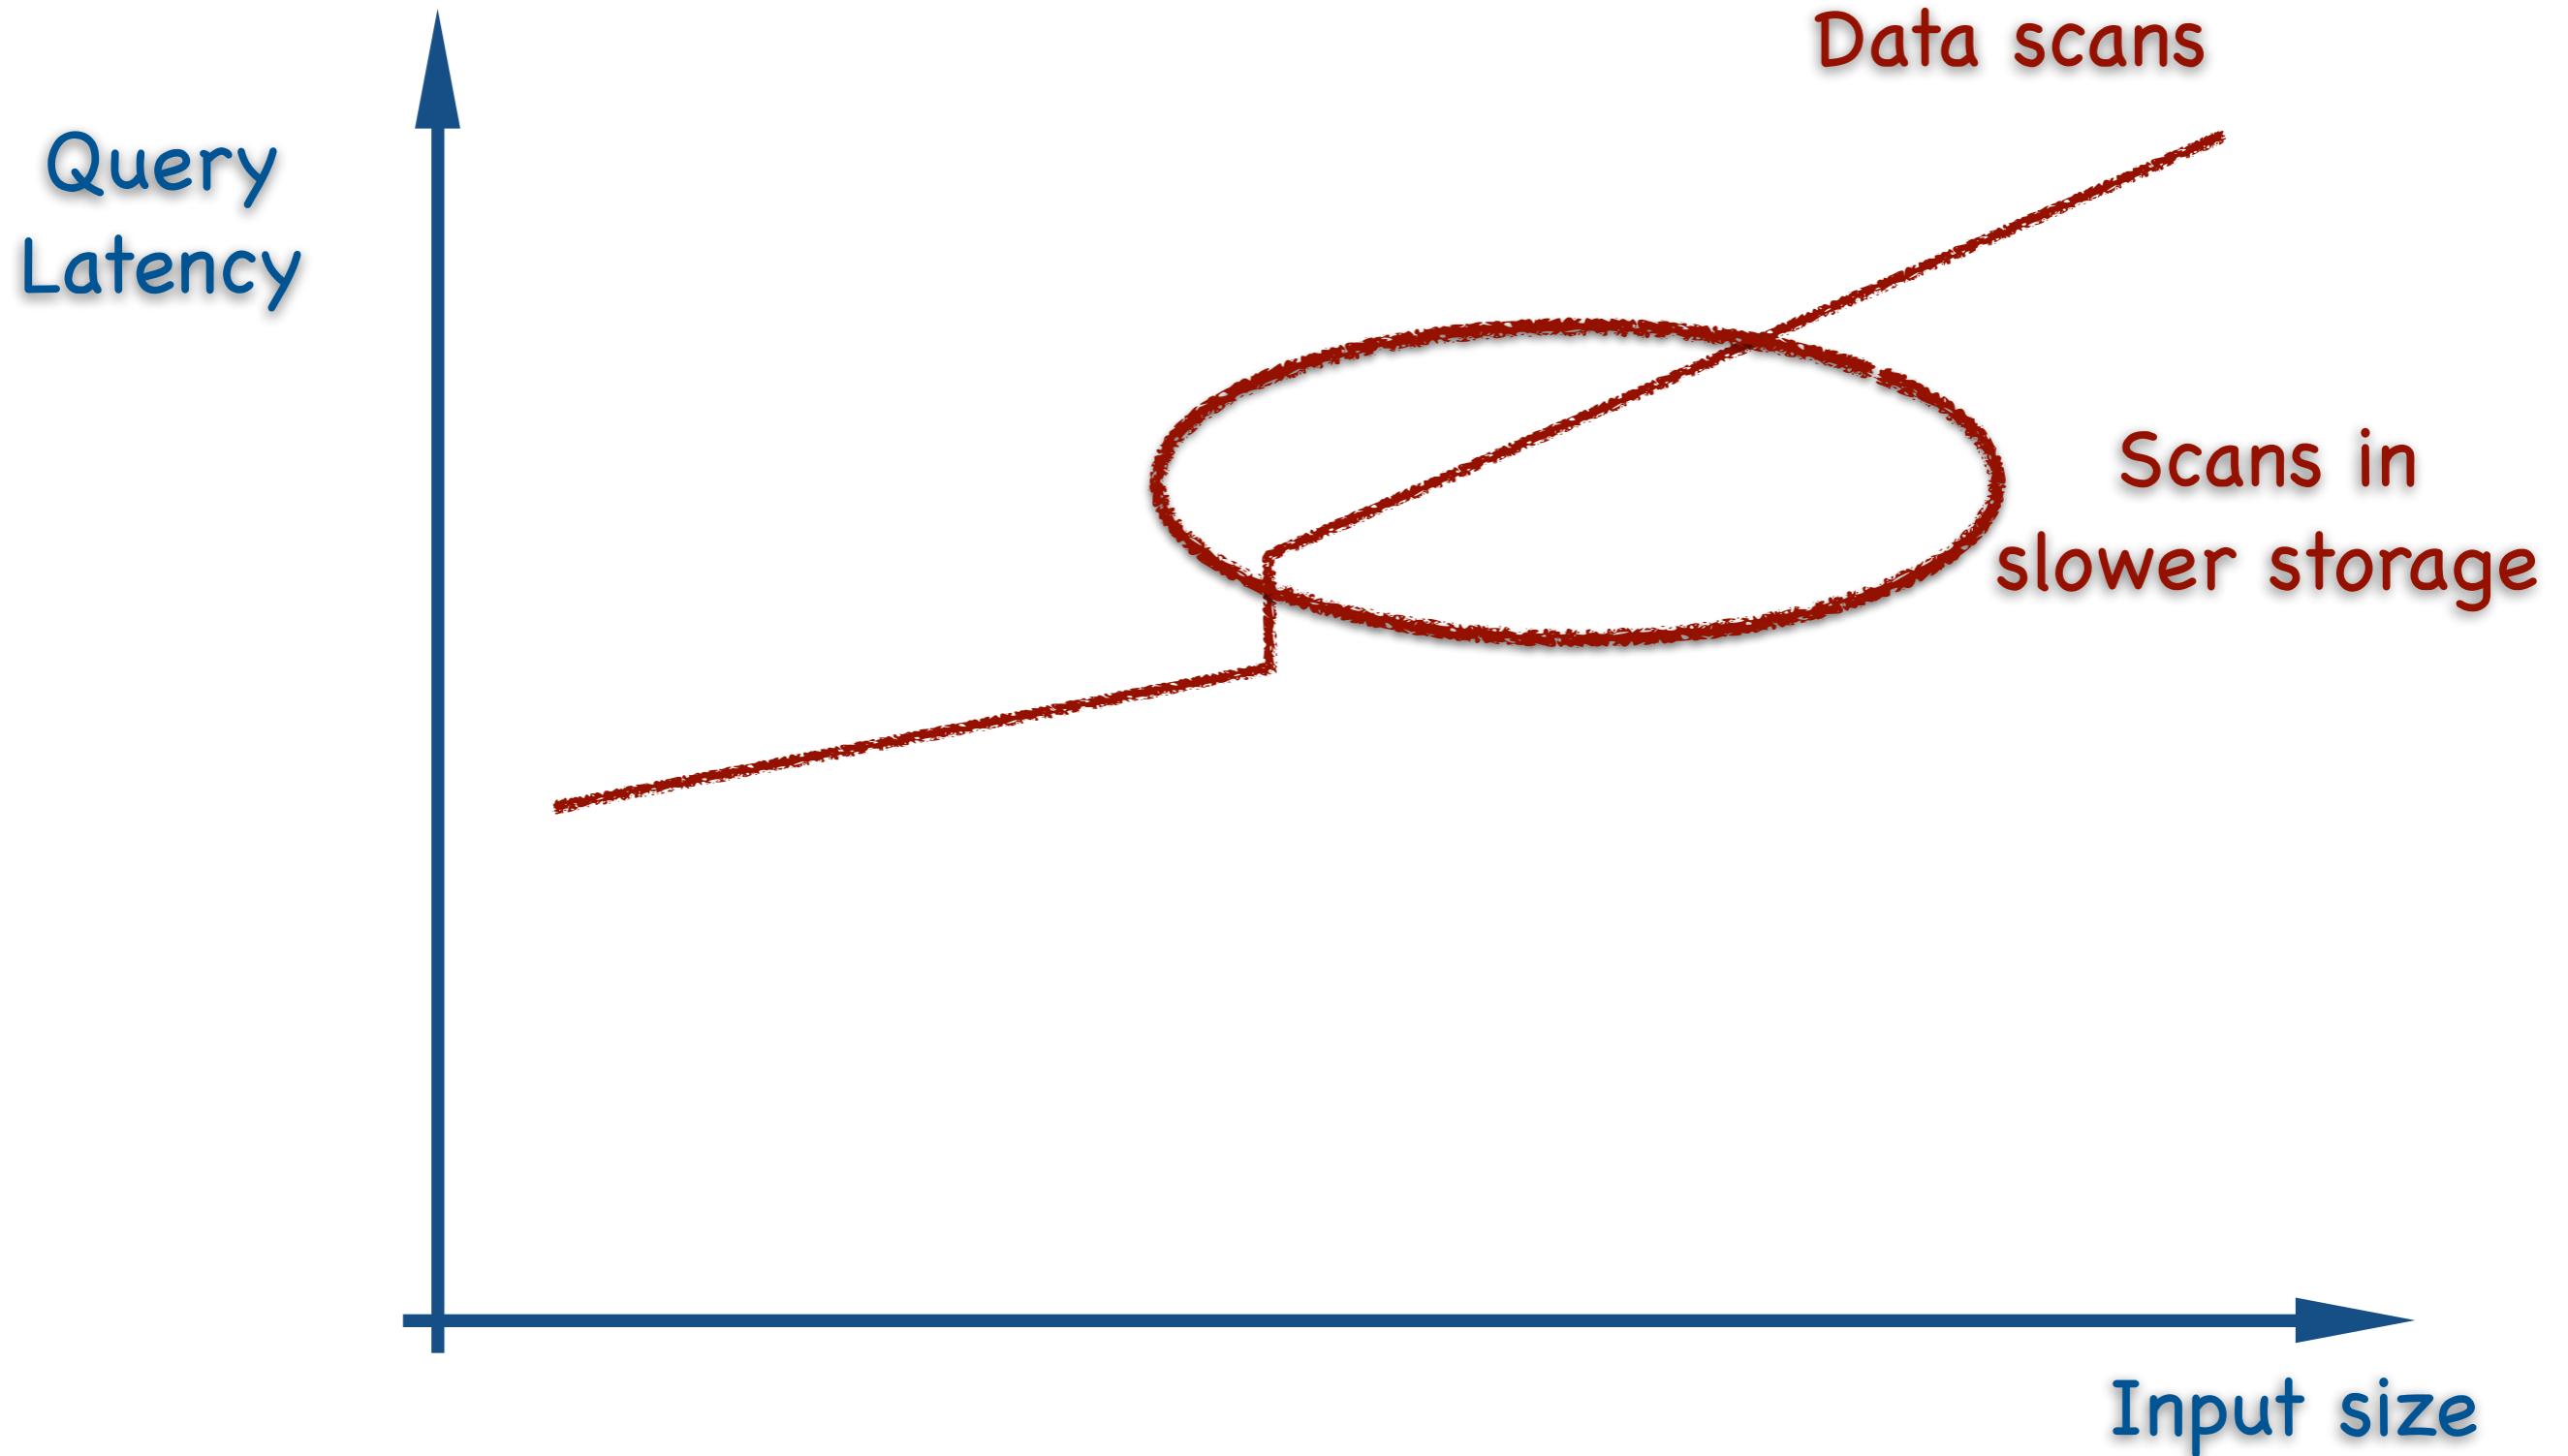
# Latency with increasing input size

---



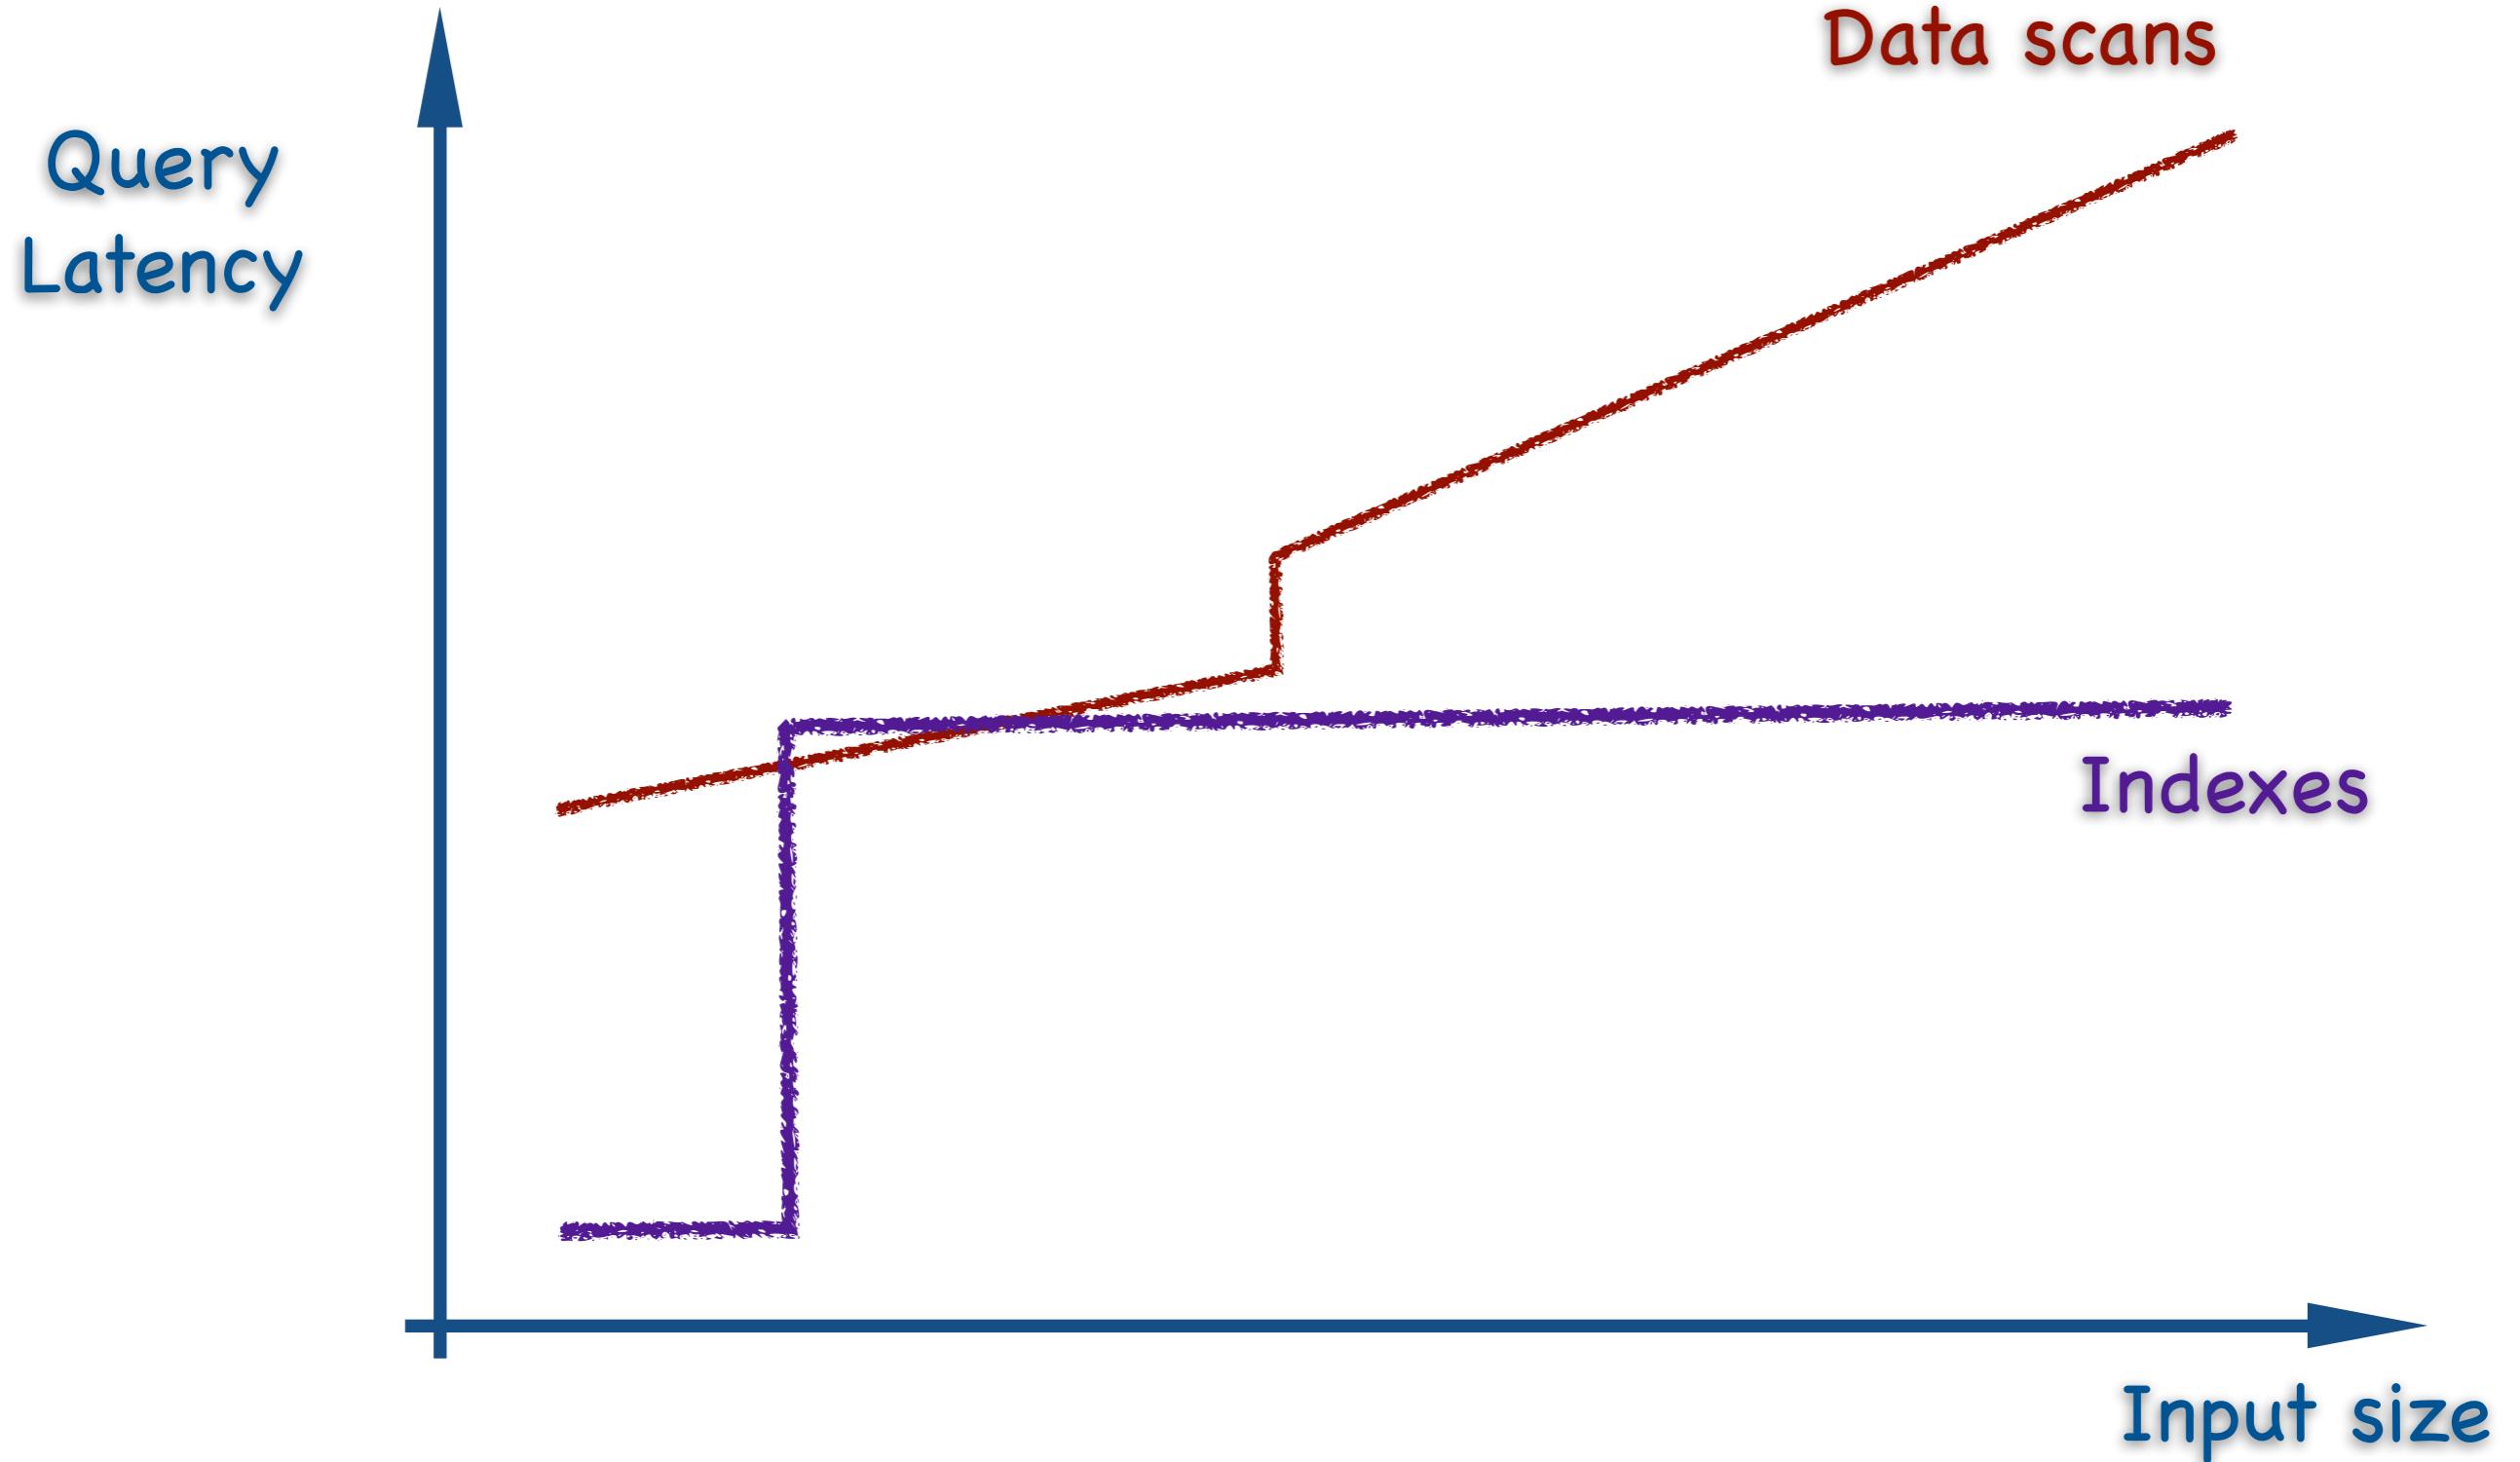
# Latency with increasing input size

---

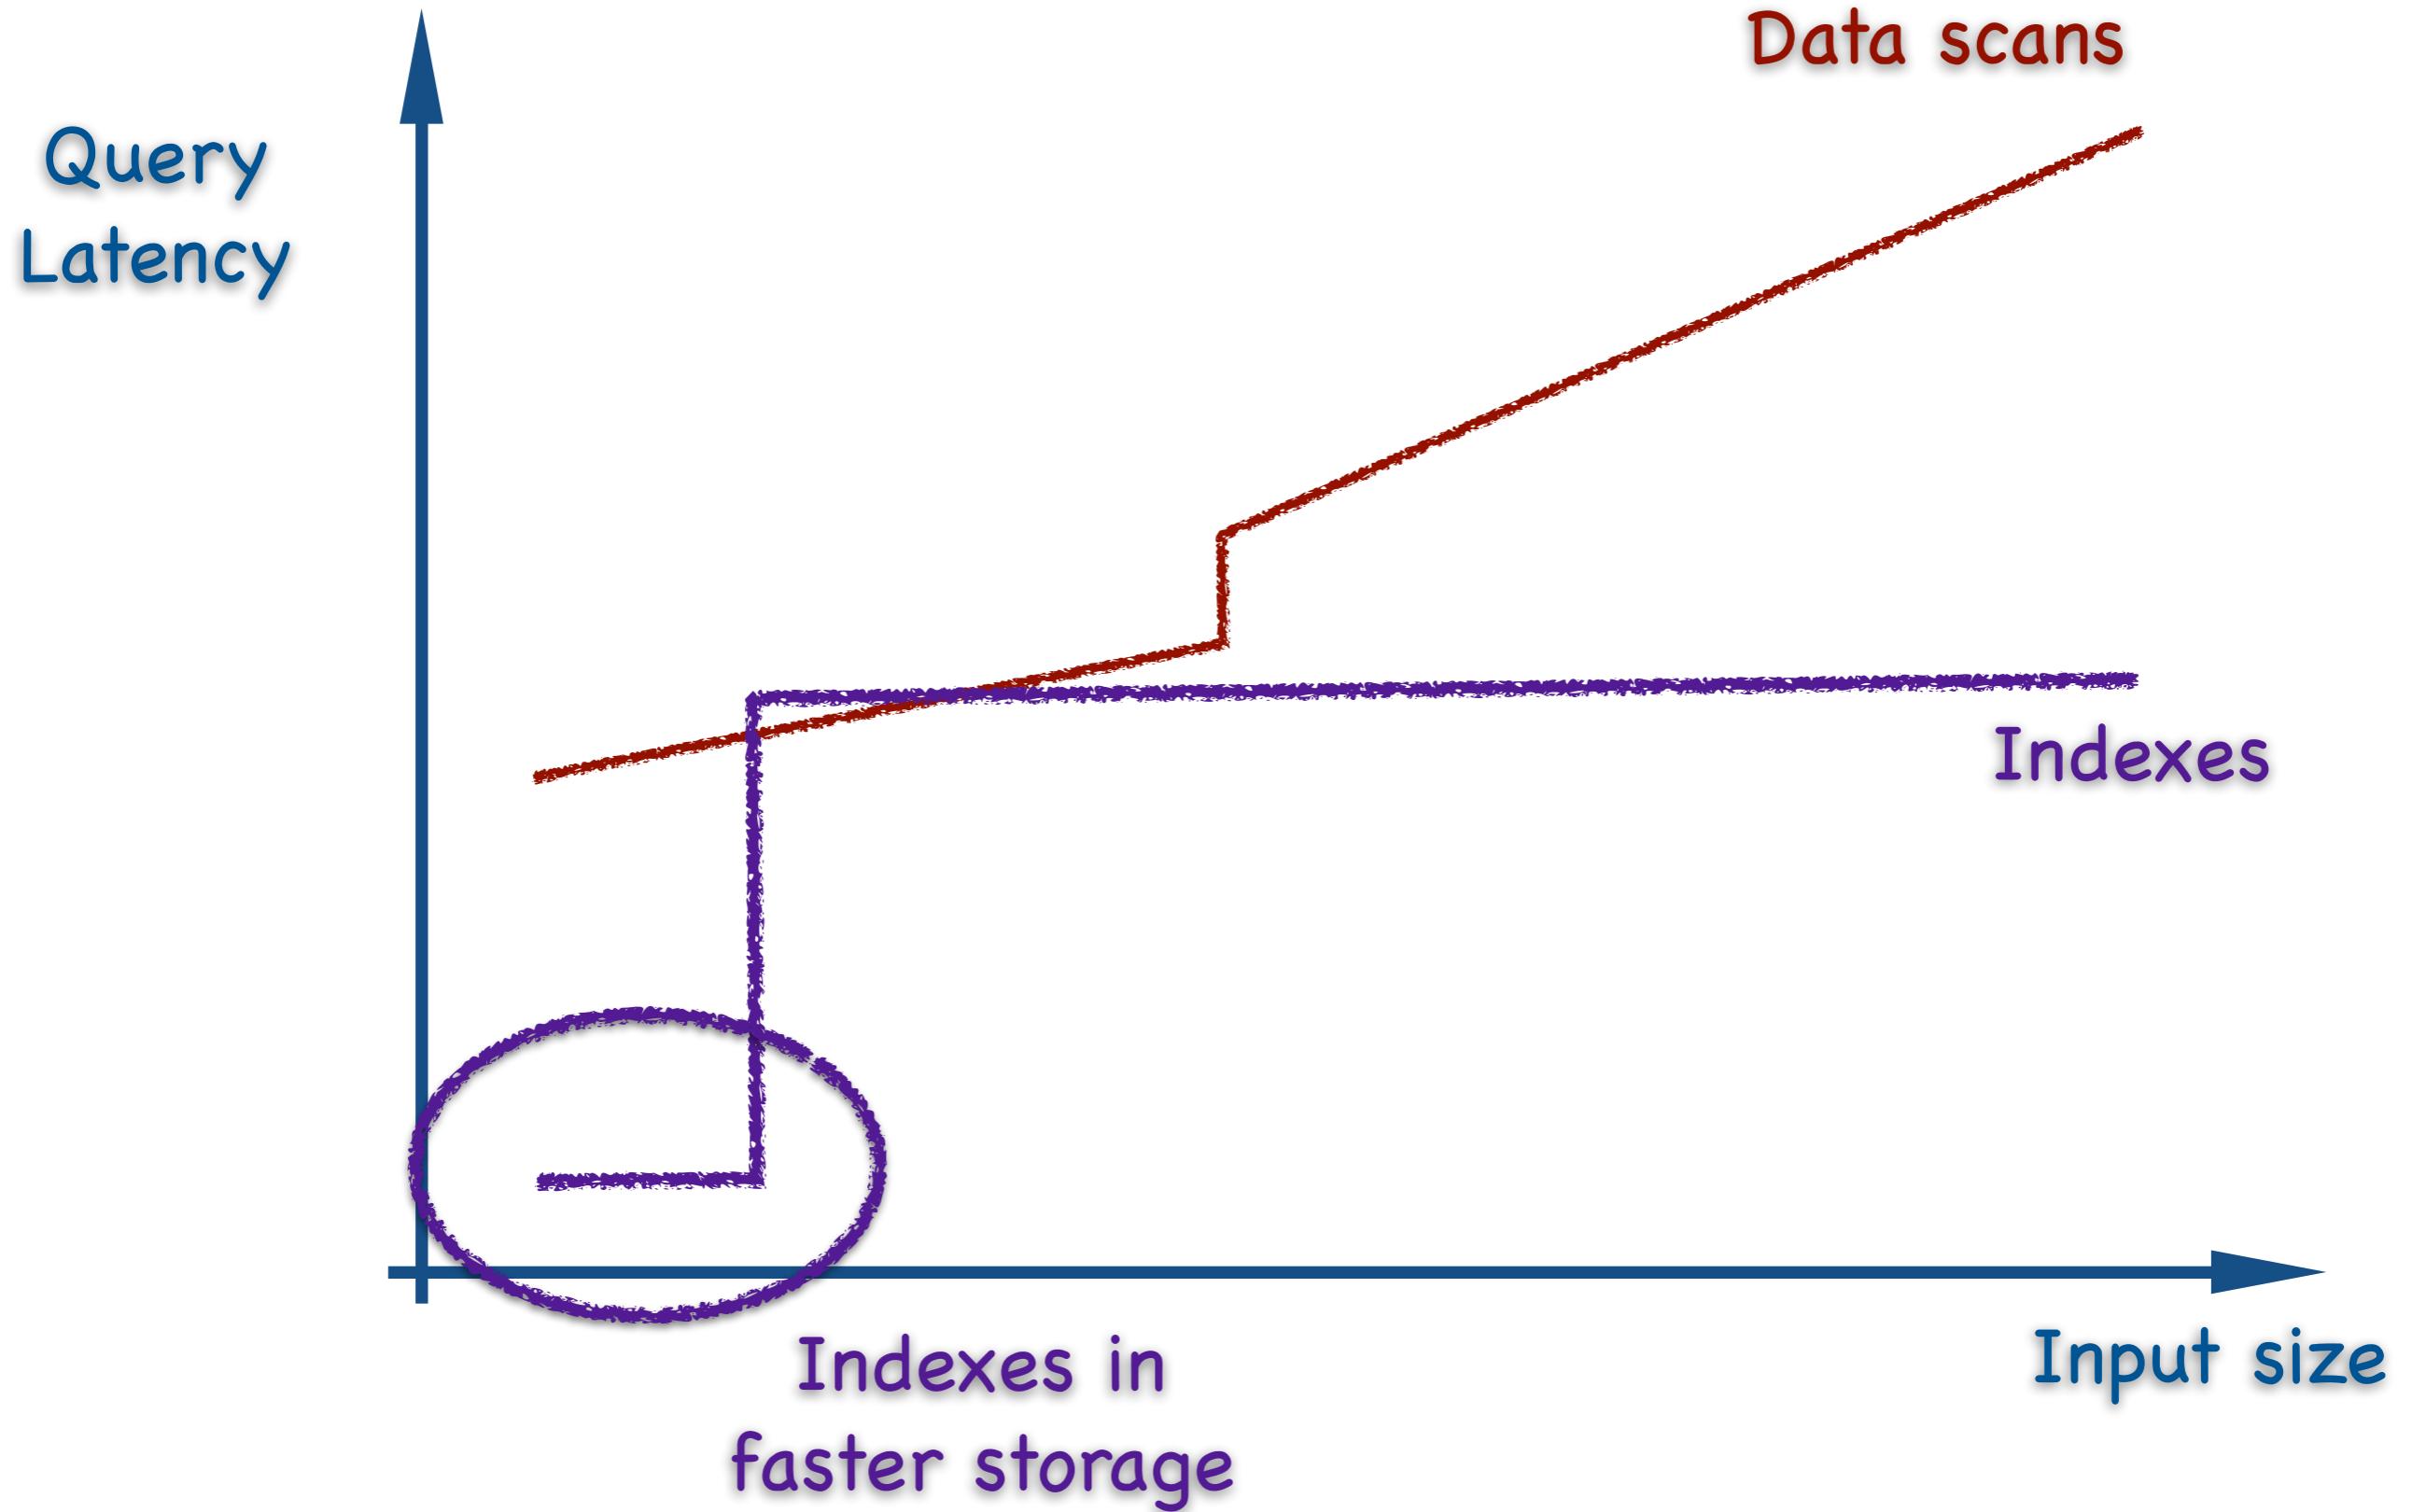


# Latency with increasing input size

---

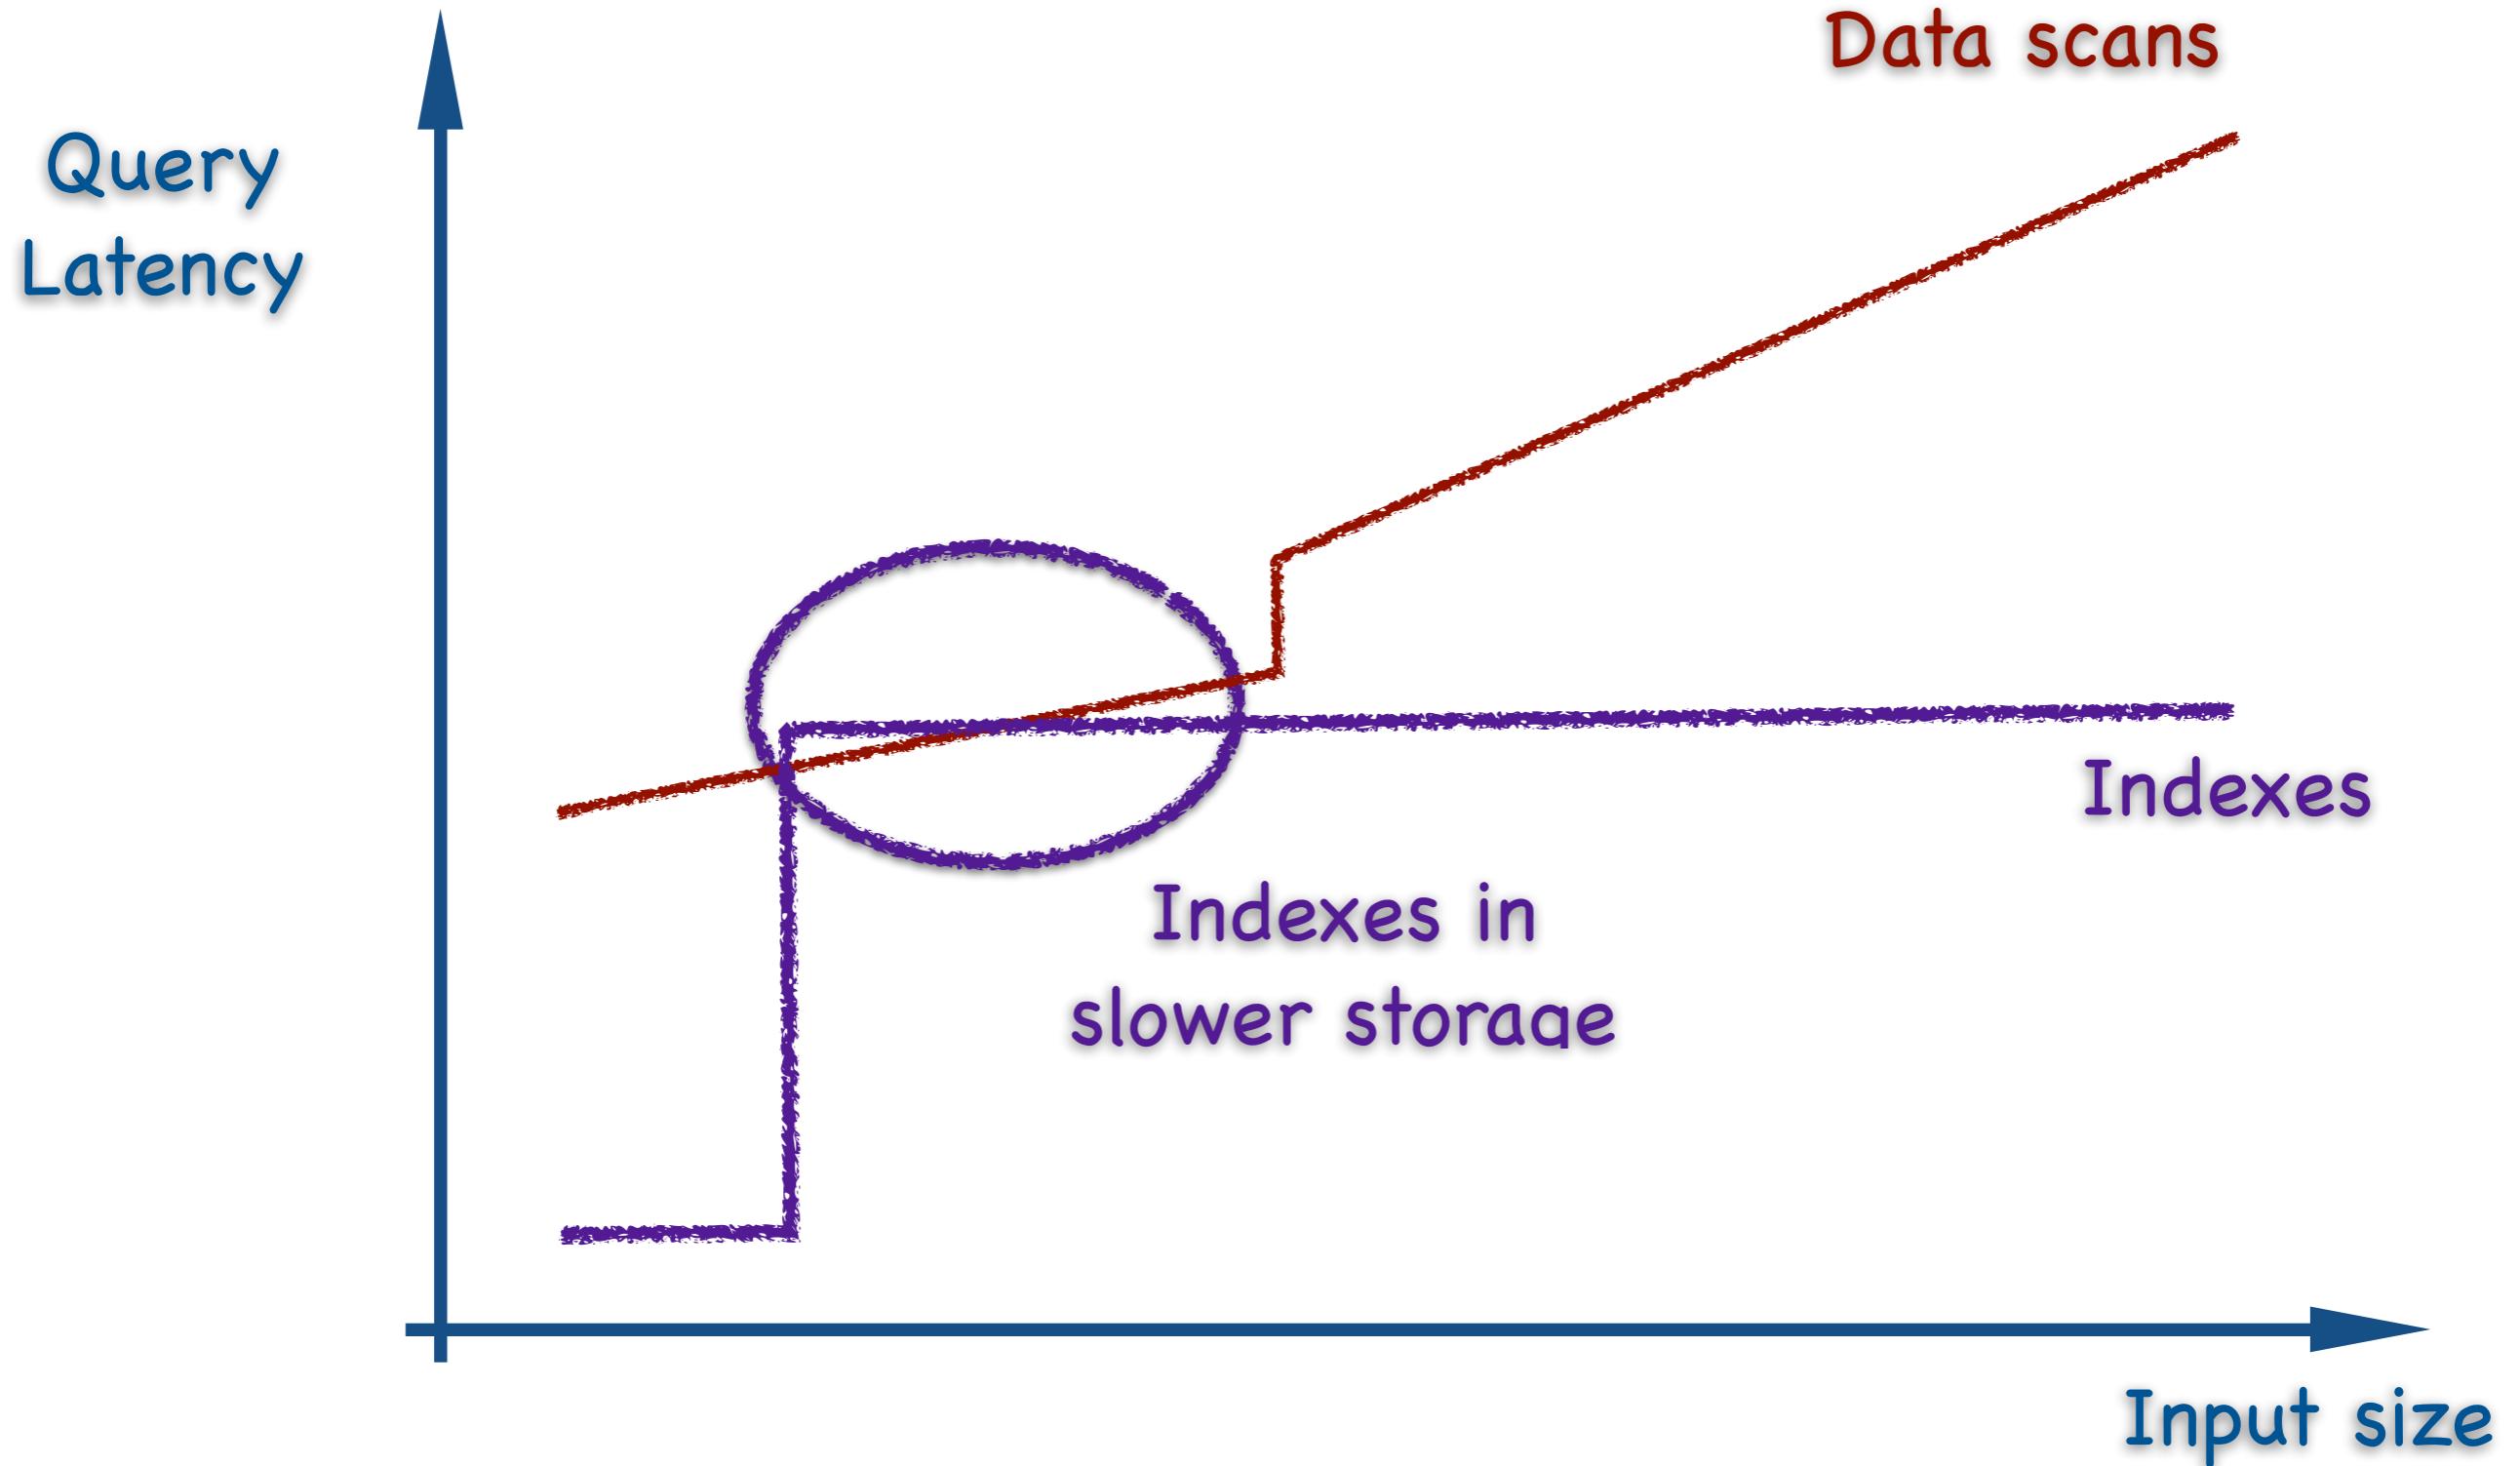


# Latency with increasing input size

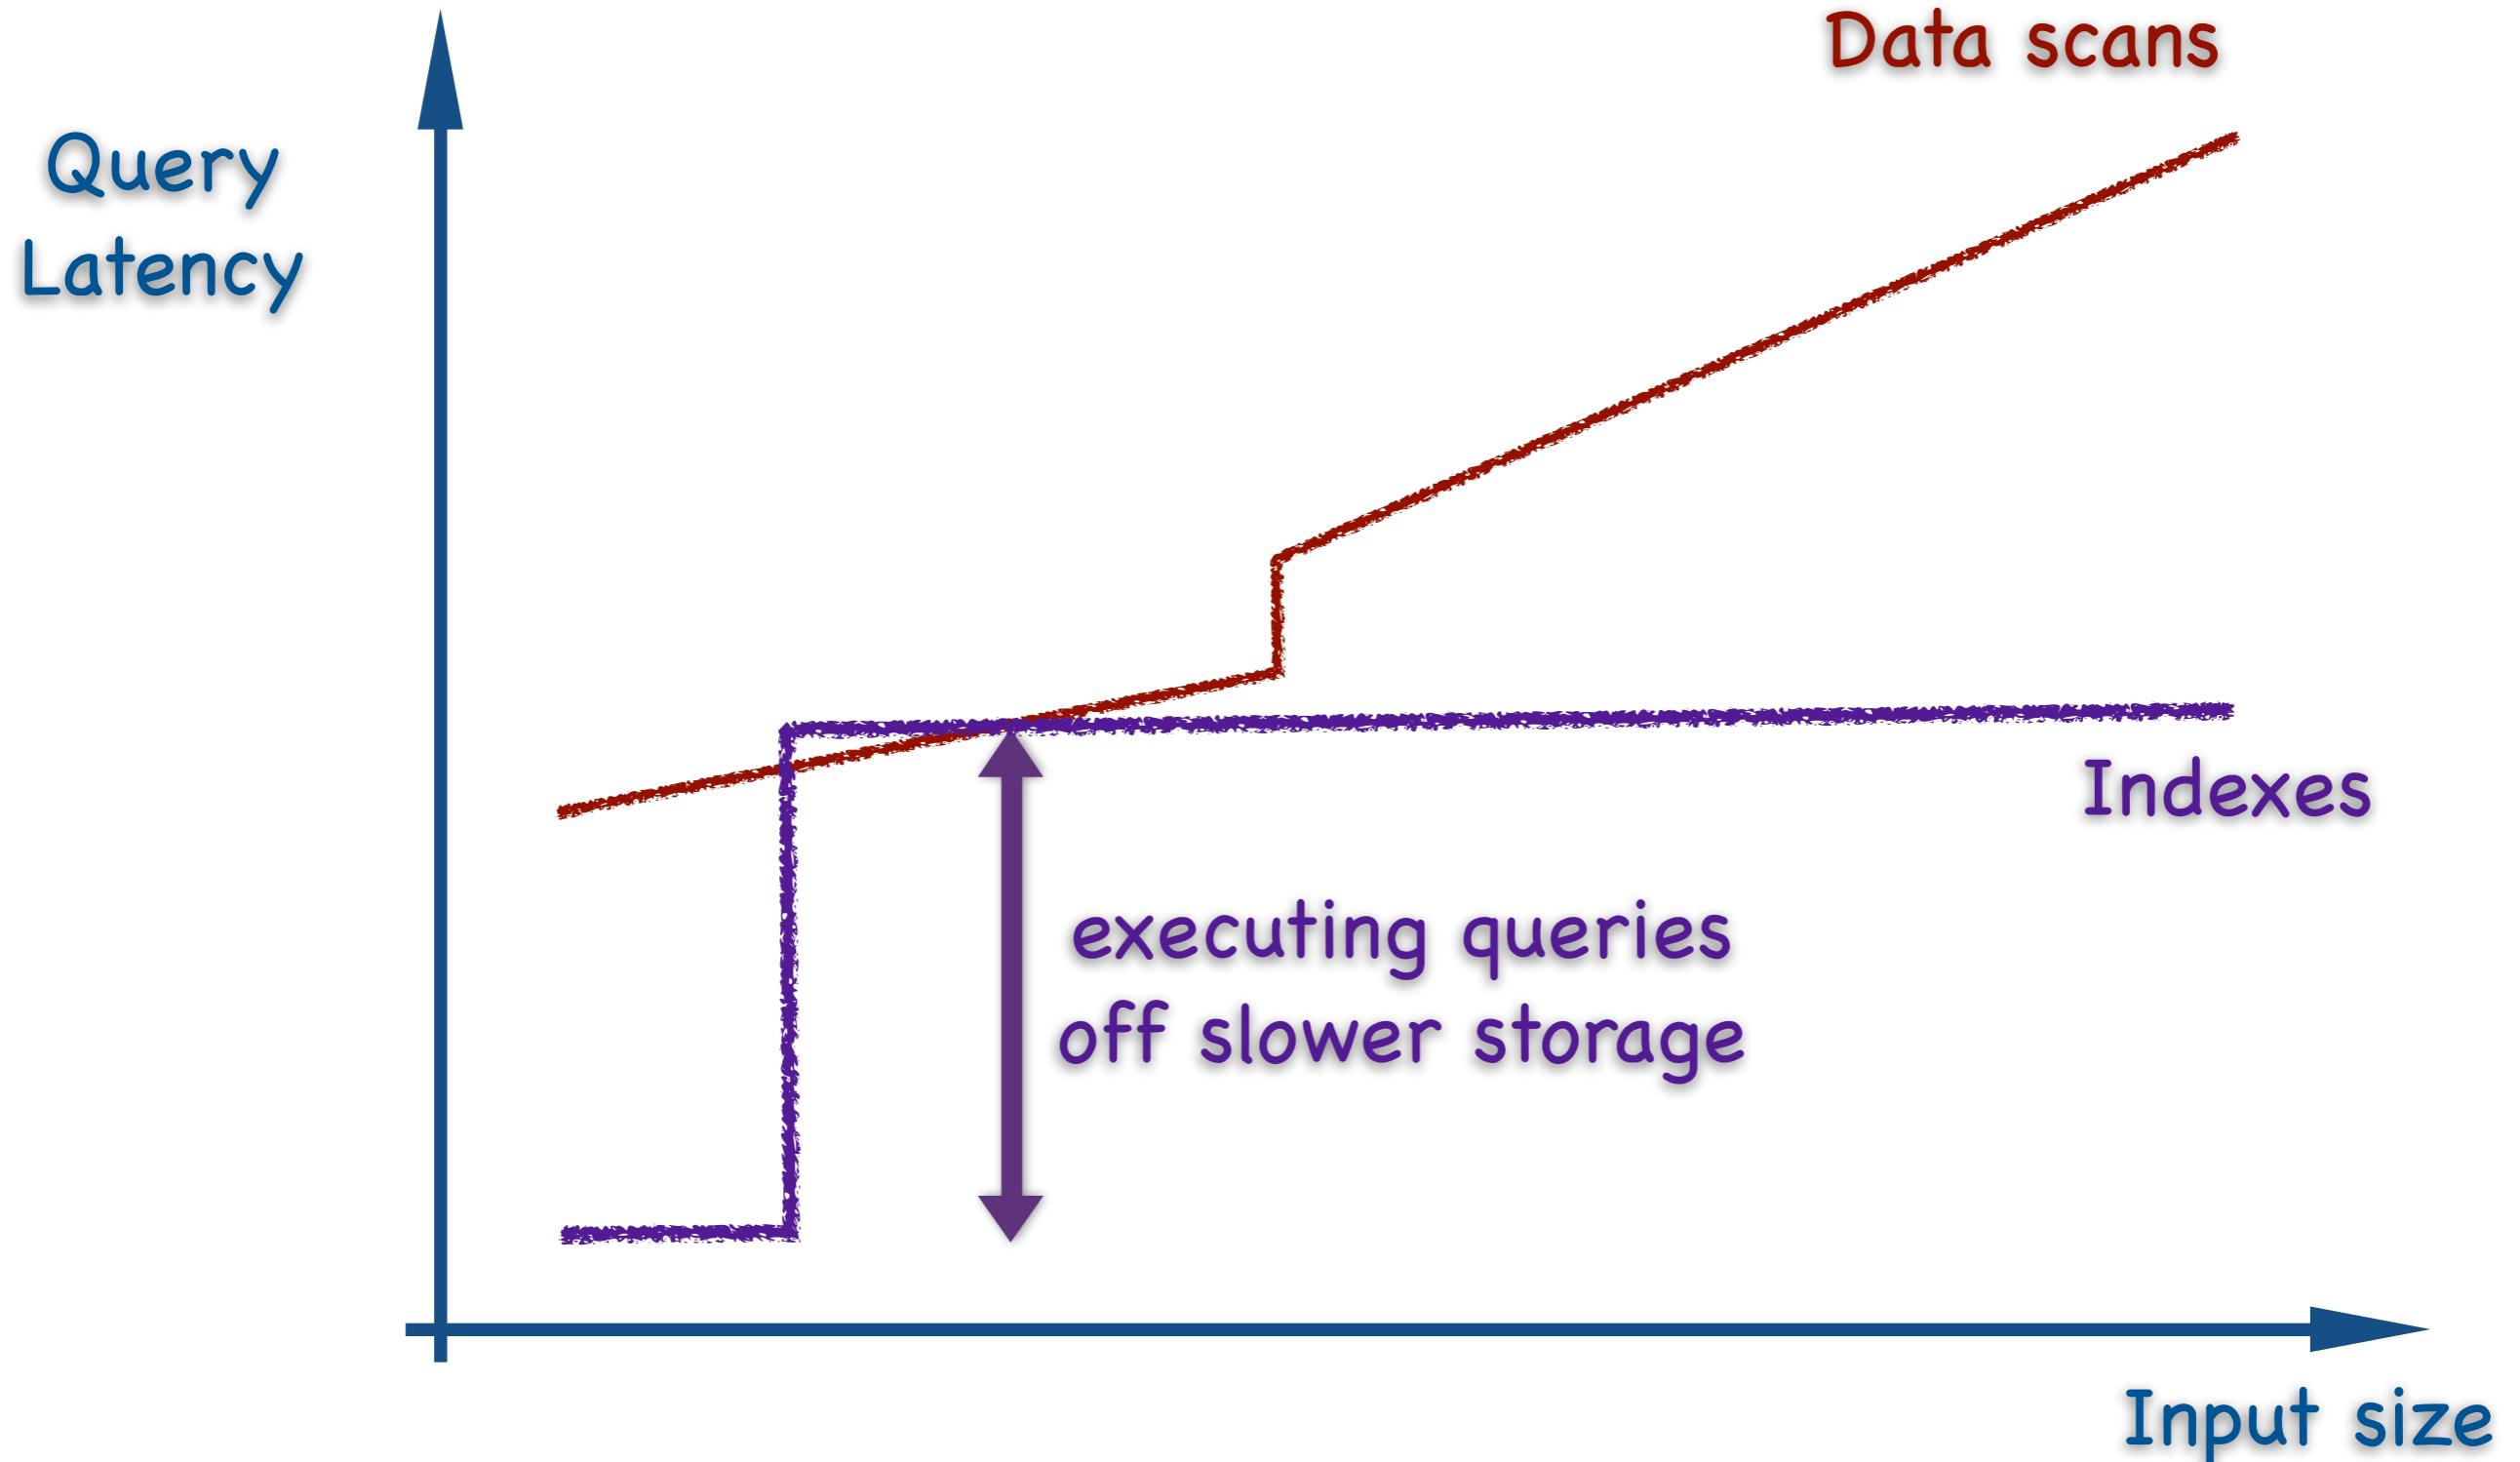


# Latency with increasing input size

---



# Latency with increasing input size



# Succinct

---

Push more data in faster storage

- Storage < Input data size

# Succinct

---

Push more data in faster storage

- Storage < Input data size

Execute queries directly on compressed data

- No data scans
- No data decompression

# Succinct

---

Push more data in faster storage

- Storage < Input data size

Execute queries directly on compressed data

- No data scans
- No data decompression

Target: read-heavy append-only workloads

# Succinct

---

Push more data in faster storage

- Storage < Input data size

Execute queries directly on compressed data

- No data scans
- No data decompression

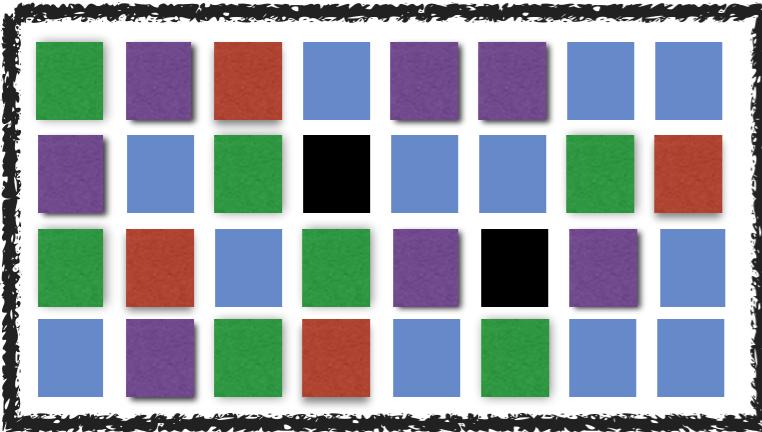
Target: read-heavy append-only workloads

Non-goals: Fault tolerance, data consistency

# Example: Search( , file)

---

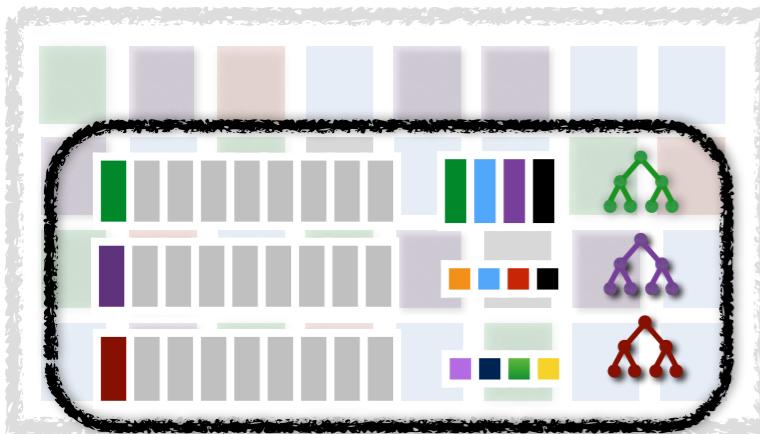
Succinct



# Example: Search( , file)

---

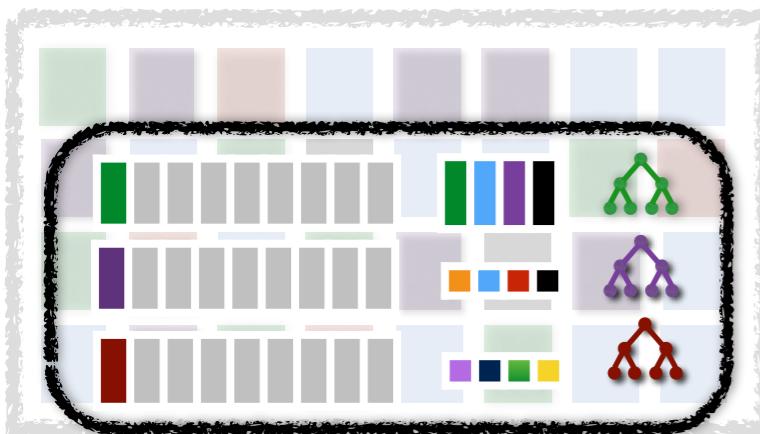
Succinct



# Example: Search( , file)

---

Succinct

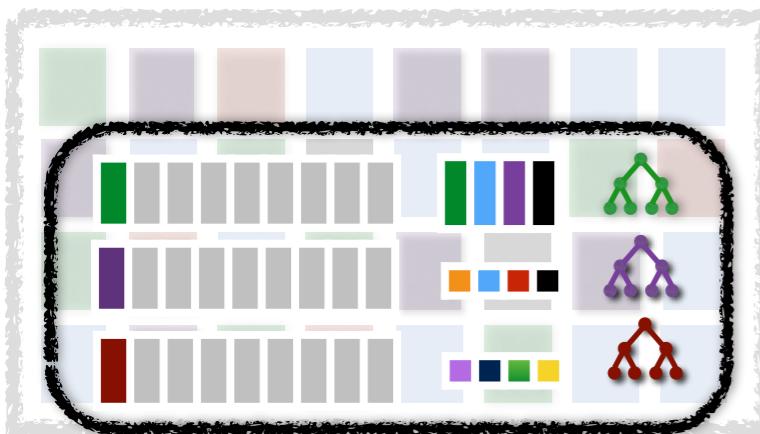


Queries executed  
directly on the  
compressed representation

# Example: Search( , file)

---

Succinct



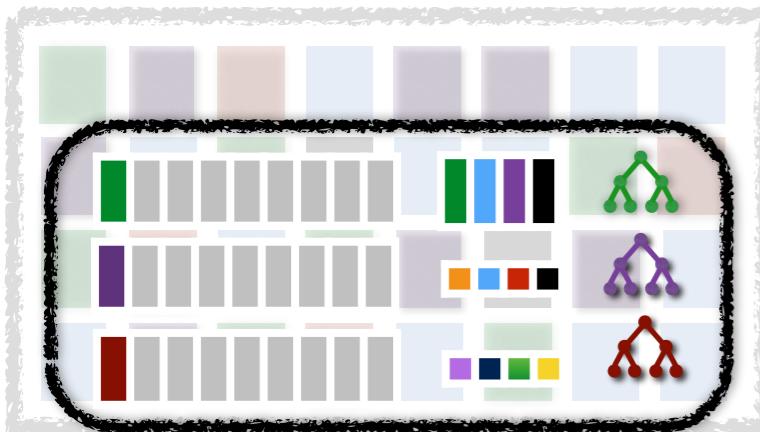
Queries executed  
directly on the  
compressed representation

Low storage  
Low Latency

# Example: Search( , file)

---

Succinct



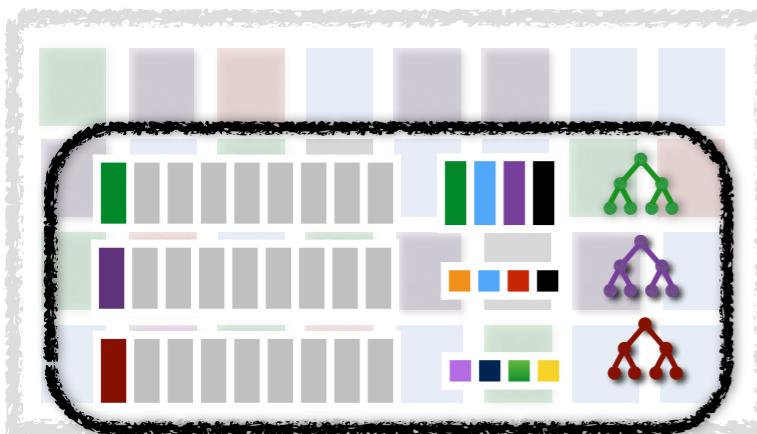
Queries executed  
directly on the  
compressed representation

Low storage  
Low Latency

What makes Succinct unique

# Example: Search( , file)

Succinct



Queries executed  
directly on the  
compressed representation

Low storage  
Low Latency

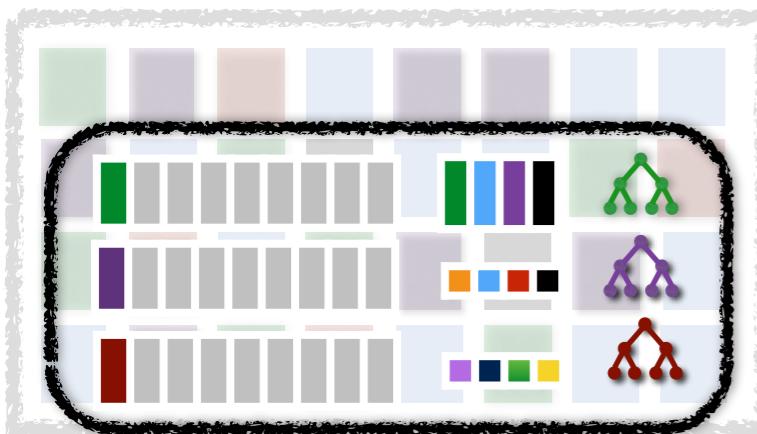
What makes Succinct unique

No additional  
indexes

Query responses  
embedded within  
the compressed representation

# Example: Search( , file)

Succinct



Queries executed  
directly on the  
compressed representation

Low storage  
Low Latency

What makes Succinct unique

No additional  
indexes

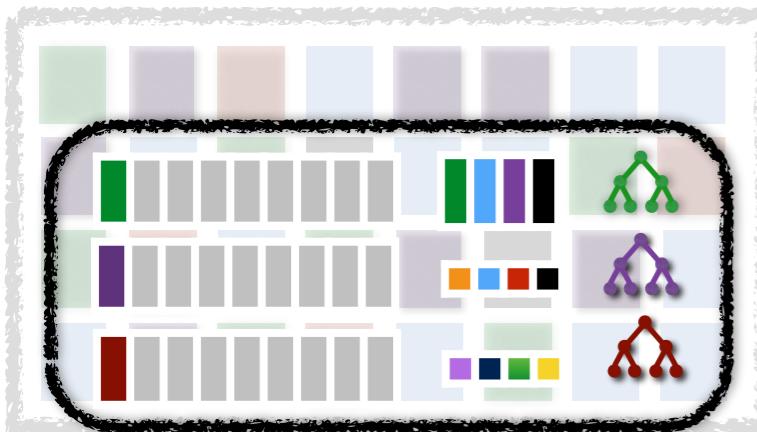
Query responses  
embedded within  
the compressed representation

No data scans

Functionality of indexes

# Example: Search( , file)

Succinct



Queries executed  
directly on the  
compressed representation

Low storage  
Low Latency

What makes Succinct unique

No additional  
indexes

Query responses  
embedded within  
the compressed representation

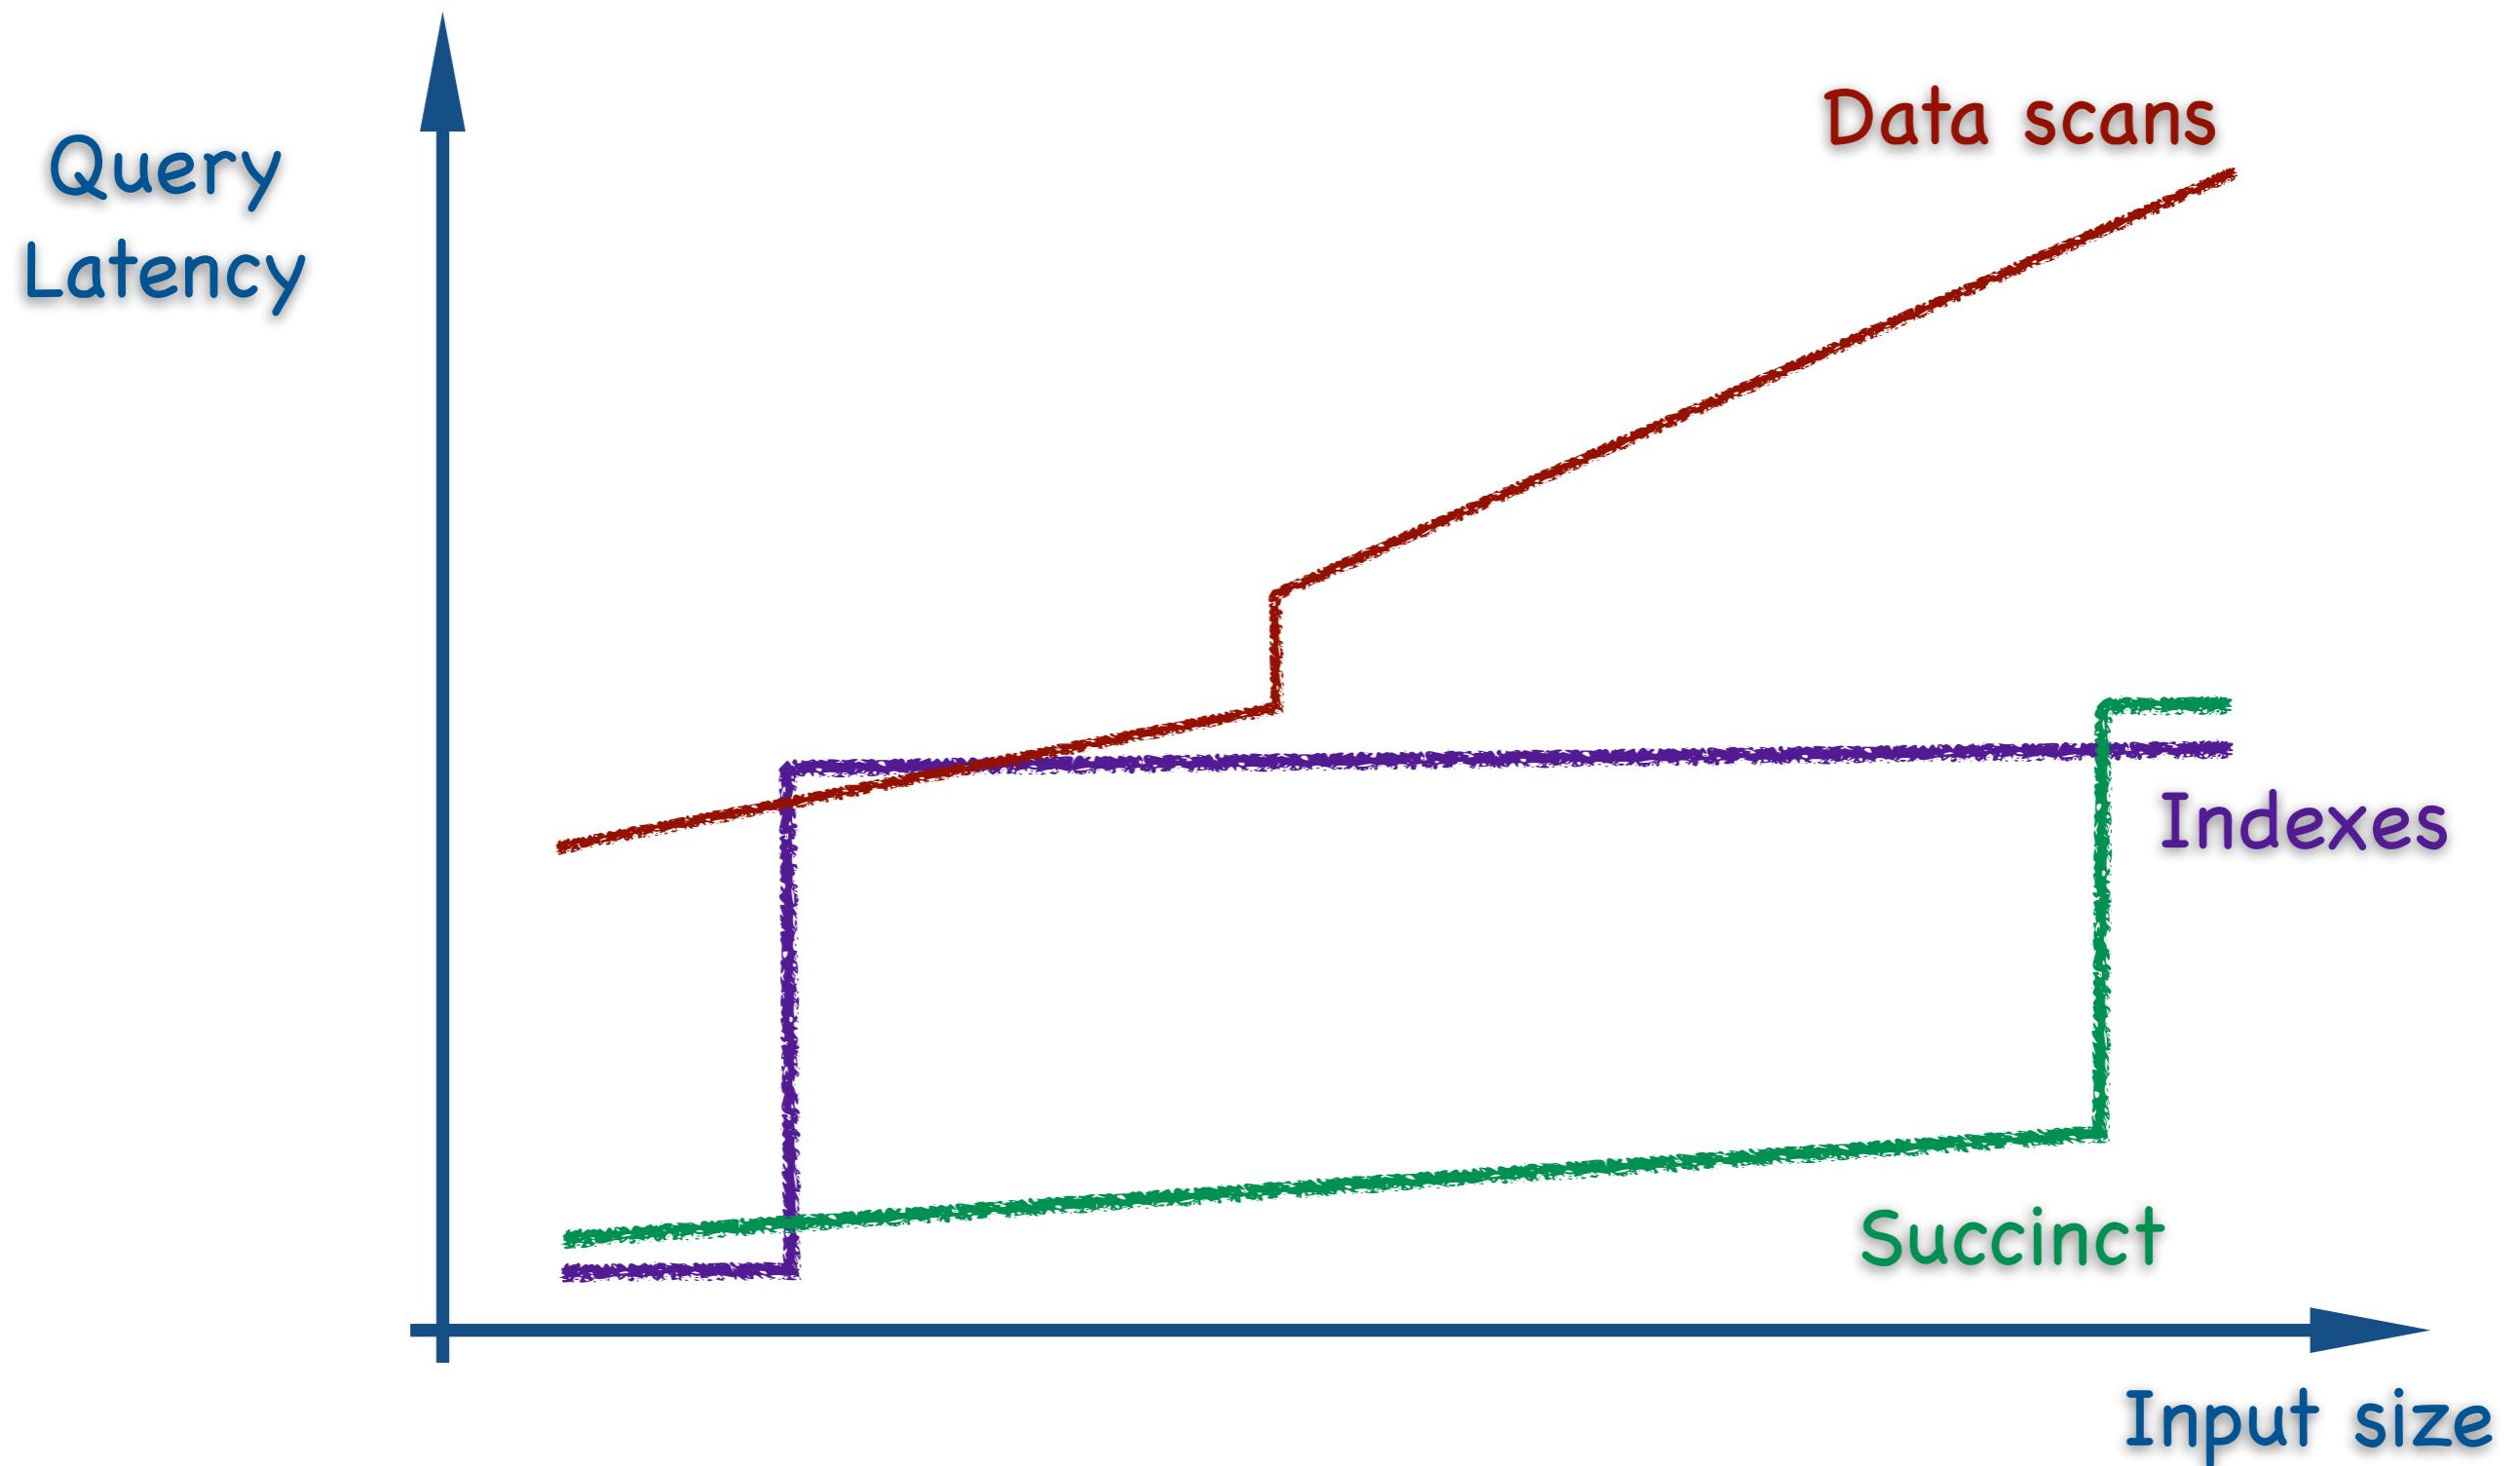
No data scans

Functionality of indexes

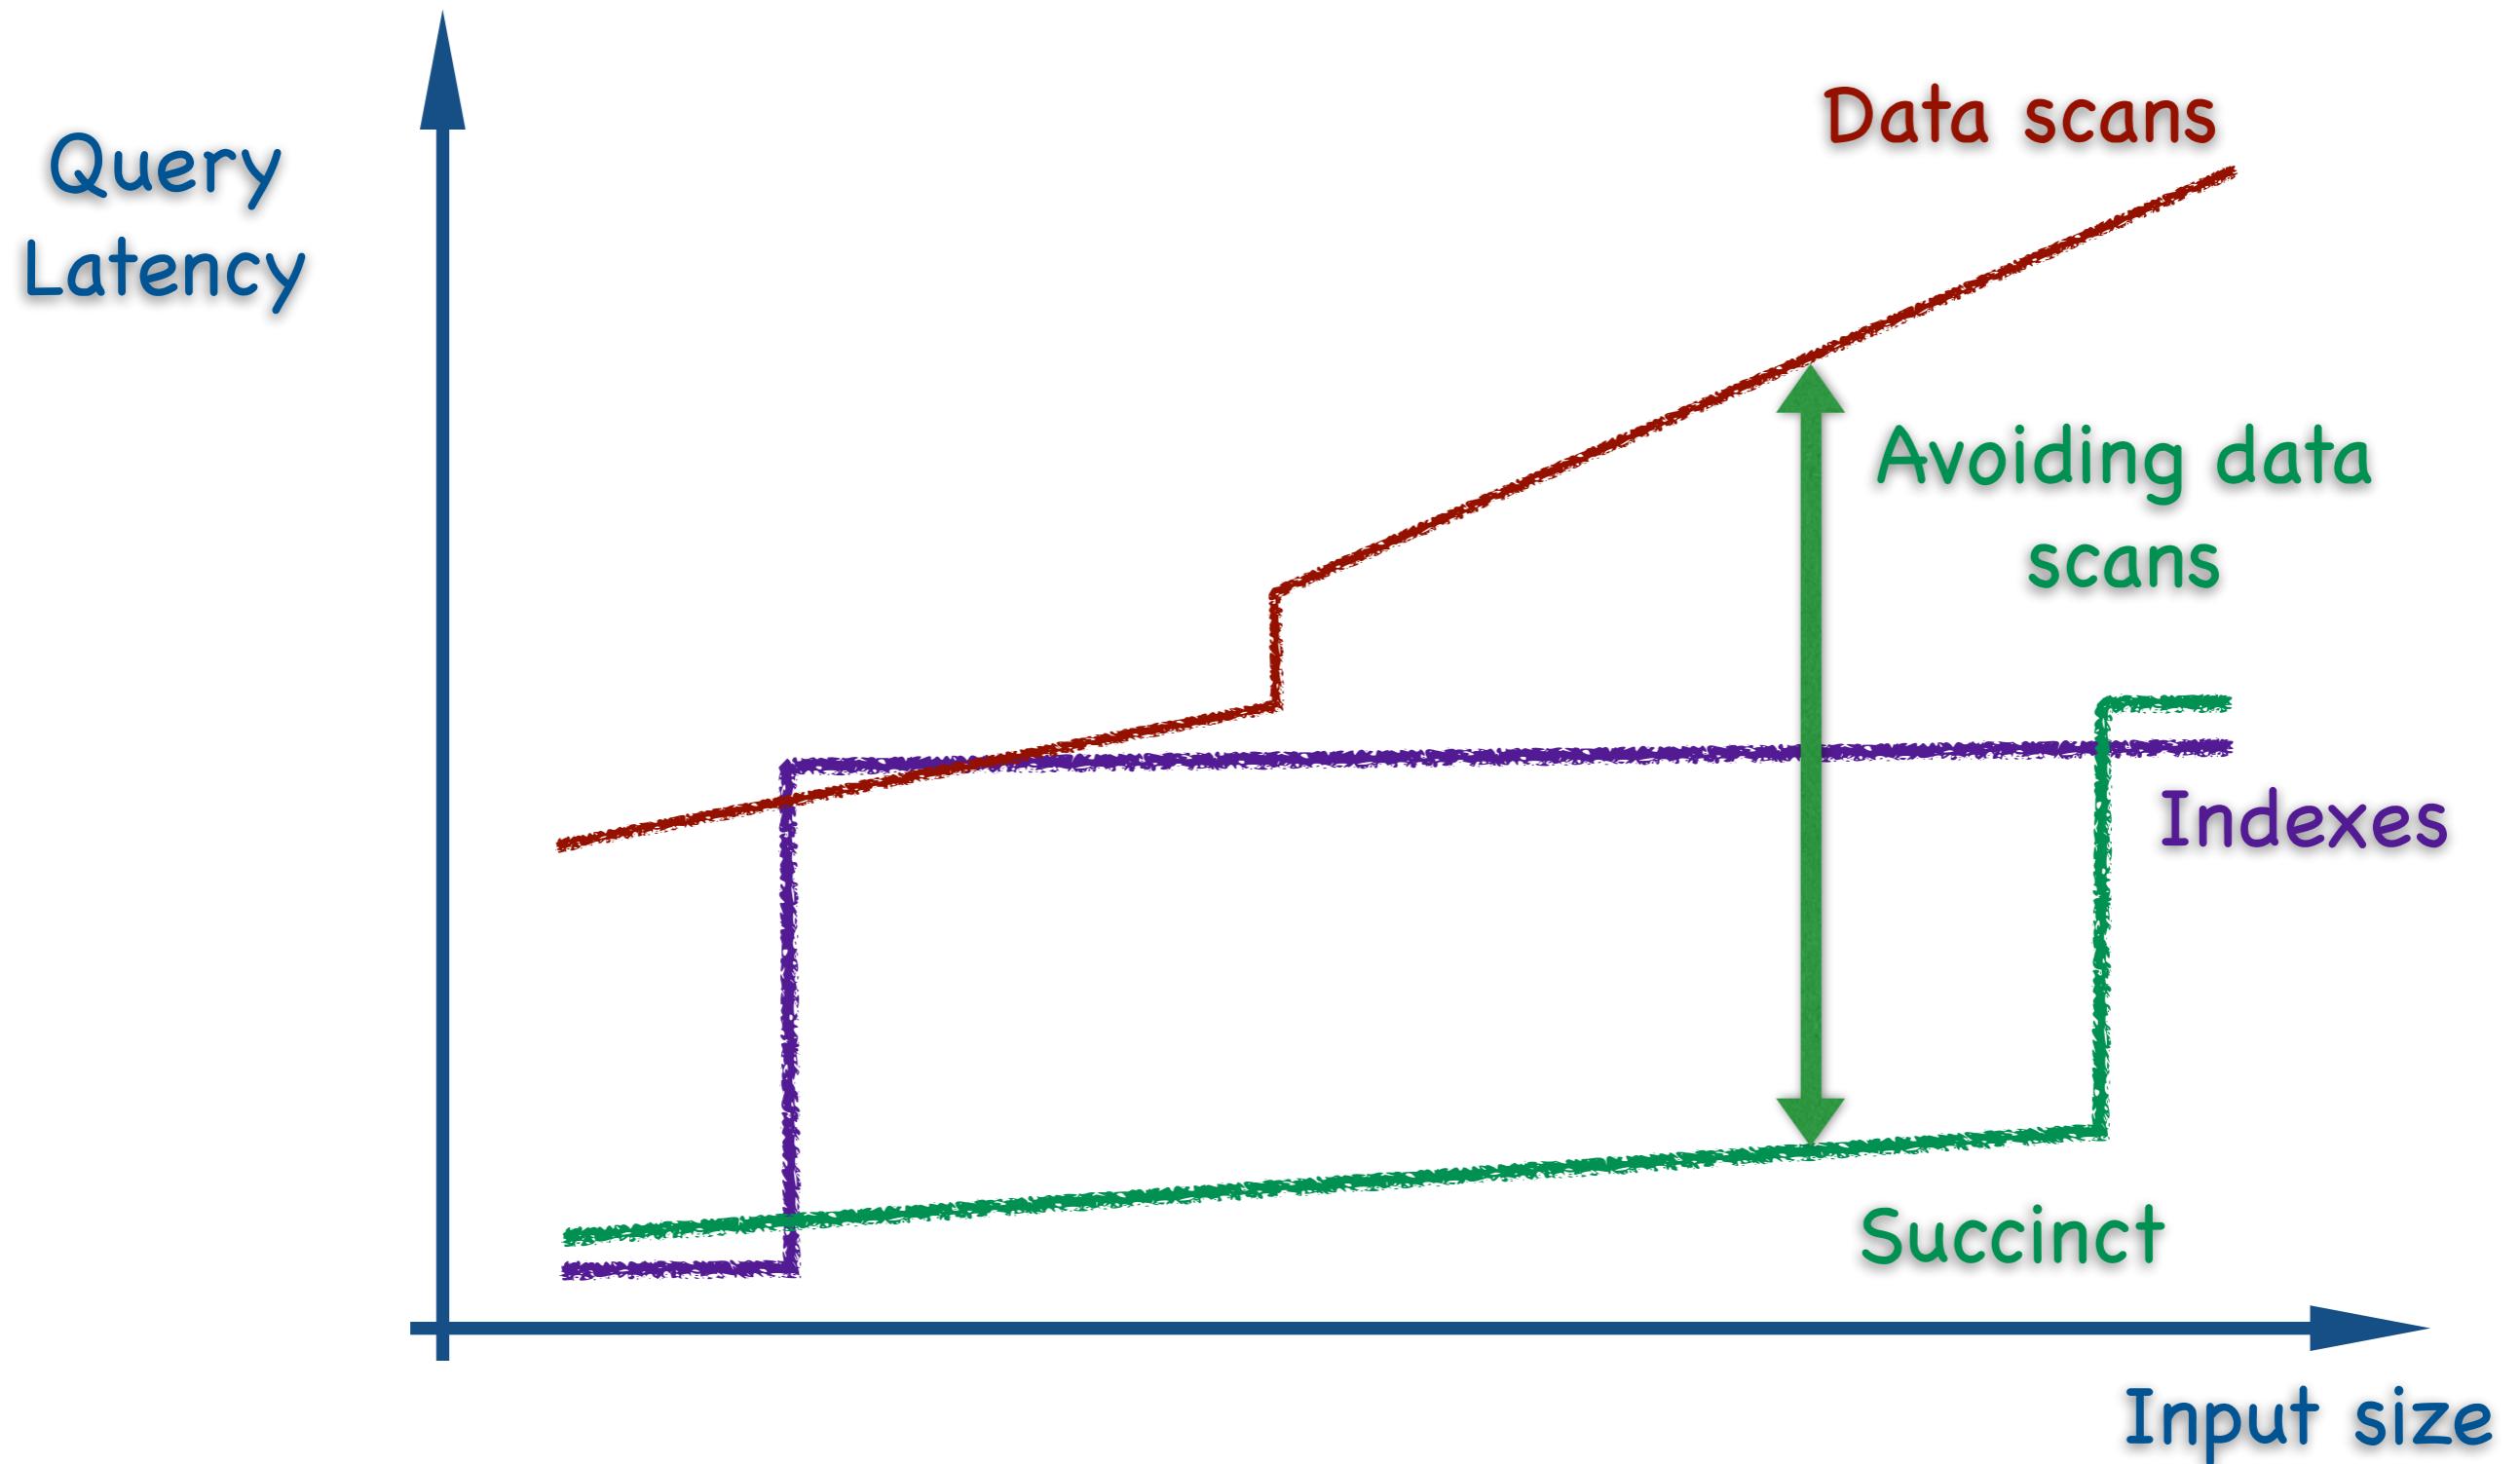
No  
decompression

Queries directly on  
the compressed representation  
(except for data access queries)

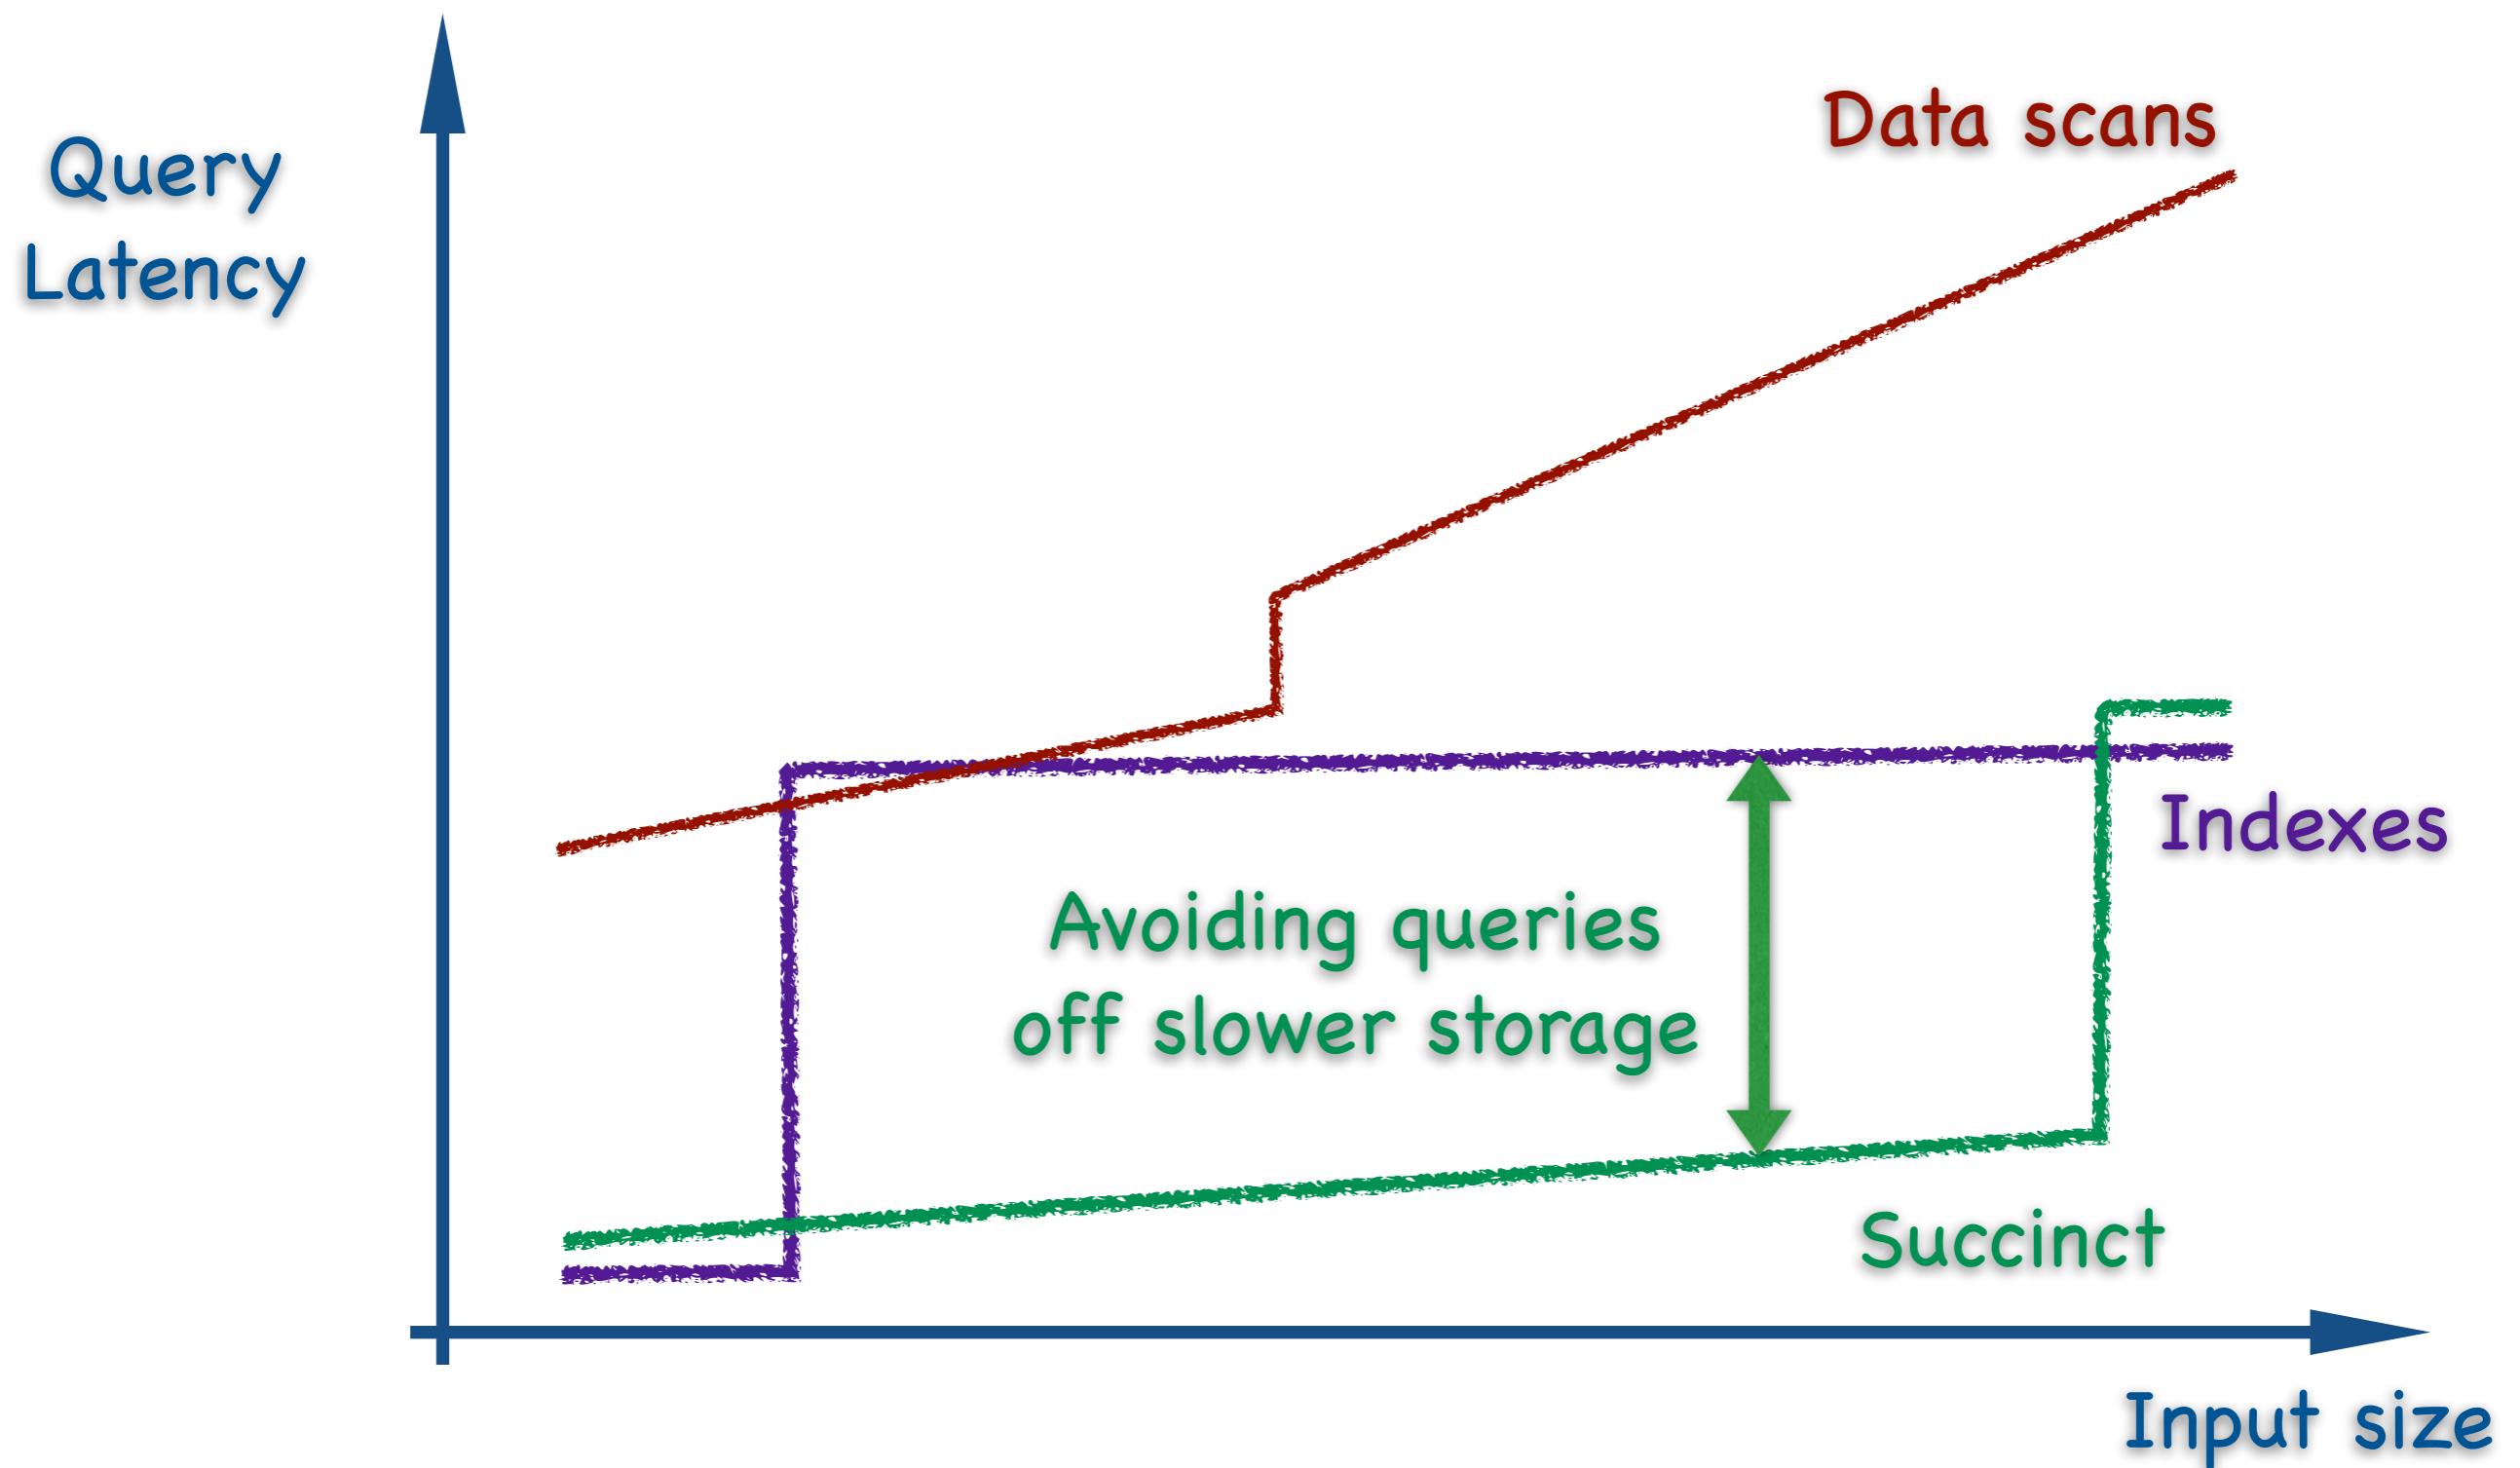
# Search latency, qualitatively



# Search latency, qualitatively



# Search latency, qualitatively



# What do we lose?

---

## Succinct

- No secondary indexes
- No data scans, no data decompression

# What do we lose?

---

## Succinct

- No secondary indexes
- No data scans, no data decompression

## What do we lose?

# What do we lose?

---

## Succinct

- No secondary indexes
- No data scans, no data decompression

## What do we lose?

- Preprocessing expensive (4GB/hour/core)
- CPU (for data access)
- Sequential scan throughput (13Mbps/thread)
- In-place updates

# Succinct Data Representation

---

Builds upon a large body of theory results

# Succinct Data Representation

---

Builds upon a large body of theory results

- FM-index
  - BWT-based technique
- Compressed Suffix arrays
  - Suffix array based technique

# Succinct Data Representation

---

Builds upon a large body of theory results

- FM-index
  - BWT-based technique
- Compressed Suffix arrays
  - Suffix array based technique

Succinct

- Builds upon latter due to simplicity
- Improved data structures (1.25–3x more space efficient)
- New query algorithm (2.3x faster on average)

# Array of suffixes (AoS)

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

# Array of suffixes (AoS)

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

**Step 1:**  
**Construct set of suffixes**

# Array of suffixes (AoS)

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

Step 1:  
Construct set of suffixes

# Array of suffixes (AoS)

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---	---

Step 1:  
Construct set of suffixes

# Array of suffixes (AoS)

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---	---

Step 1:  
Construct set of suffixes

# Array of suffixes (AoS)

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---	---

a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---

Step 1:  
Construct set of suffixes

# Array of suffixes (AoS)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---	---

a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---

p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---

p	y	p	u	p	p	y
---	---	---	---	---	---	---

y	p	u	p	p	y
---	---	---	---	---	---

p	u	p	p	y
---	---	---	---	---

u	p	p	y
---	---	---	---

p	p	y
---	---	---

p	y
---	---

y
---

Step 1:  
Construct set of suffixes

# Array of suffixes (AoS)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---	---

a	p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---	---

p	p	y	p	u	p	p	y
---	---	---	---	---	---	---	---

p	y	p	u	p	p	y
---	---	---	---	---	---	---

y	p	u	p	p	y
---	---	---	---	---	---

p	u	p	p	y
---	---	---	---	---

u	p	p	y
---	---	---	---

p	p	y
---	---	---

p	y
---	---

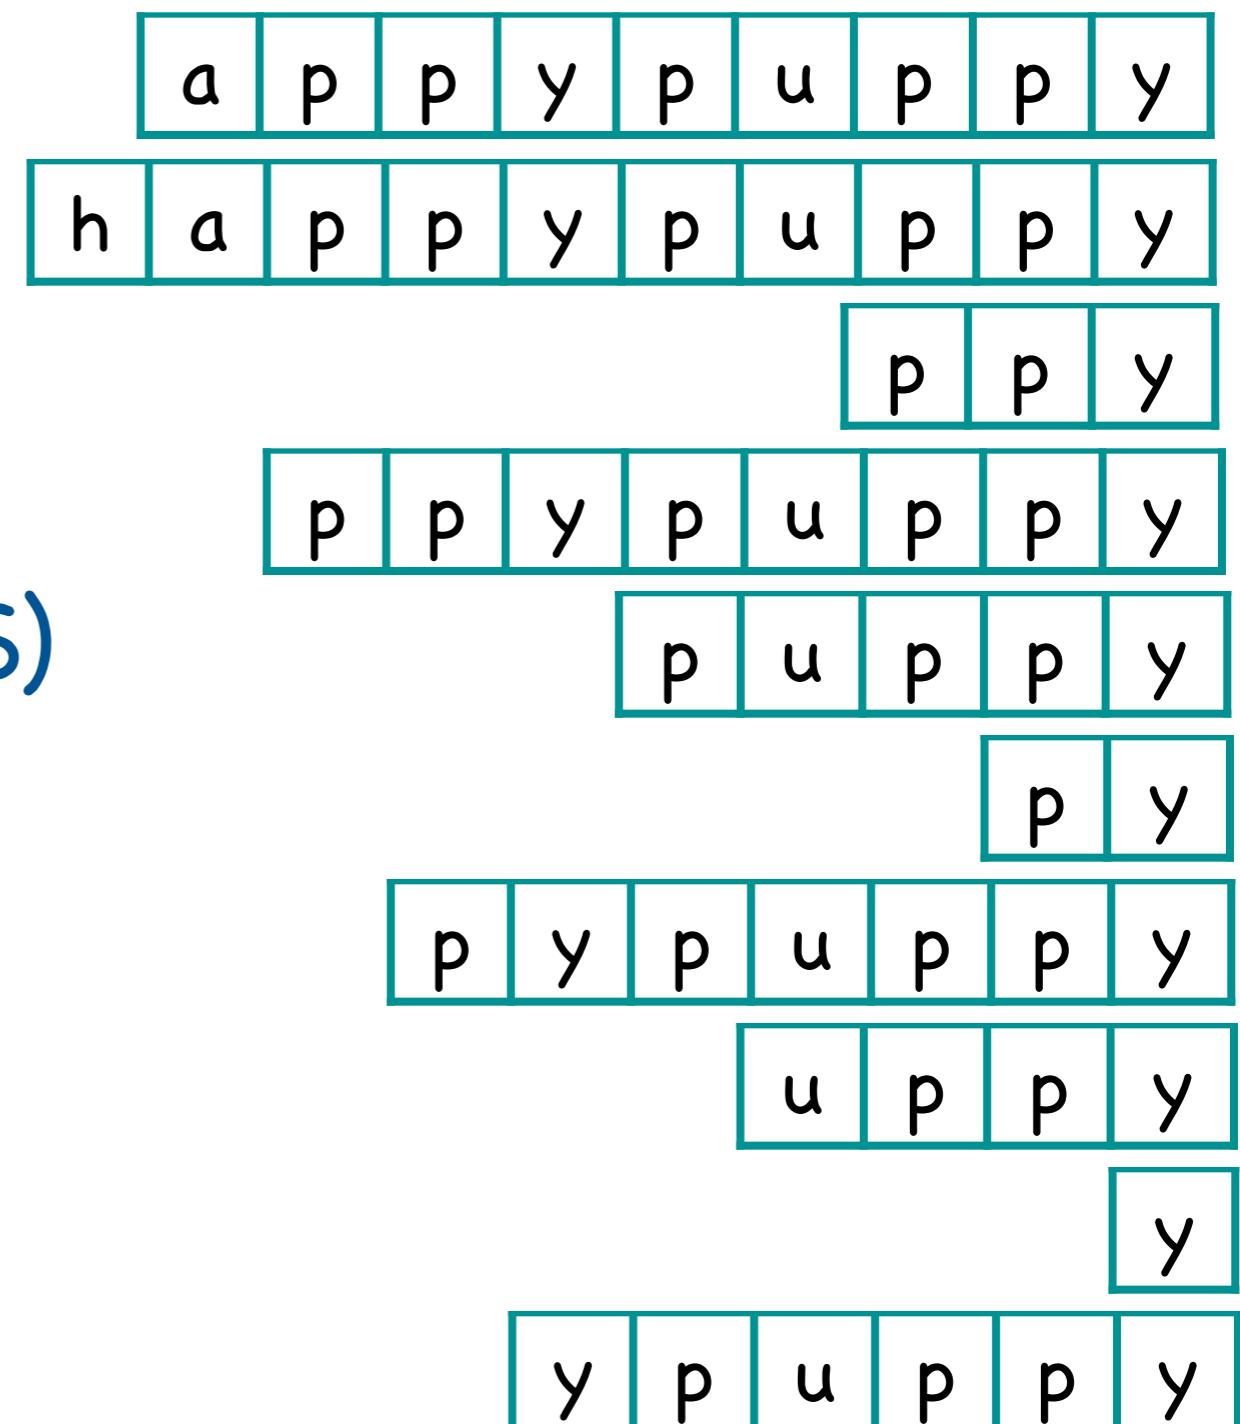
y
---

**Step 2:**  
**Sort suffixes**  
**(lexicographically)**

# Array of suffixes (AoS)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

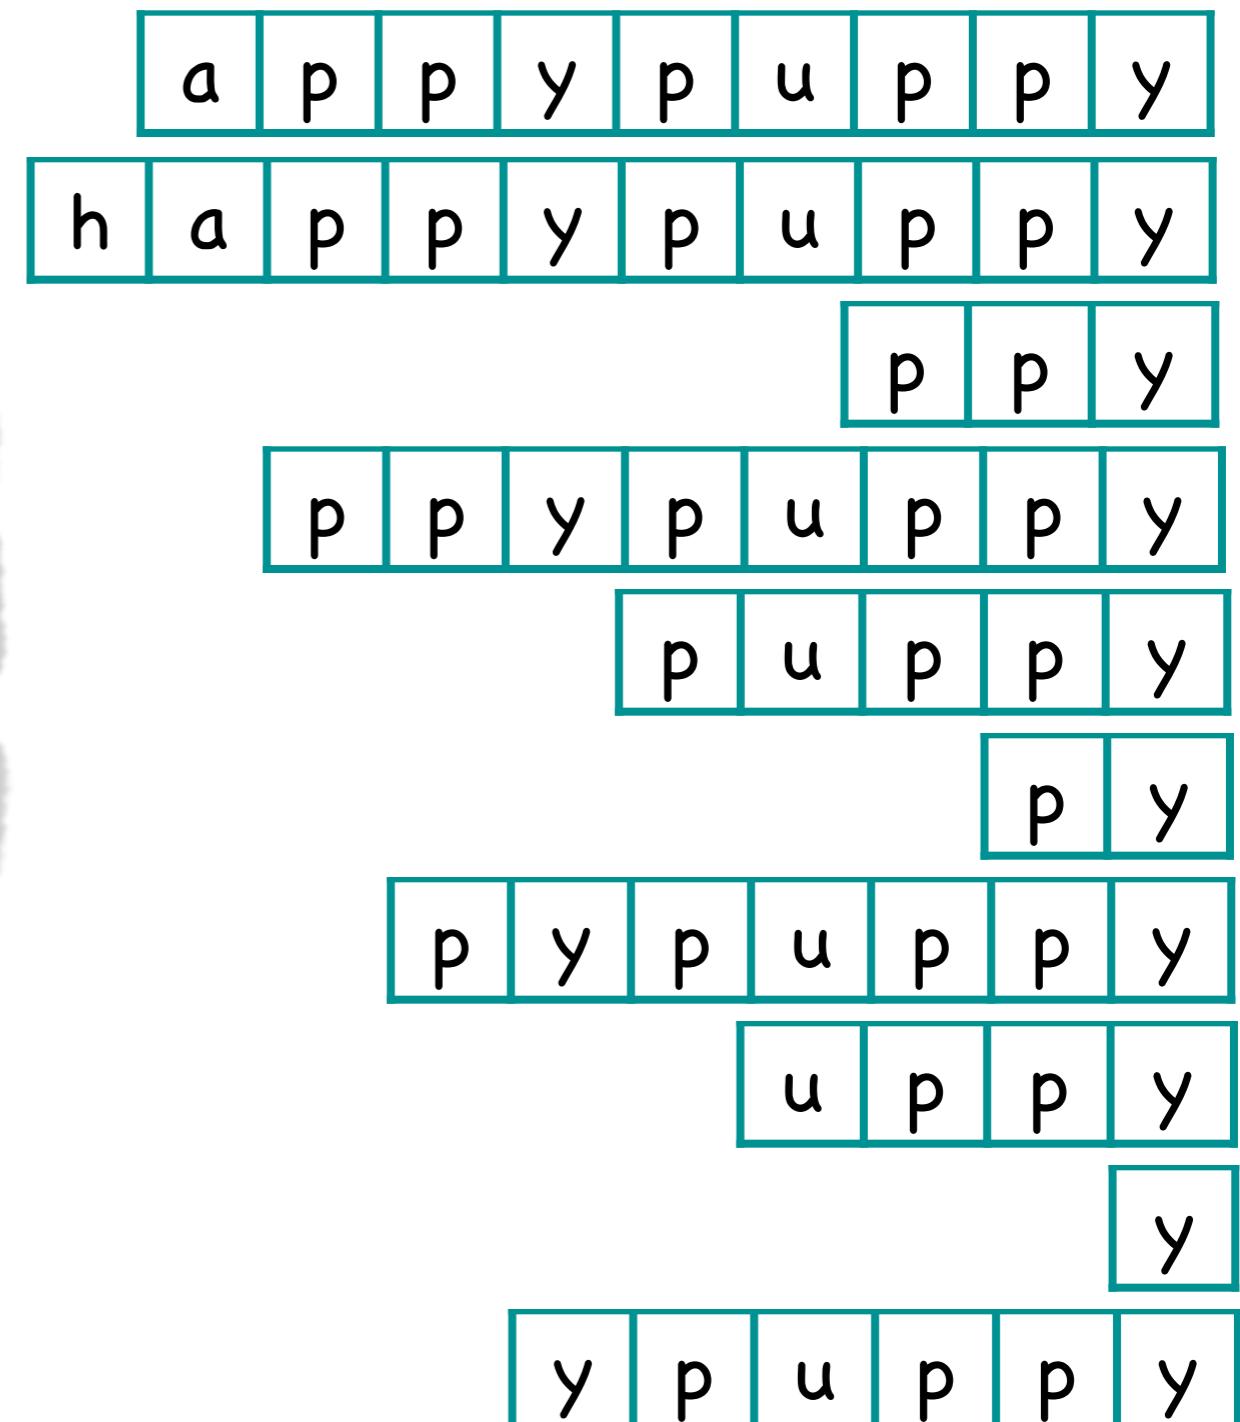
# Array of suffixes (AoS)



# Suffix array (AoS2Input)

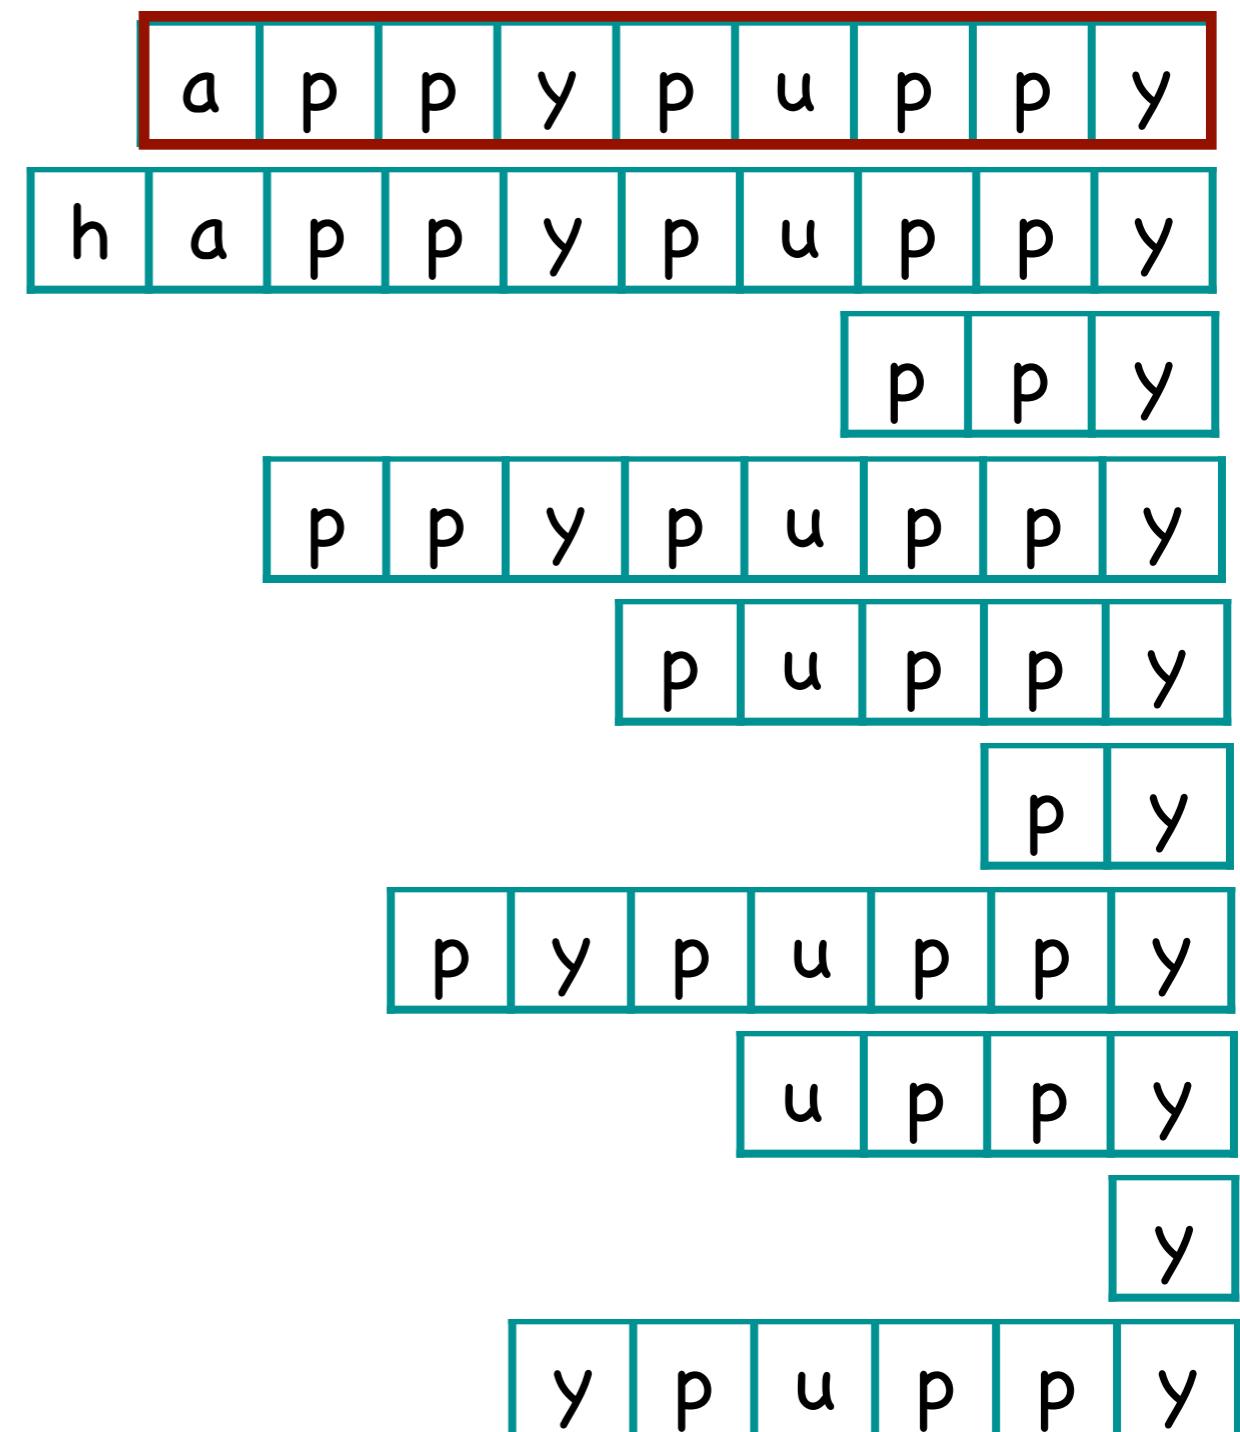
0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

**Step 3:**  
**Mark location of  
each sorted suffix  
(in input file)**



# Suffix array (AoS2Input)

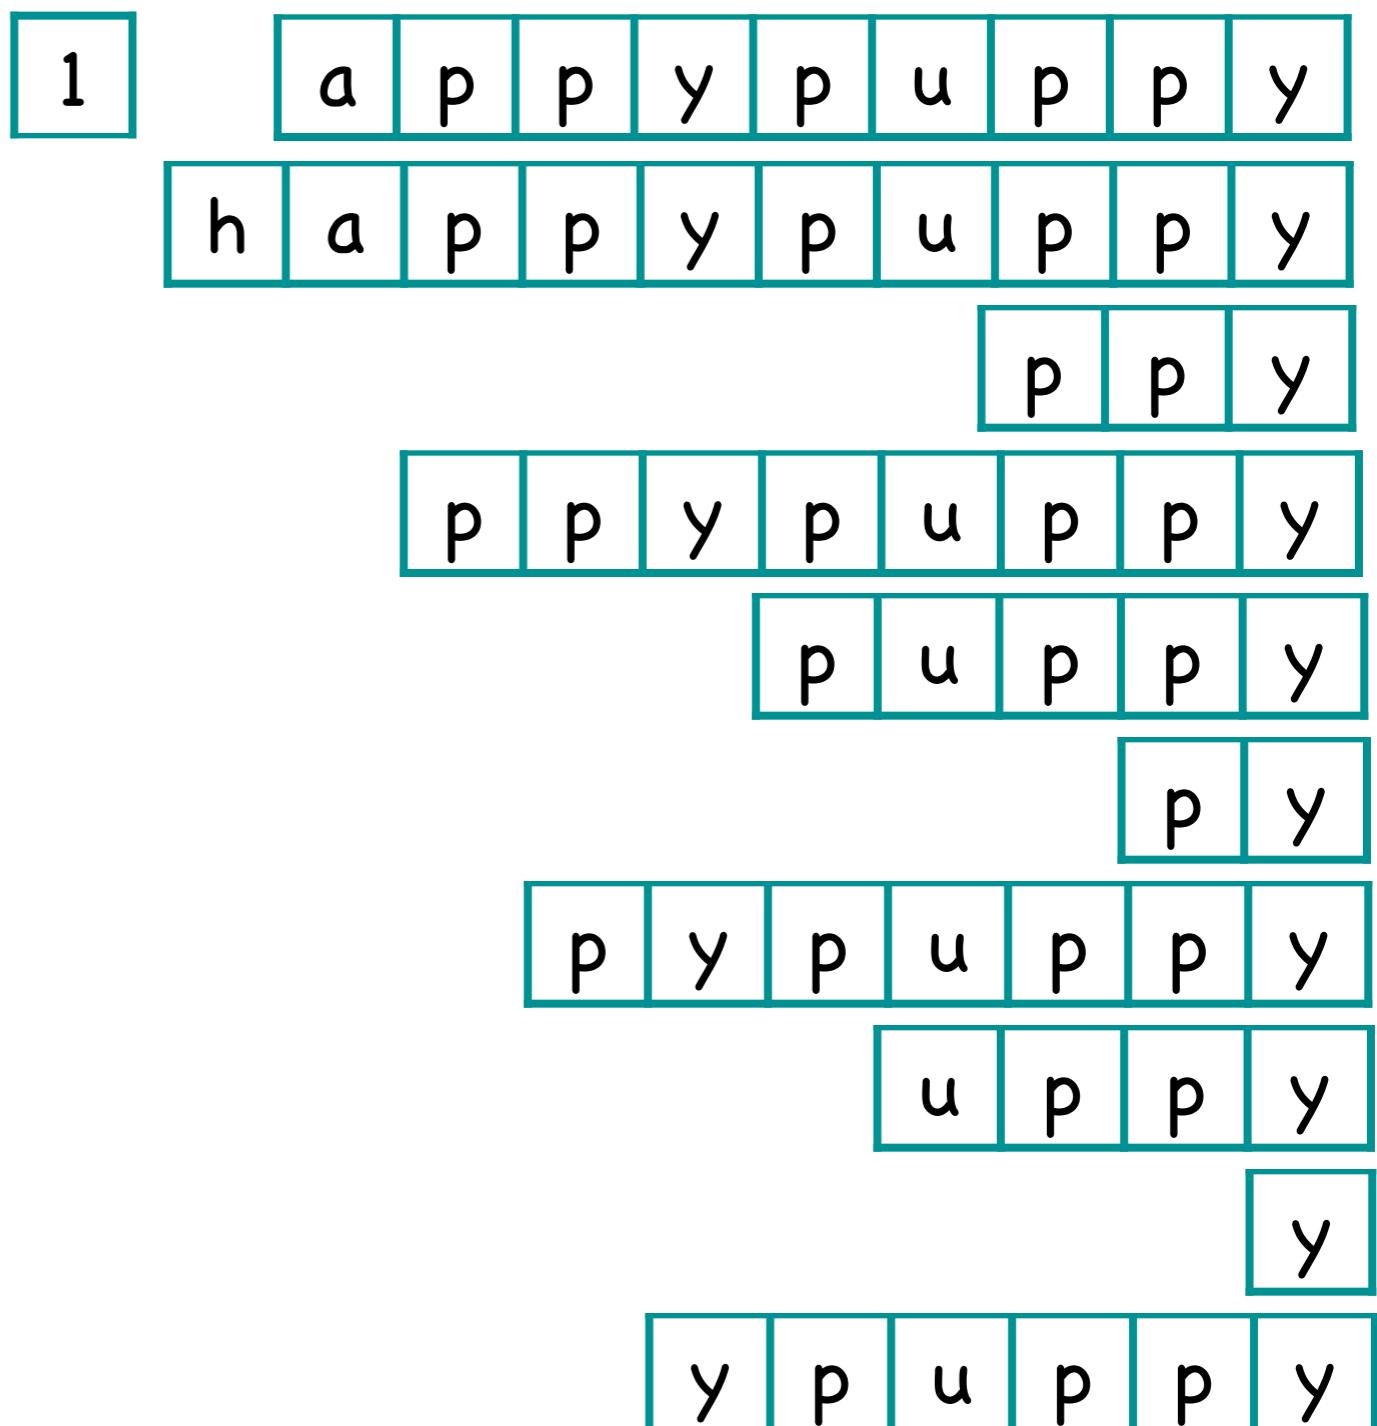
0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



Step 3:  
Mark location of  
each sorted suffix  
(in input file)

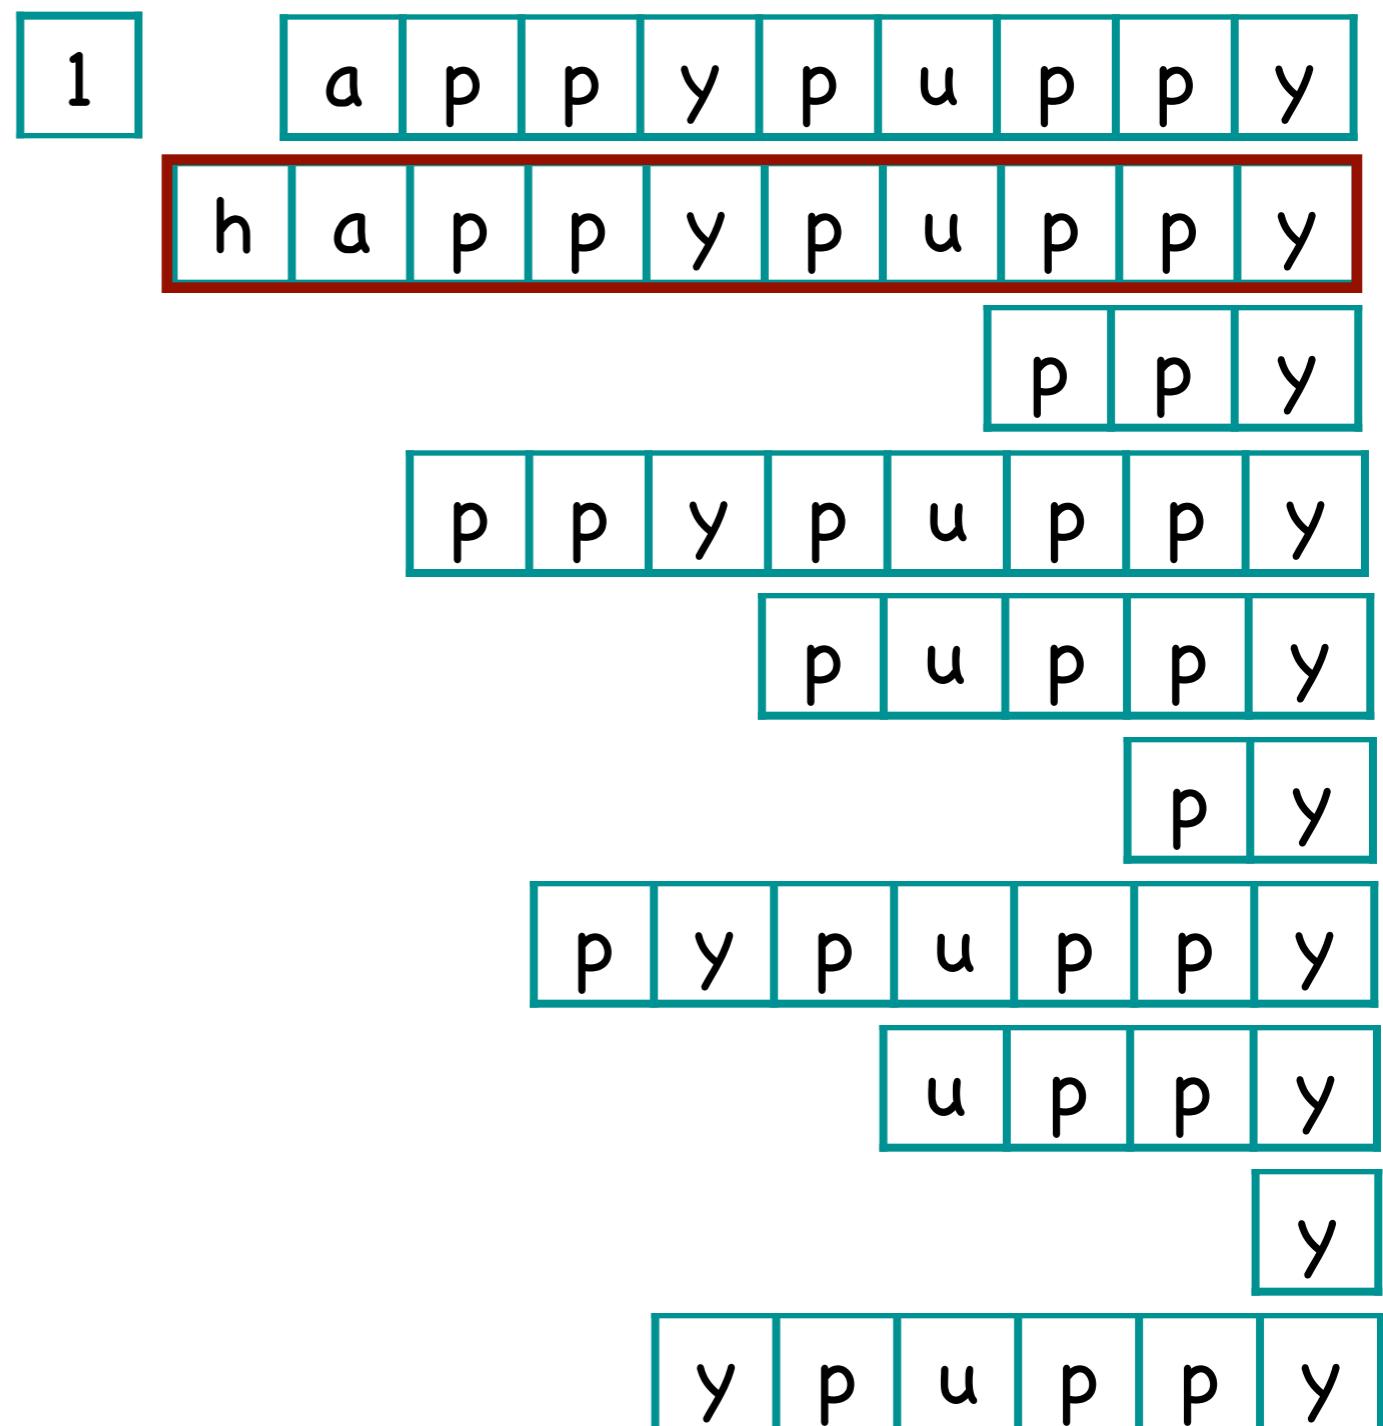
# Suffix array (AoS2Input)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



# Suffix array (AoS2Input)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

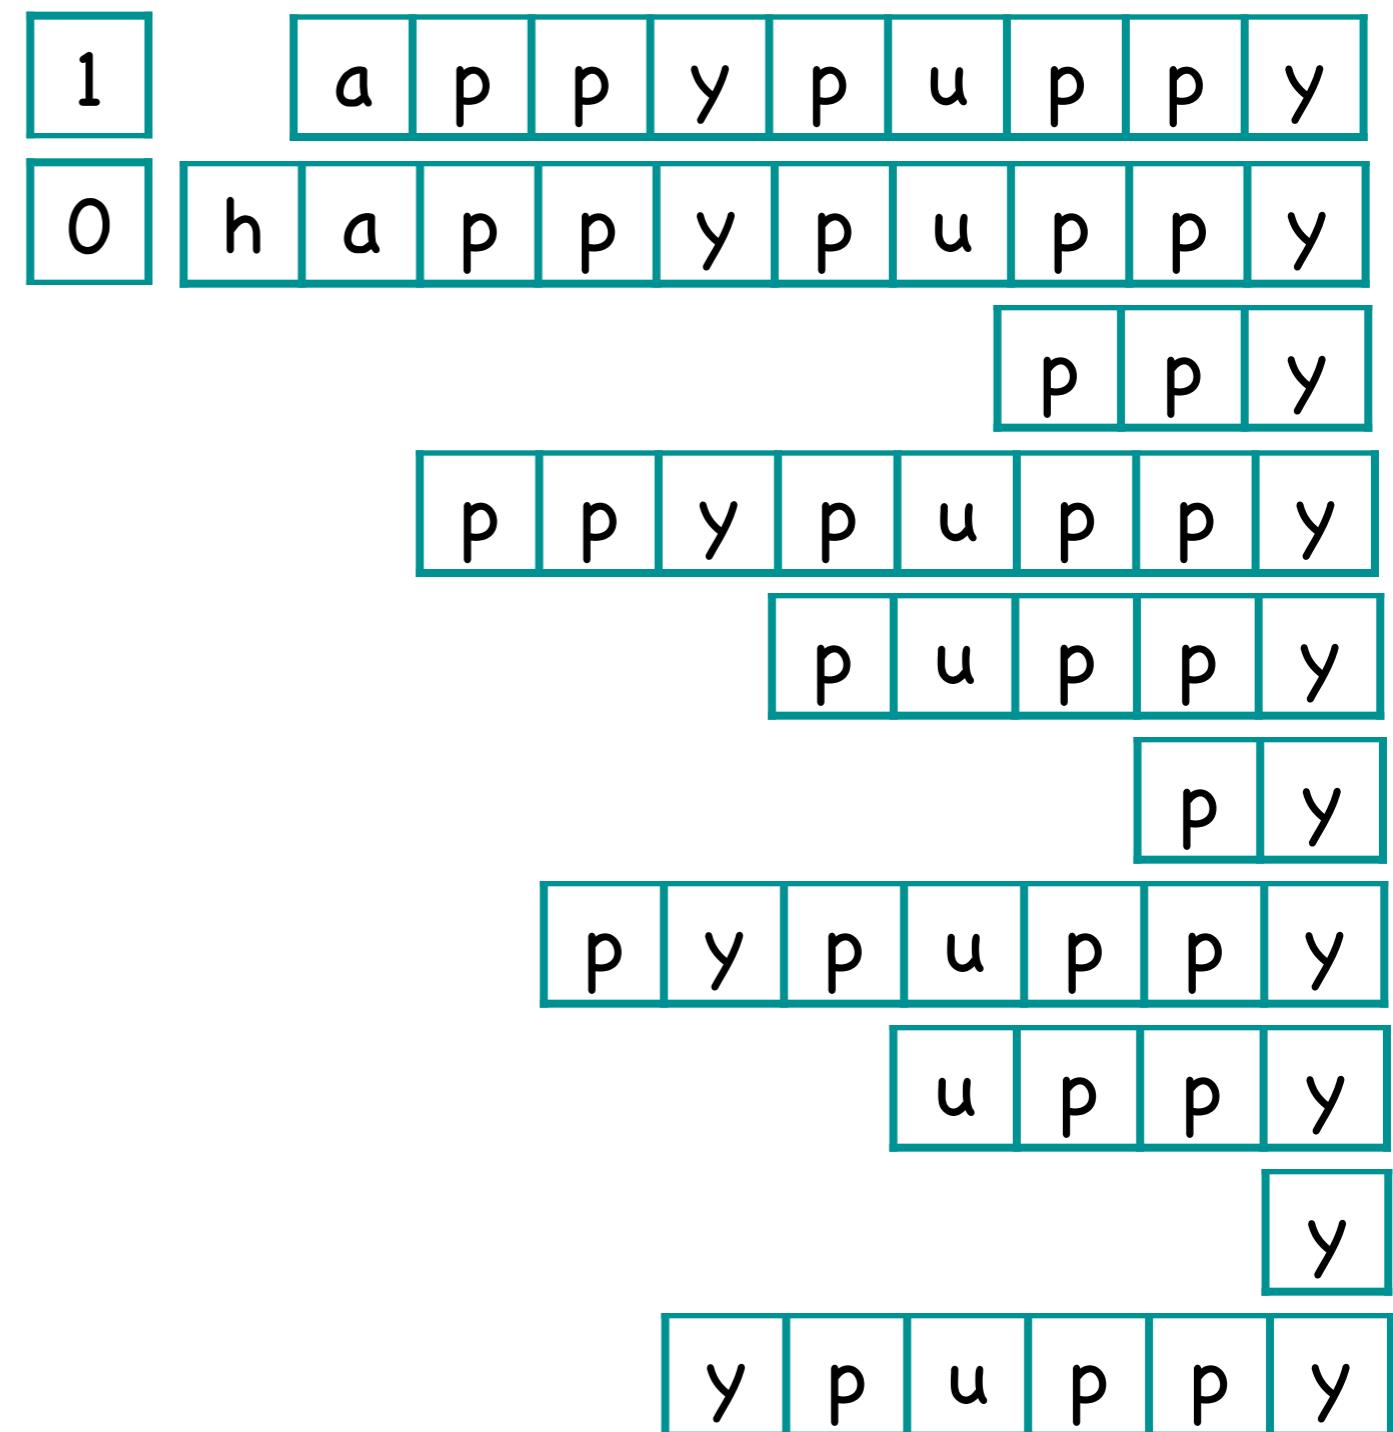


Step 3:  
Mark location of  
each sorted suffix  
(in input file)

# Suffix array (AoS2Input)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

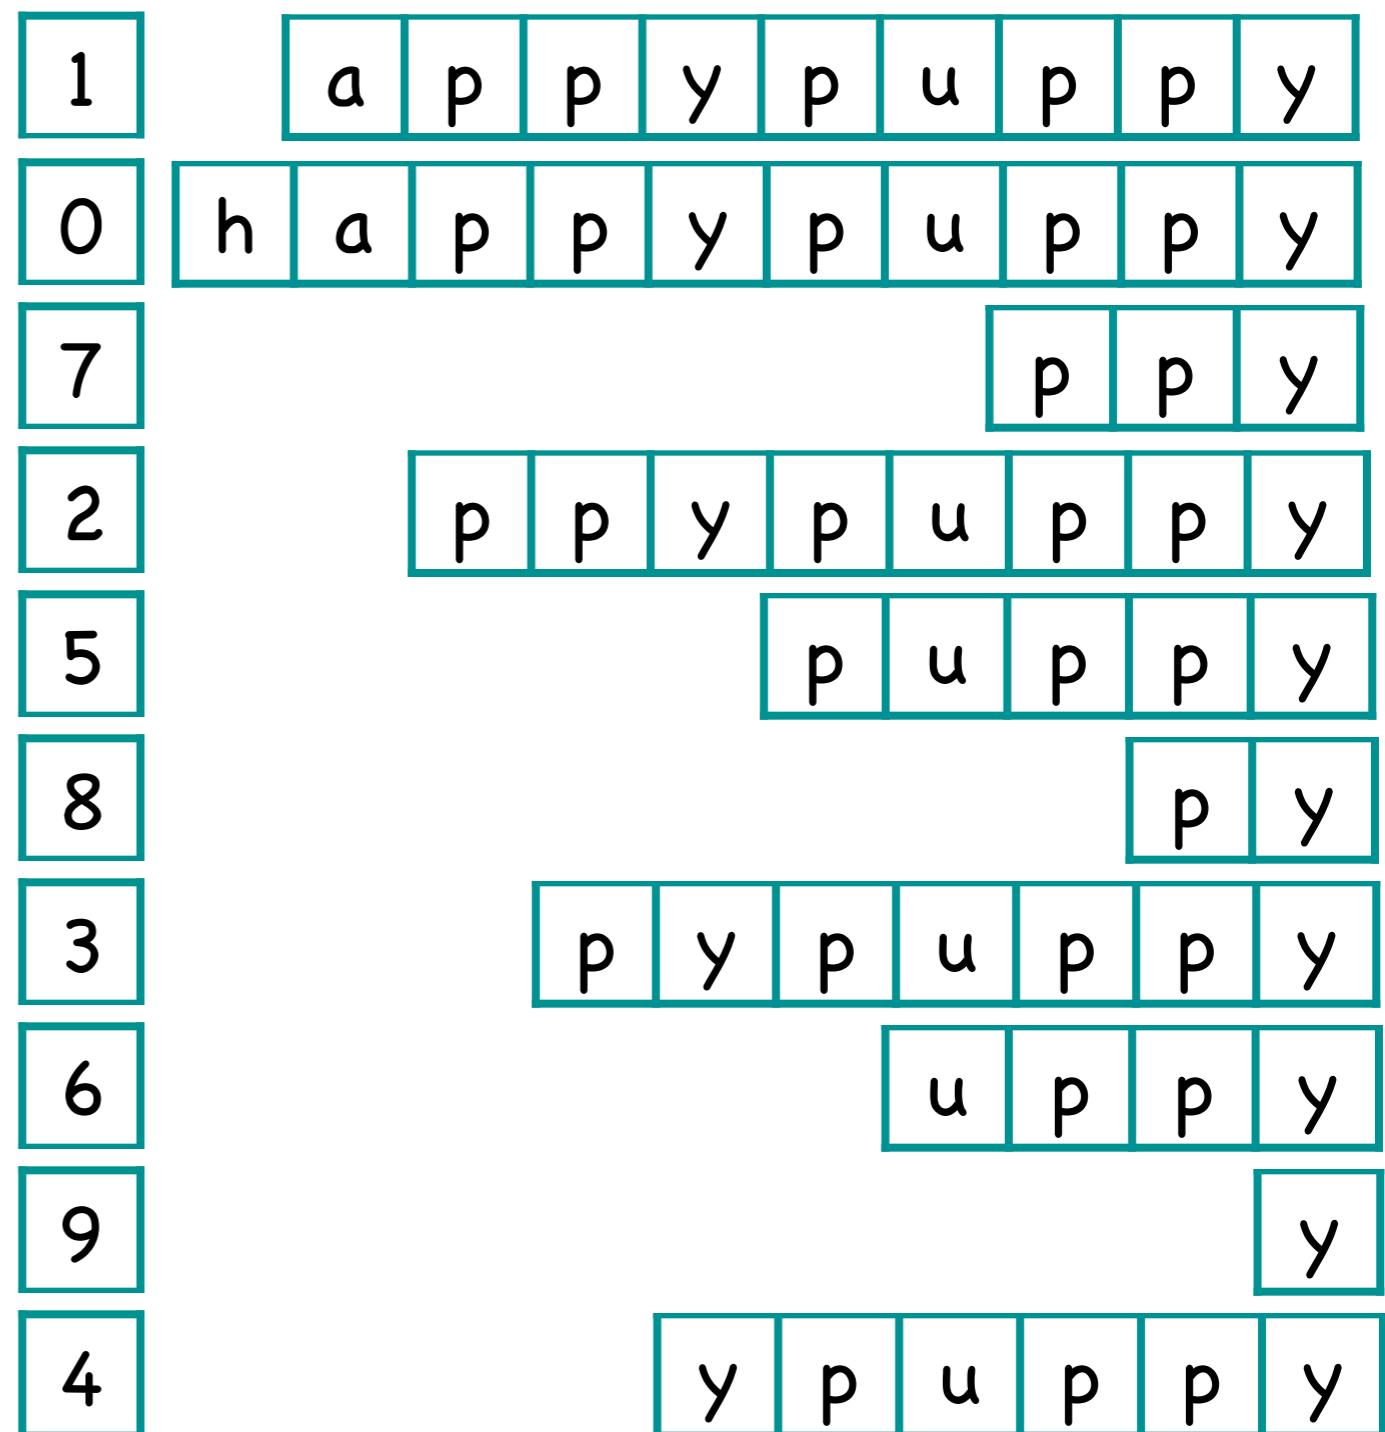
## Step 3: Mark location of each sorted suffix (in input file)



# Suffix array (AoS2Input)

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

AoS2Input



# Succinct Data Representation

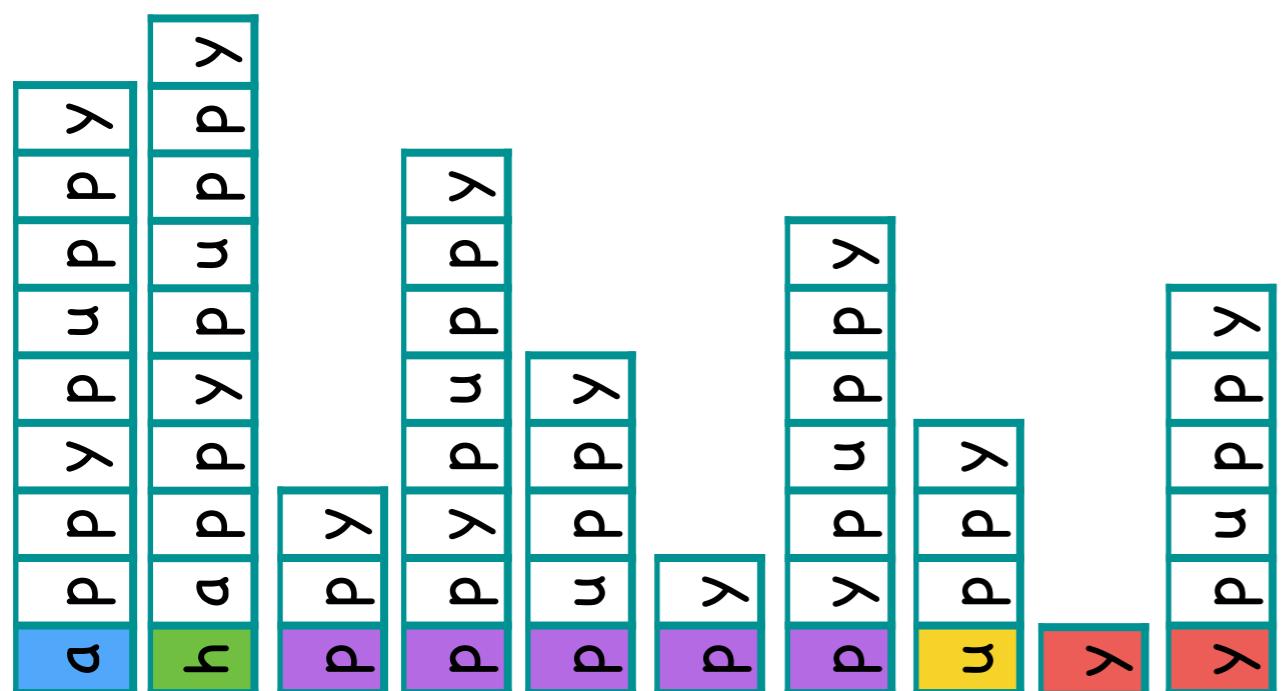
---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



Array of Suffixes (AoS)  
(suffixes in sorted order)



# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

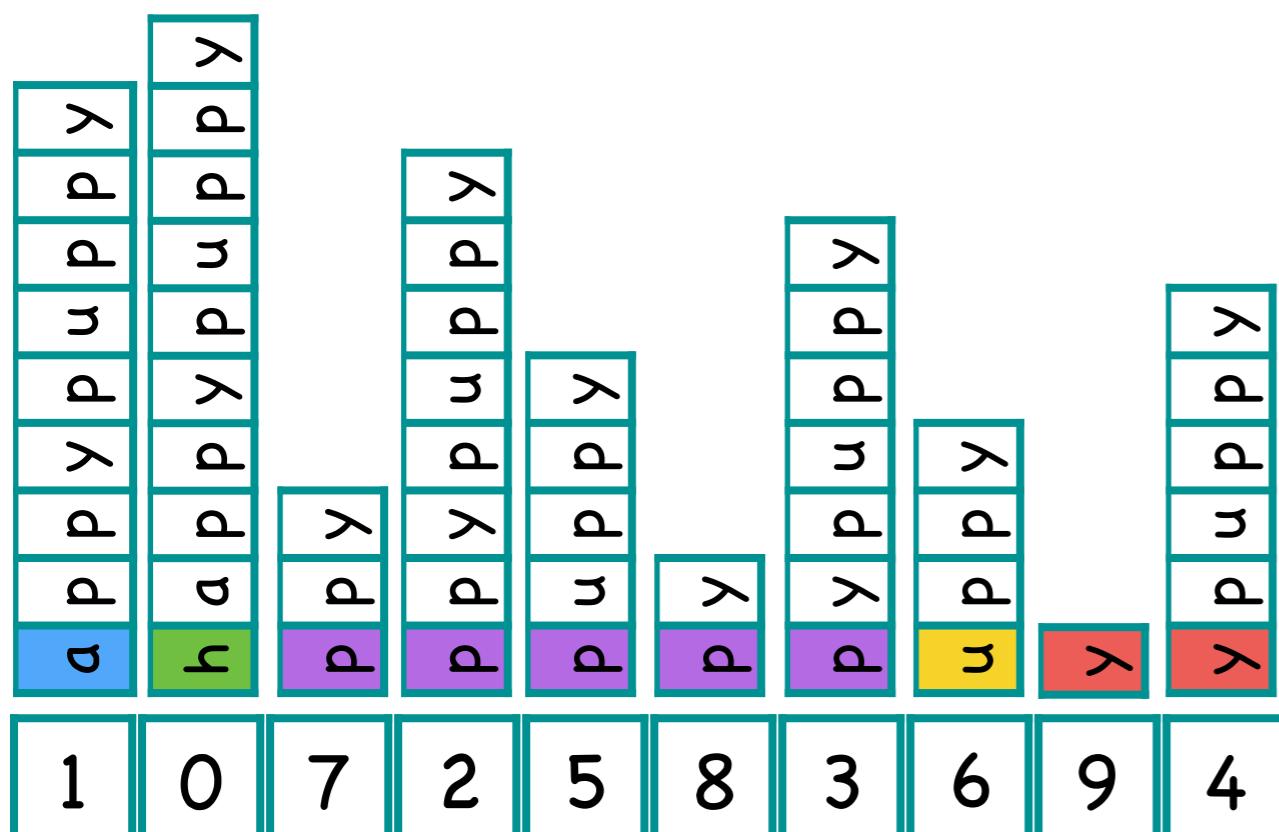
h	a	p	p	y	p	u	p	p	y
a	p	p	y	p	u	p	p	y	
p	p	y	p	u	p	p	y		
y	p	u	p	p	y				
1	0	7	2	5	8	3	6	9	4

Array of Suffixes (AoS)  
(suffixes in sorted order)

AoS2Input  
(location of suffix in file)

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



Arbitrary substring search  
via binary search

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
a	p	p	y	p	u	p	p	y	
p	p	y	p	u	p	p	y		
y	p	u	p	p	y				
1	0	7	2	5	8	3	6	9	4

Problem1: Size of AoS  
 $O(n^*n)$  bits

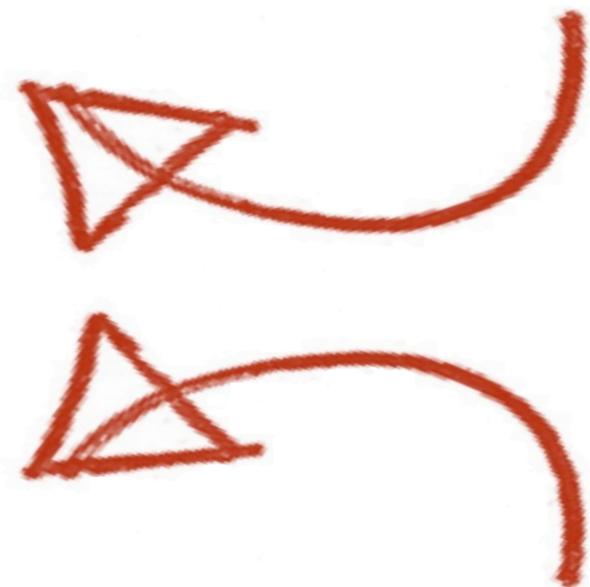


# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

h	a	p	p	y	p	u	p	p	y
a	p	p	y	p	u	p	p	y	p
p	p	y	p	u	p	p	y	p	y
y	p	y	p	u	p	p	y	p	y
p	u	p	u	p	u	p	p	y	p
u	p	u	p	u	p	p	y	p	y
p	u	p	u	p	u	p	p	y	p
u	p	u	u	p	u	p	p	y	p
p	u	p	u	u	p	p	p	y	y
h	a	p	p	y	p	u	p	p	y
1	0	7	2	5	8	3	6	9	4

Problem1: Size of AoS  
 $O(n^*n)$  bits

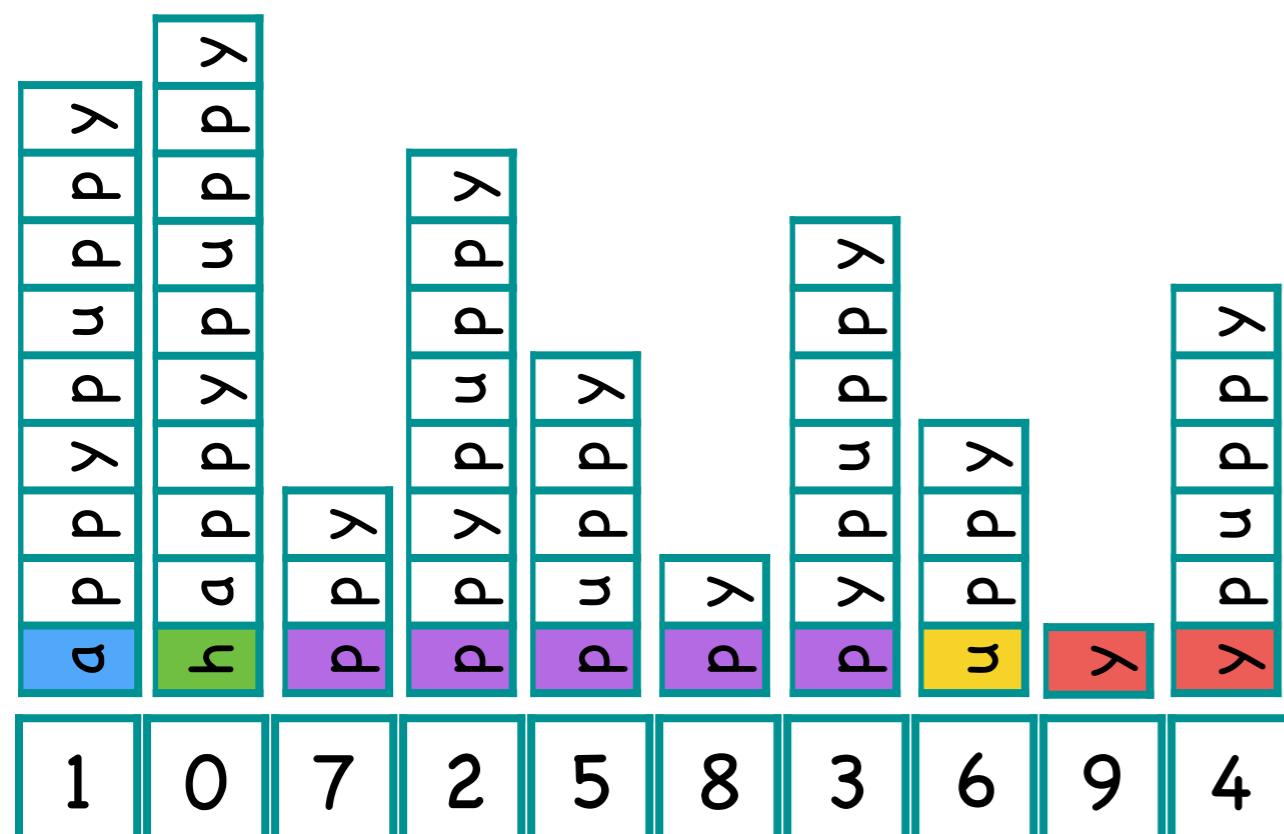


Problem2: Size of AoS2Input  
each entry takes  $\log(n)$  bits

# Succinct Data Representation

---

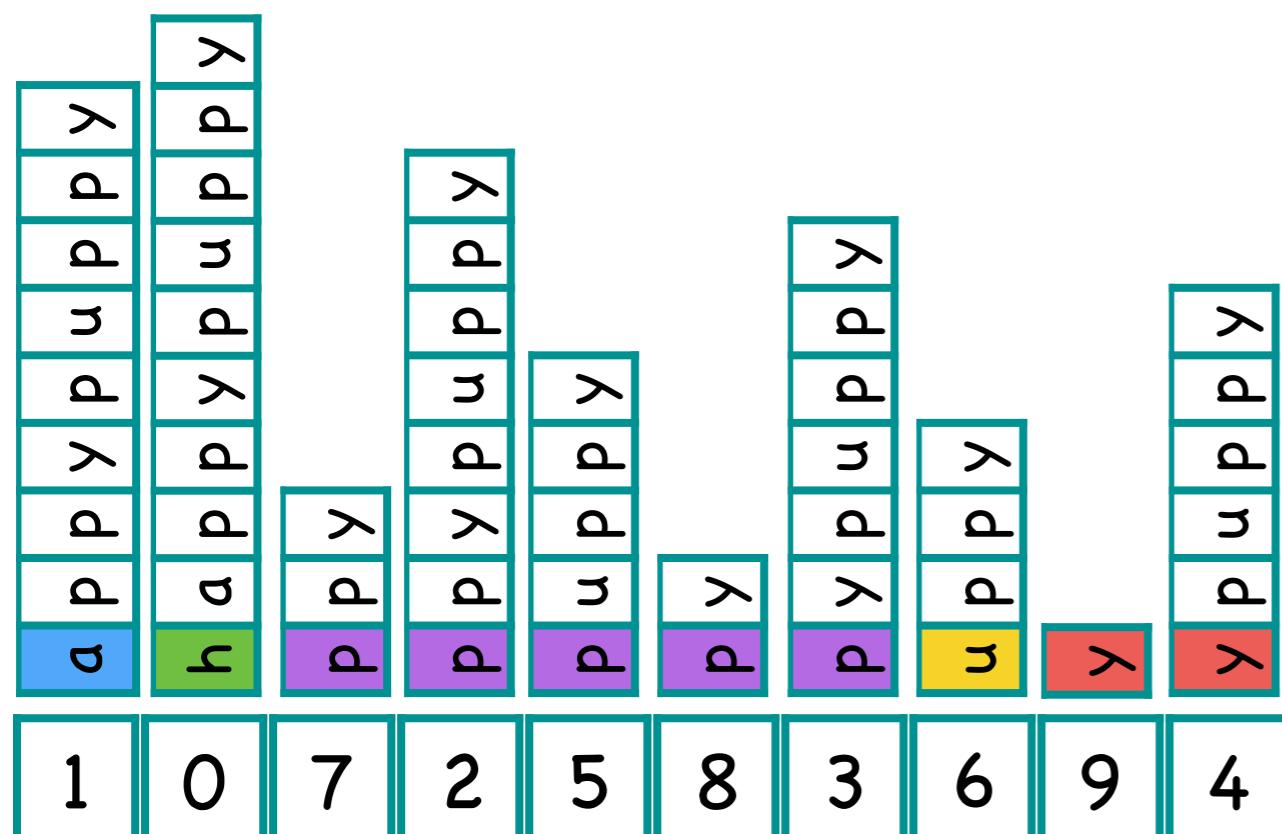
0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



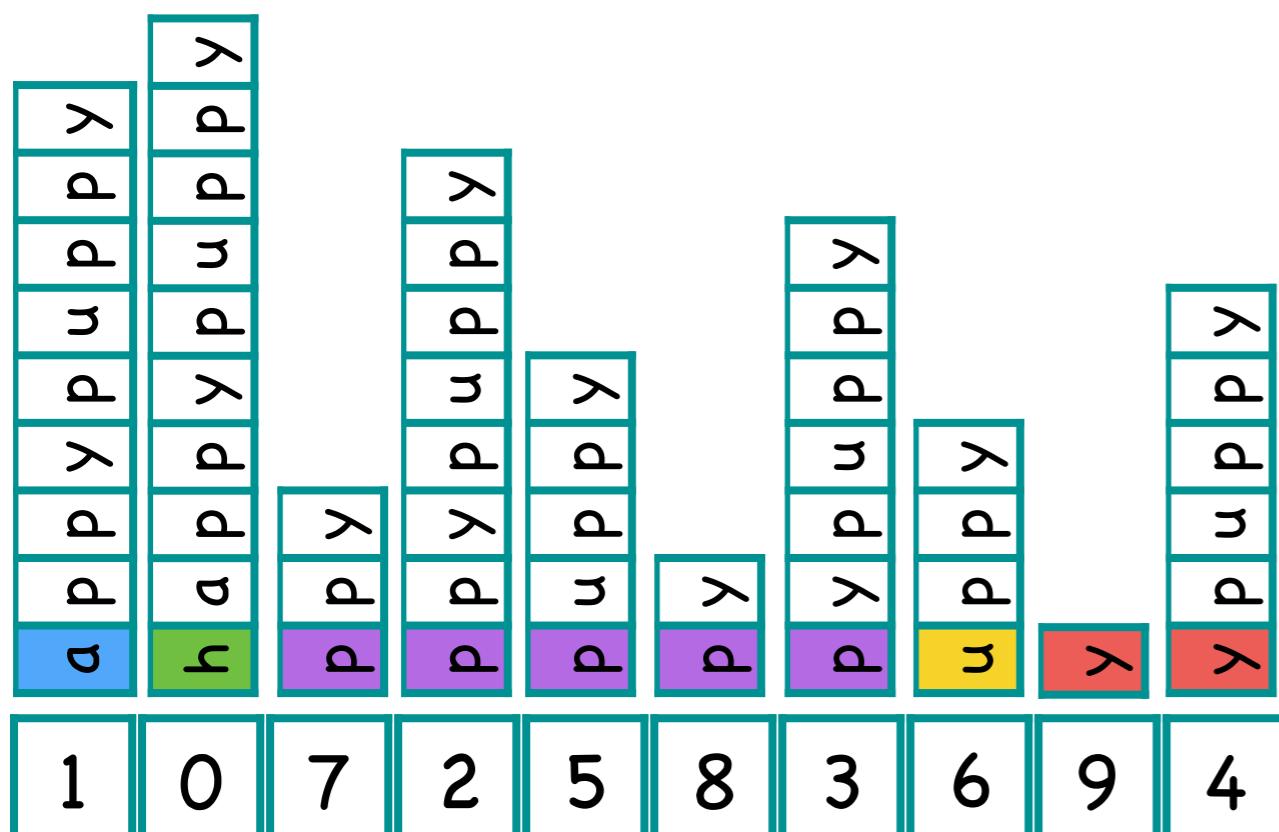
↗

0

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

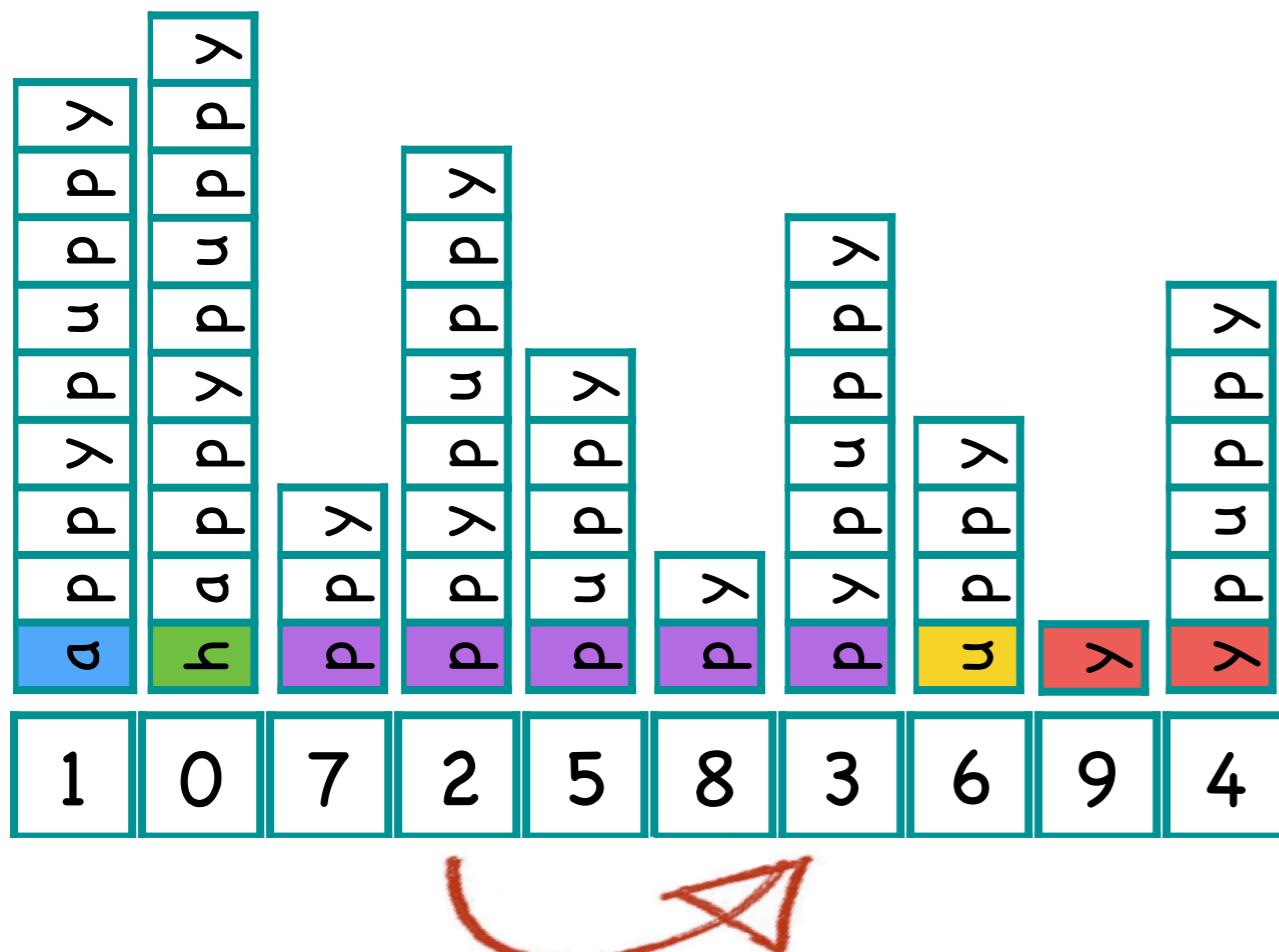


3	0
---	---

# Succinct Data Representation

---

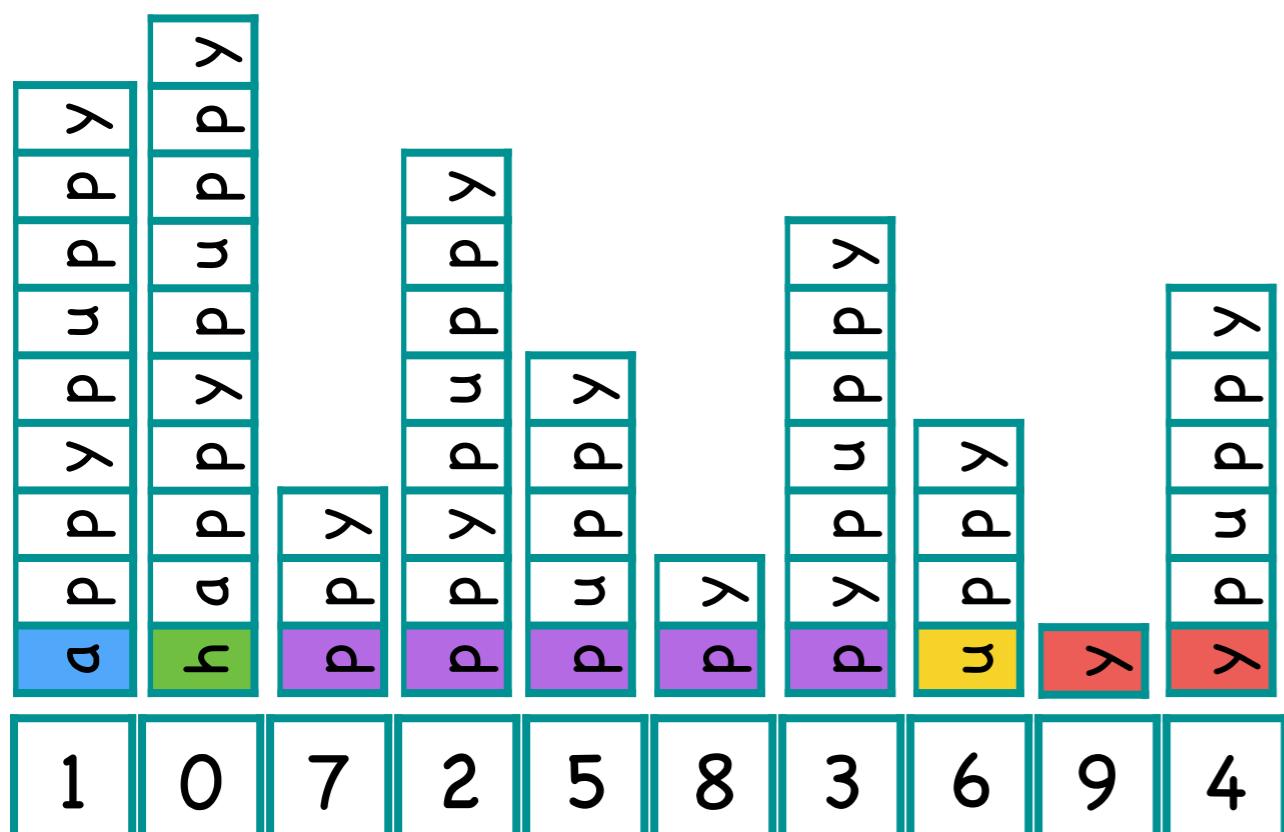
0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



3	0
6	

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

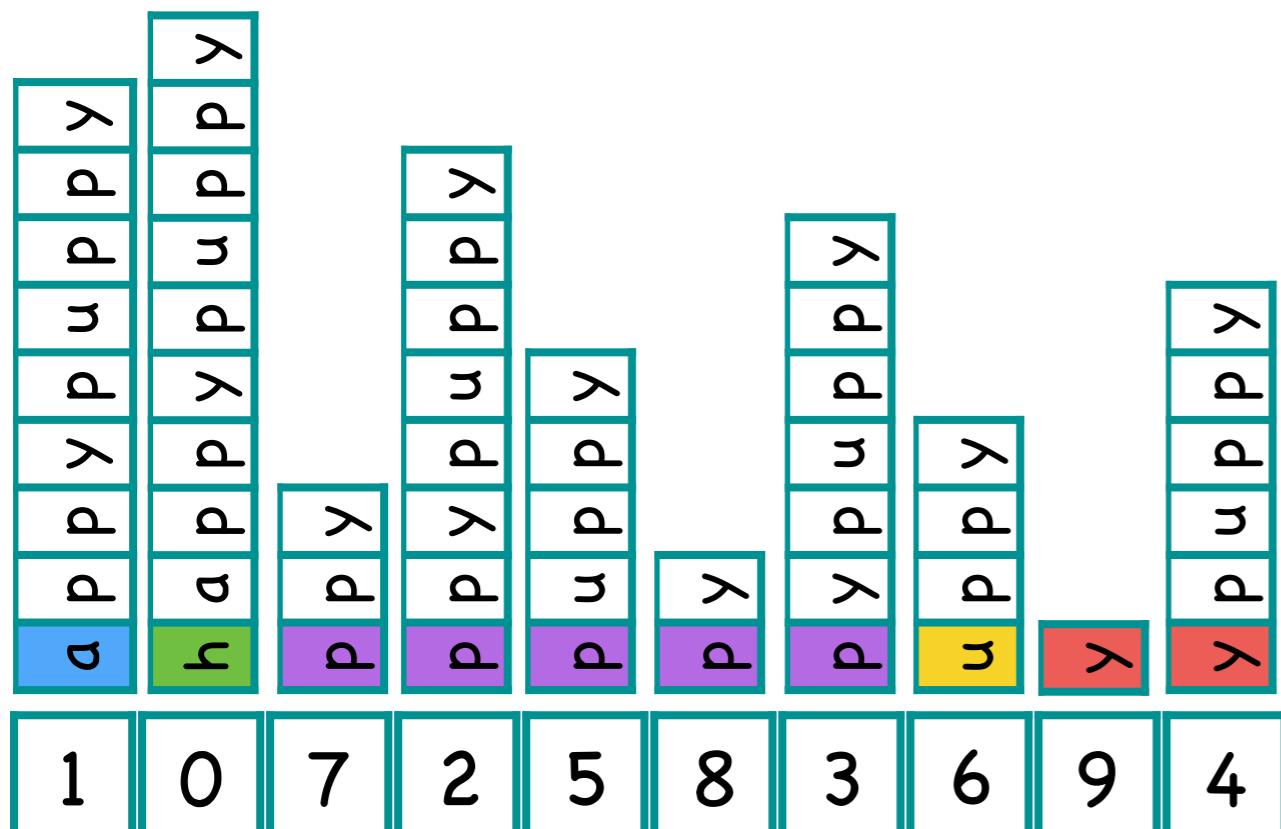


3 0      6      9

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



3 0

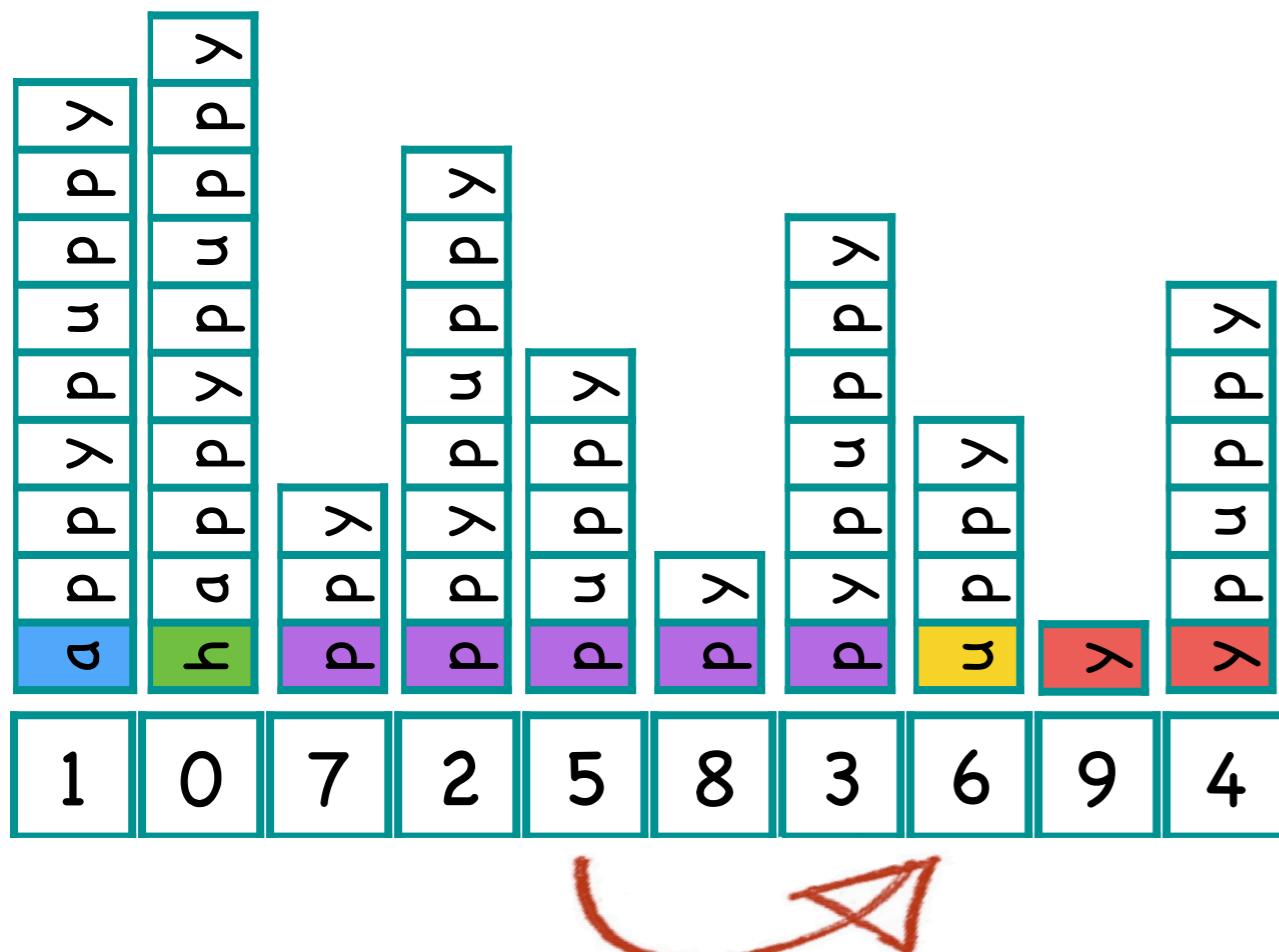
6

9

4

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



3	0
---	---

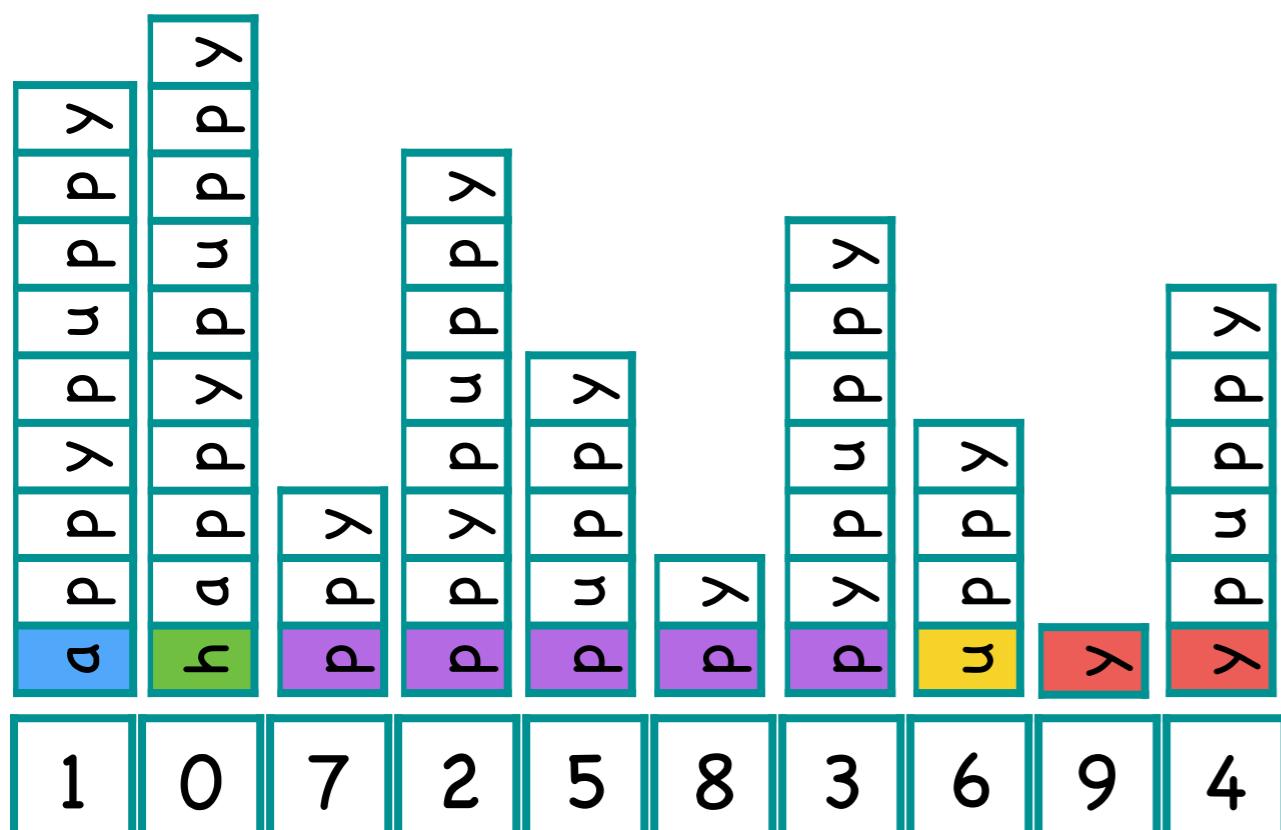
6	7
---	---

9
---

4
---

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

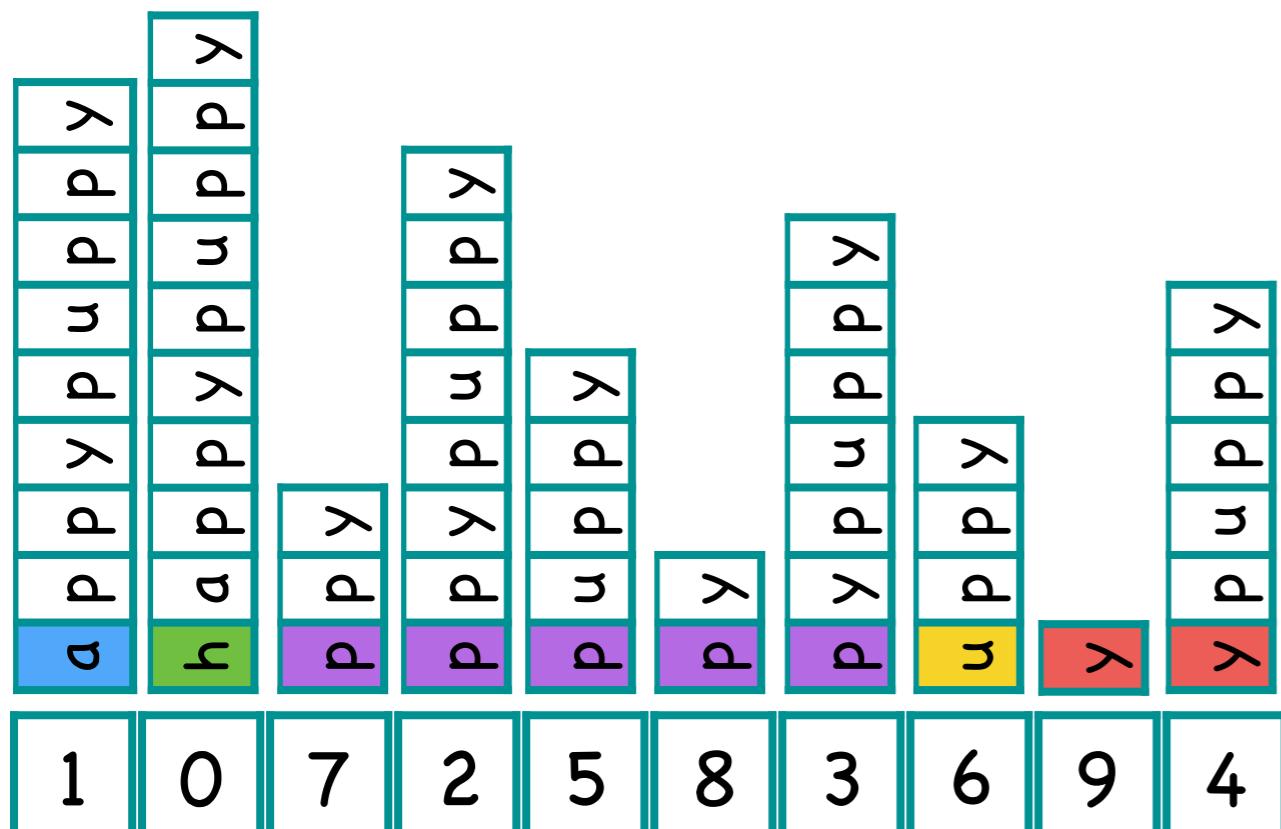


3	0
6	7
9	2
4	

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

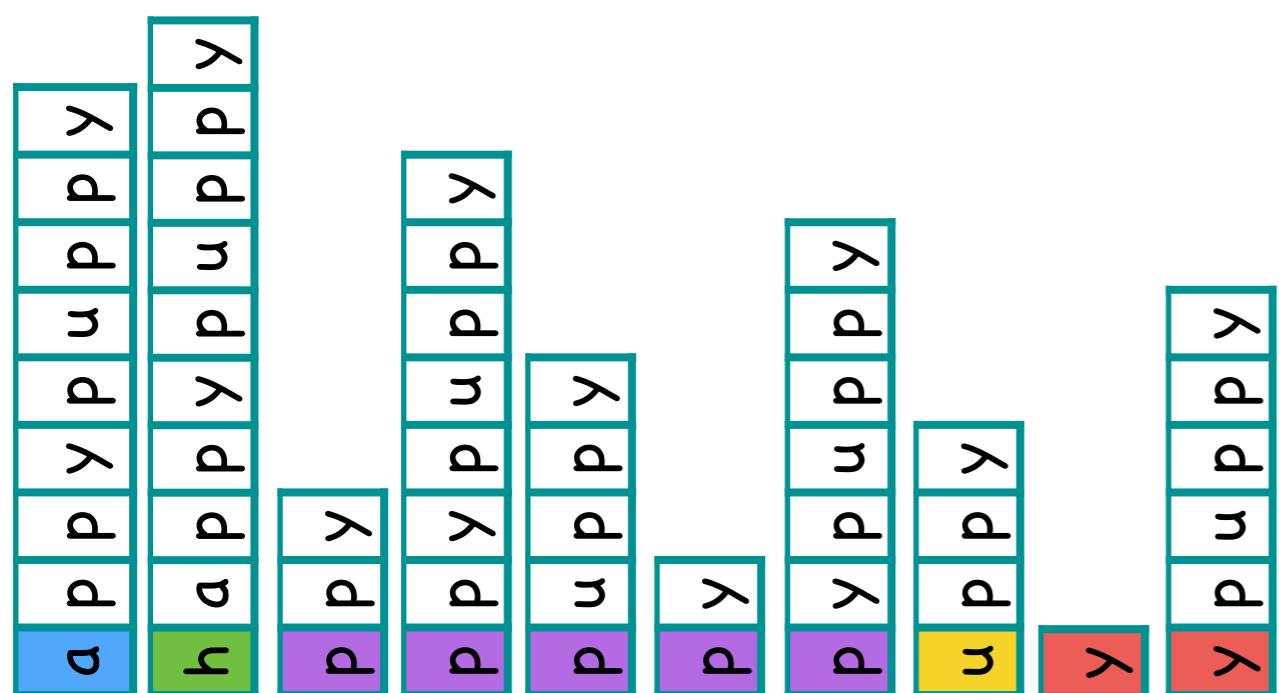


3	0	5	6	7
---	---	---	---	---

9	2
4	

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



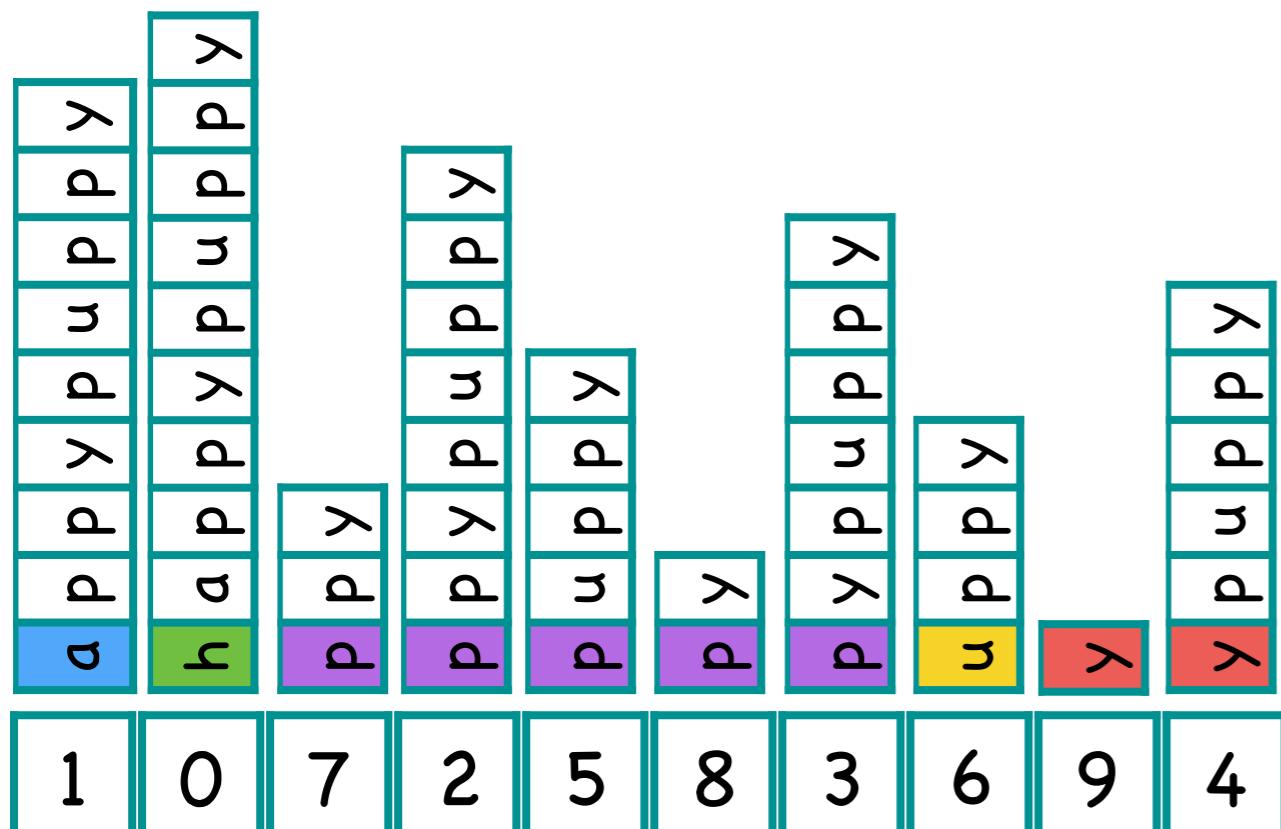
↗

3	0	5	6	7	8	9	2	4
---	---	---	---	---	---	---	---	---

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

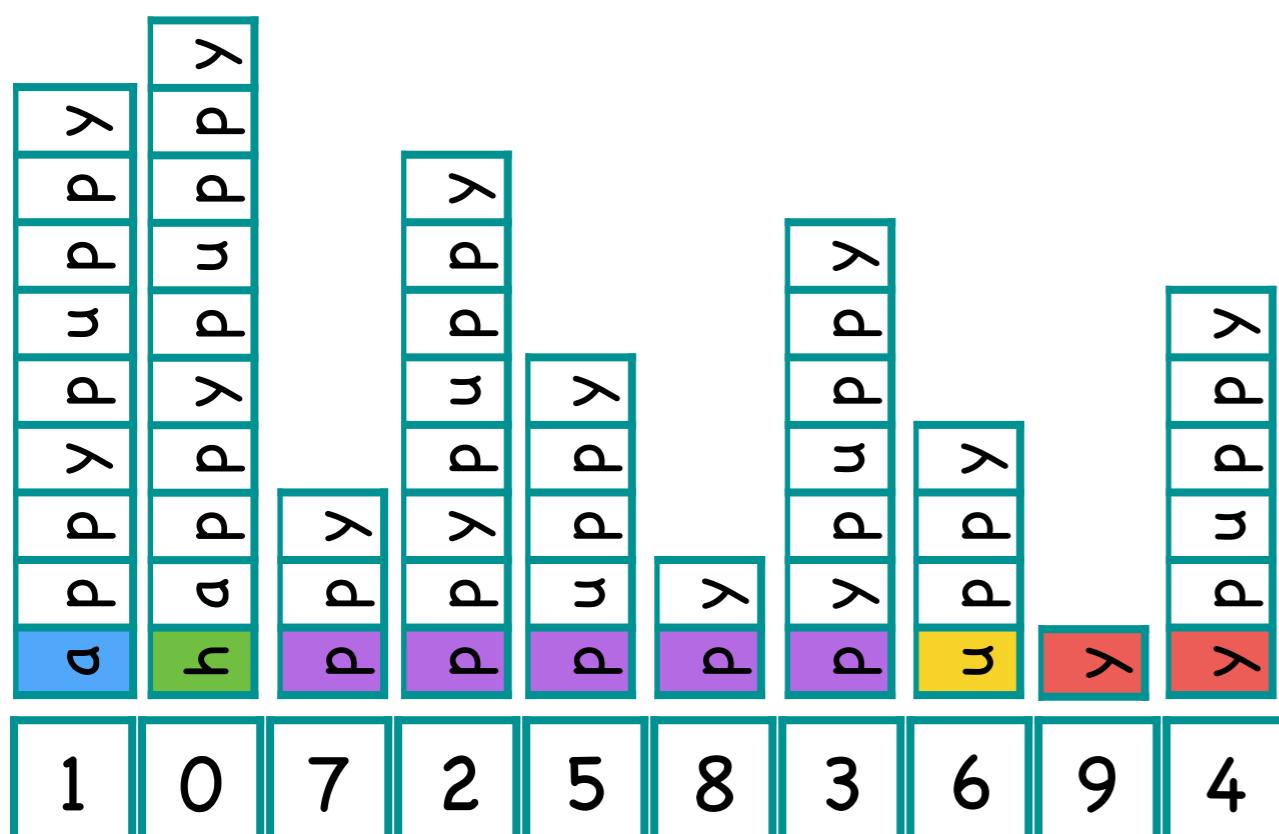


3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---

# Succinct Data Representation

---

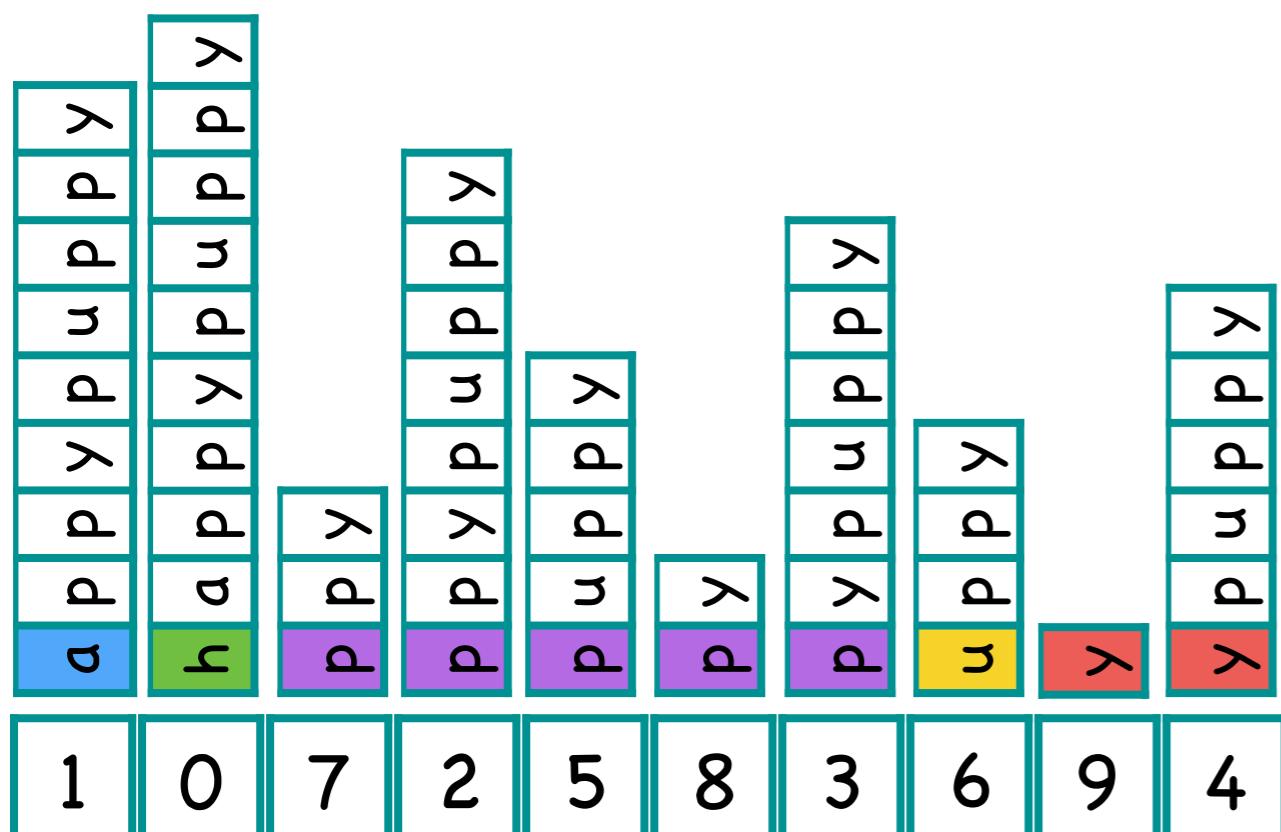
0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



NextCharIdx

(index of suffix after removing  
the first character)

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

σ	τ	ω	α	α	ω	ω	ω	τ	λ	λ
1	0	7	2	5	8	3	6	9	4	

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---



NextCharIdx  
(index of suffix after removing  
the first character)

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

σ	τ	ρ							
1	0	7	2	5	8	3	6	9	4

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---



NextCharIdx  
(index of suffix after removing  
the first character)

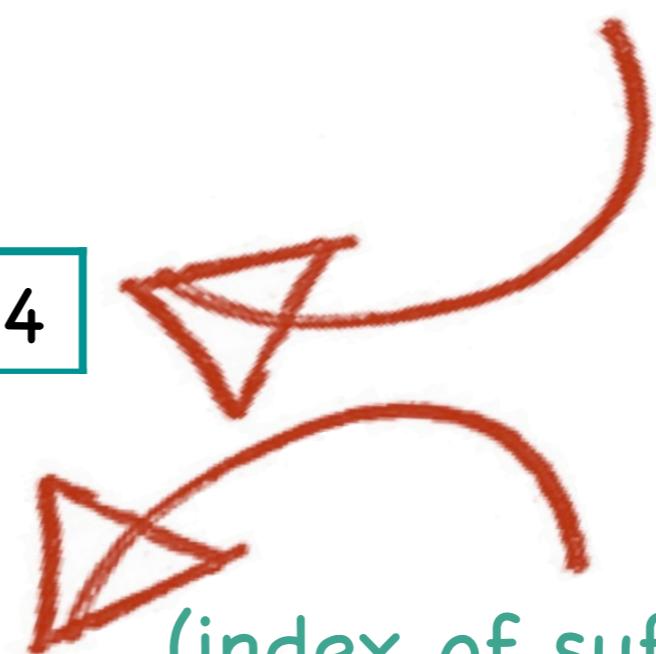
# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

No redundancy in AoS2Input

σ	ε	ρ							
1	0	7	2	5	8	3	6	9	4

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---



NextCharIdx  
(index of suffix after removing  
the first character)

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

σ	ε	ρ							
1	0	7	2	5	8	3	6	9	4

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---

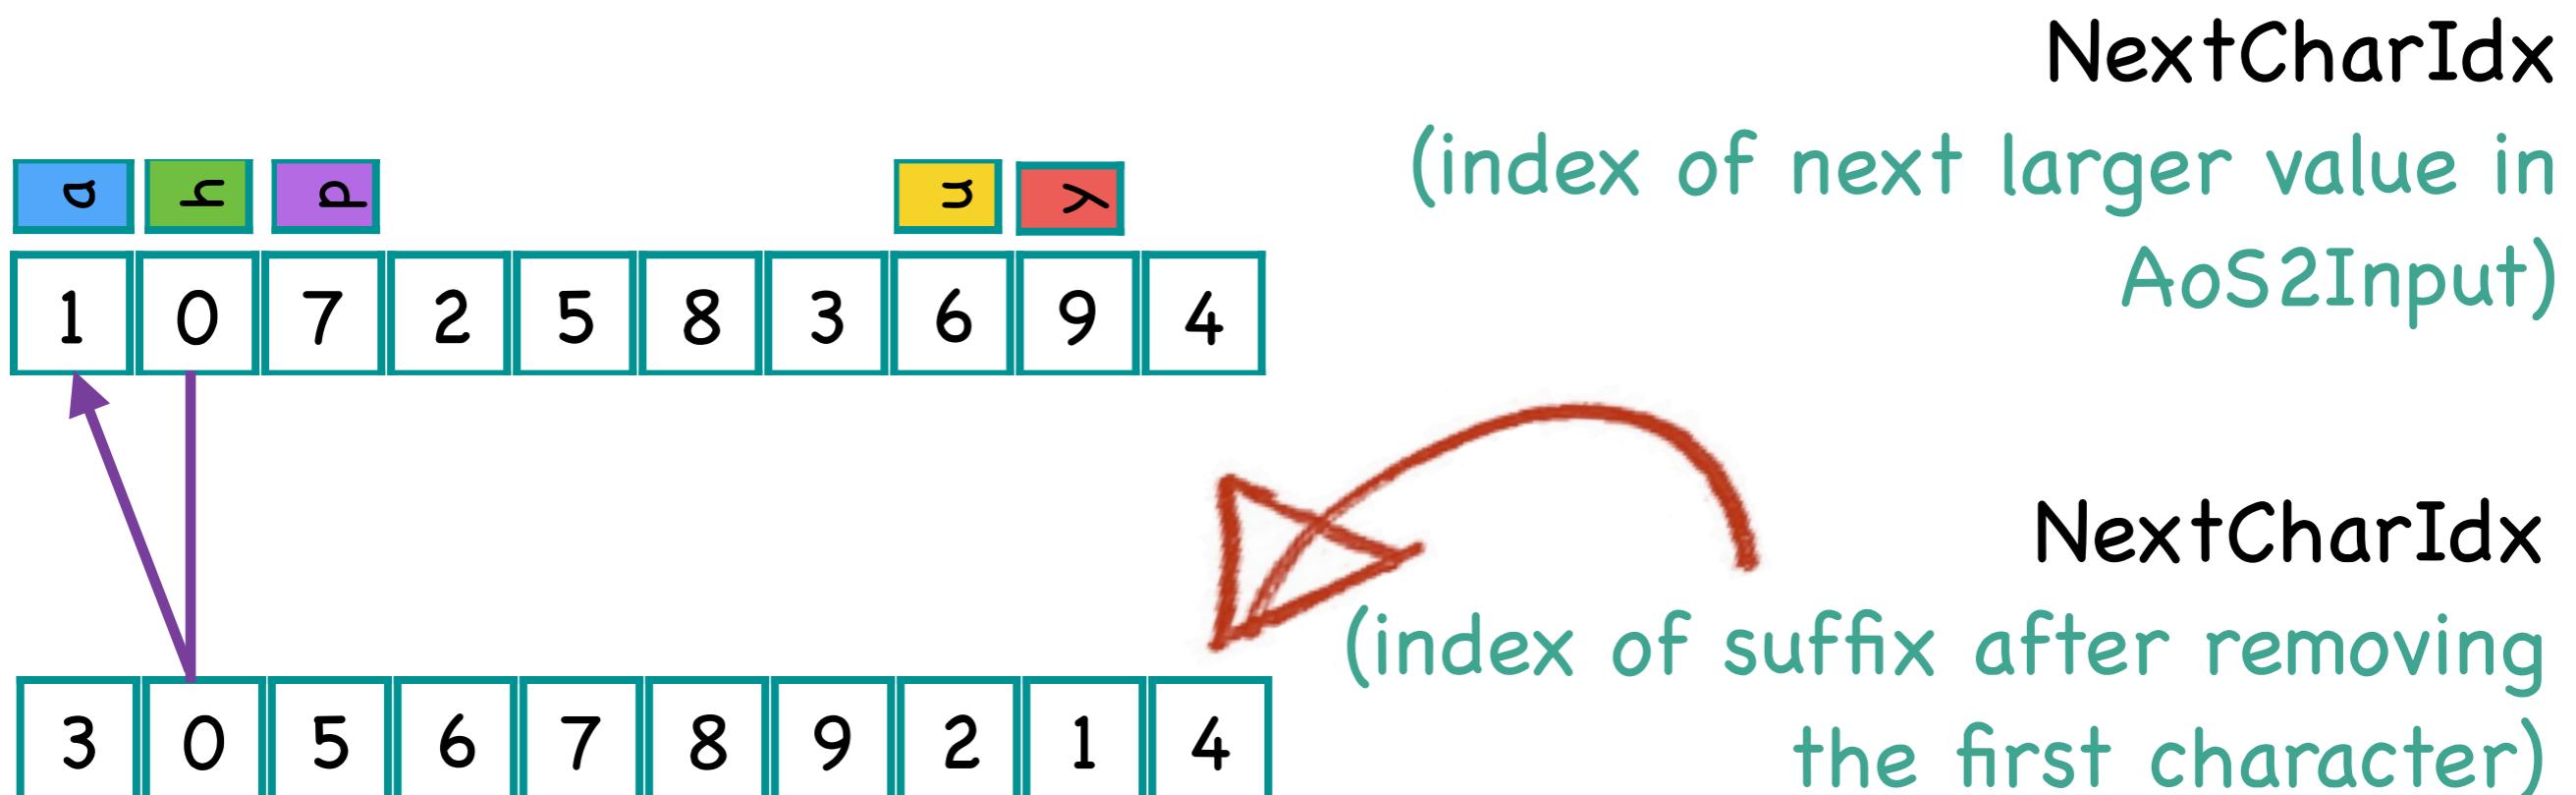
NextCharIdx  
(index of next larger value in AoS2Input)



NextCharIdx  
(index of suffix after removing the first character)

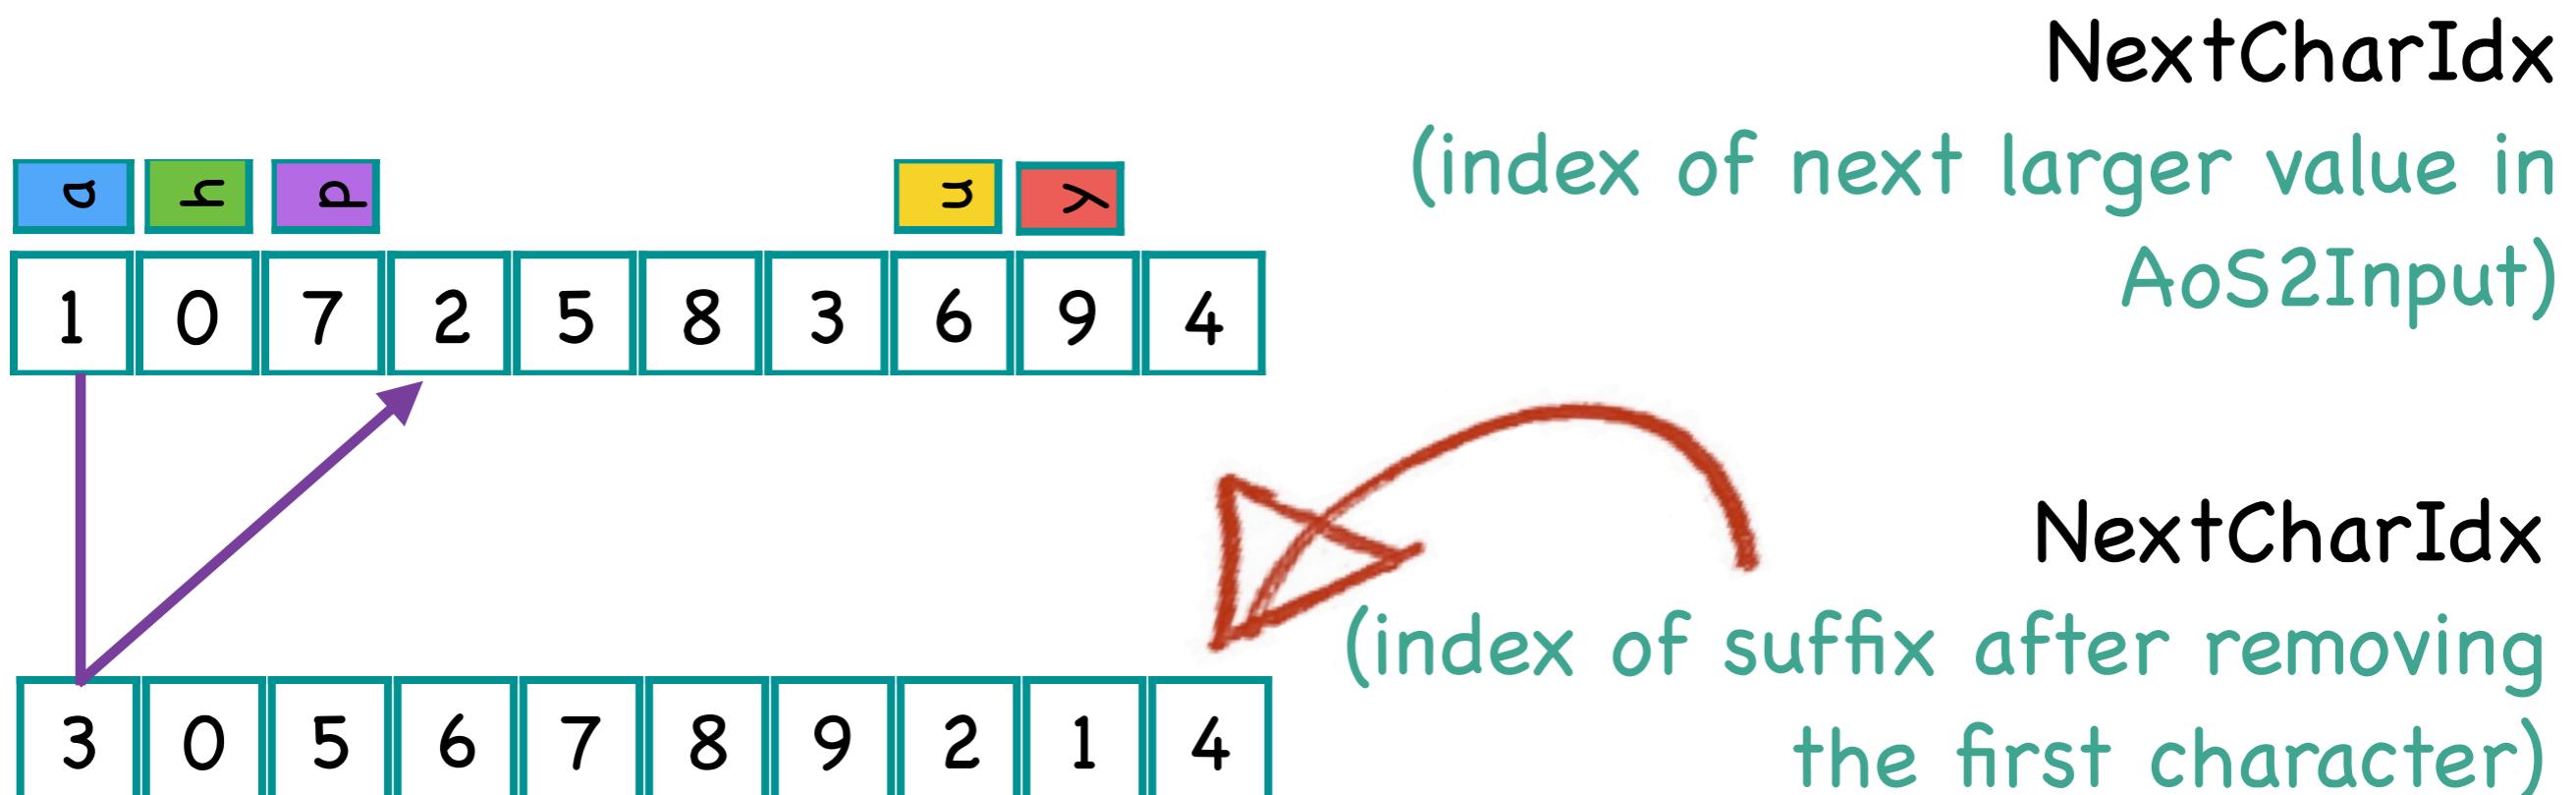
# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y



# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

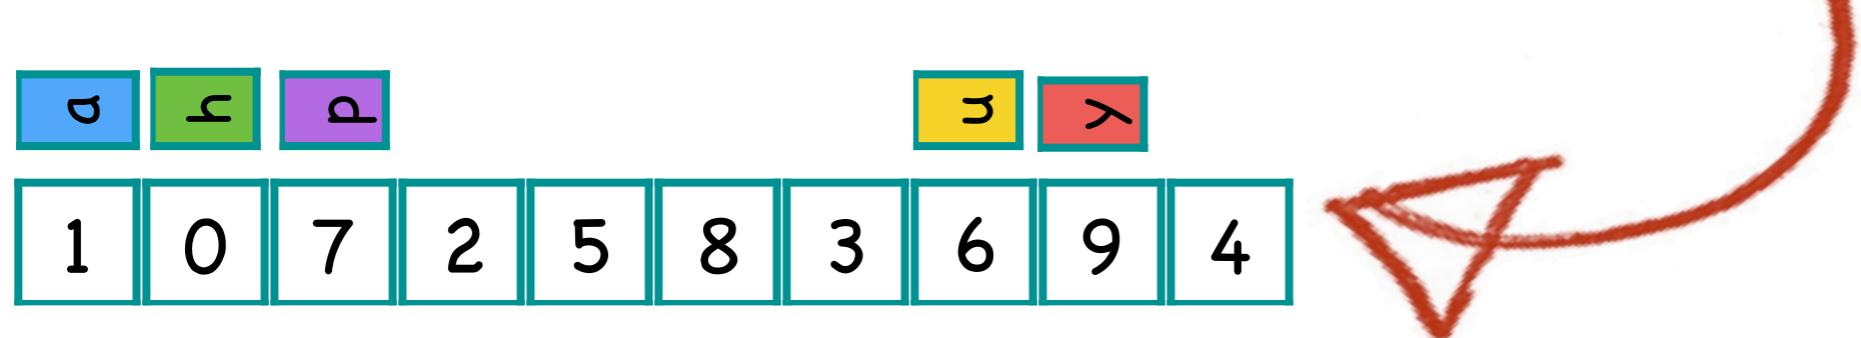


# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

No redundancy in AoS2Input



# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

σ	τ	ρ		δ	λ				
1	0	7	2	5	8	3	6	9	4

Solution: Sample!

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--

σ	τ	ρ							
1	0	7	2	5	8	3	6	9	4

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--

σ	ε	ρ							
1	0	7	2	5	8	3	6	9	4

Problem: How to find  
unsampled values?



# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--

σ	ε	ρ							
1	0	7	2	5	8	3	6	9	4

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---

Problem: How to find  
unsampled values?



Solution: Follow  
NextCharIdx pointers until  
you hit a sampled value

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--

σ	ε	ρ							
1	0	7	2	5	8	3	6	9	4

3	0	5	6	7	8	9	2	1	4
---	---	---	---	---	---	---	---	---	---

Problem: How to find  
unsampled values?

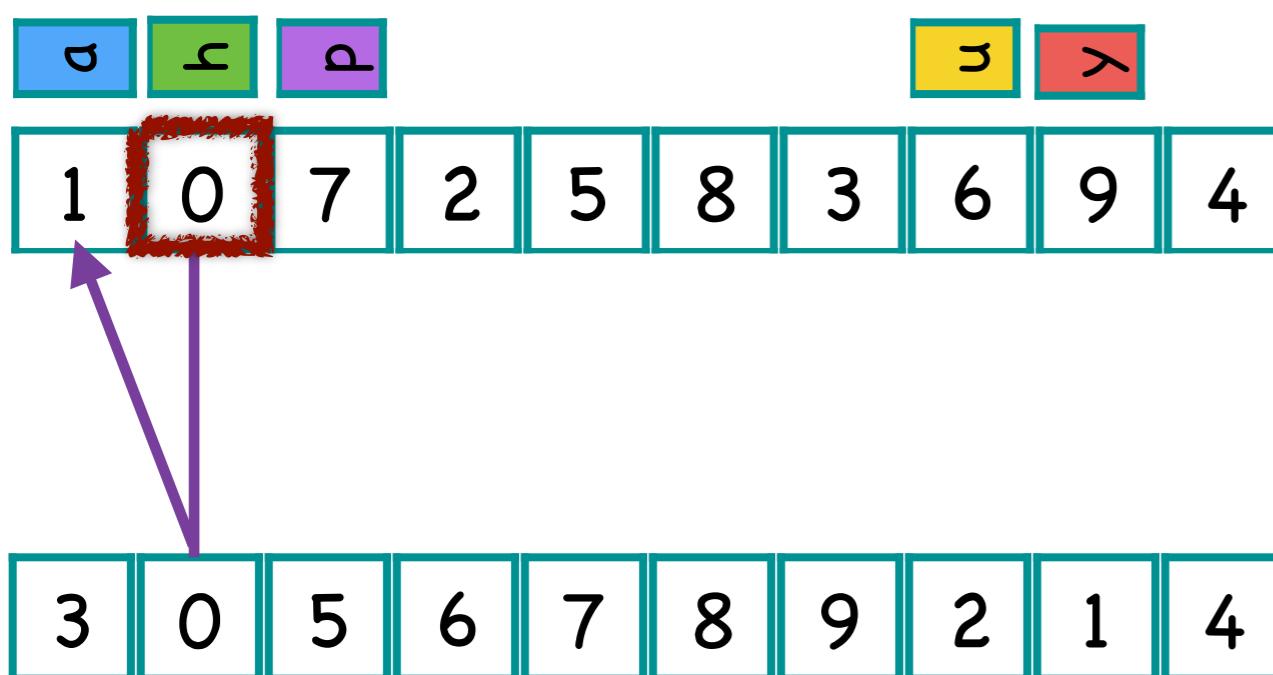


Solution: Follow  
NextCharIdx pointers until  
you hit a sampled value

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--



Problem: How to find  
unsampled values?

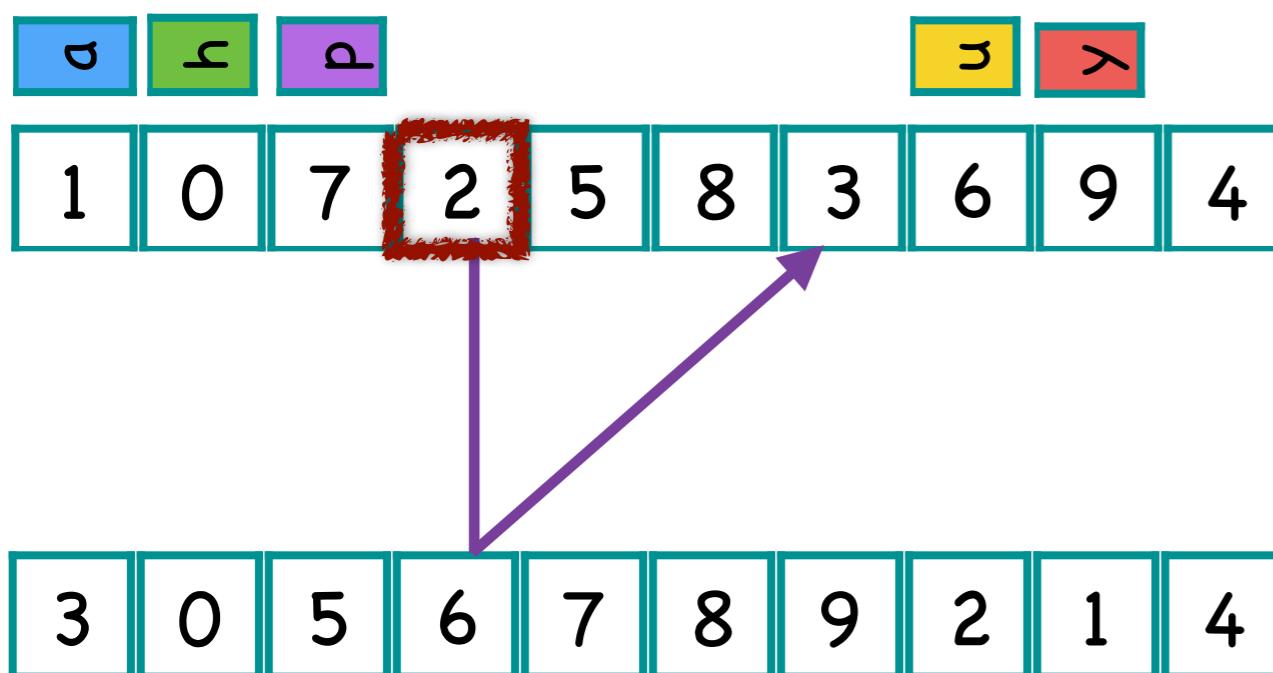


Solution: Follow  
NextCharIdx pointers until  
you hit a sampled value

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--



Problem: How to find  
unsampled values?



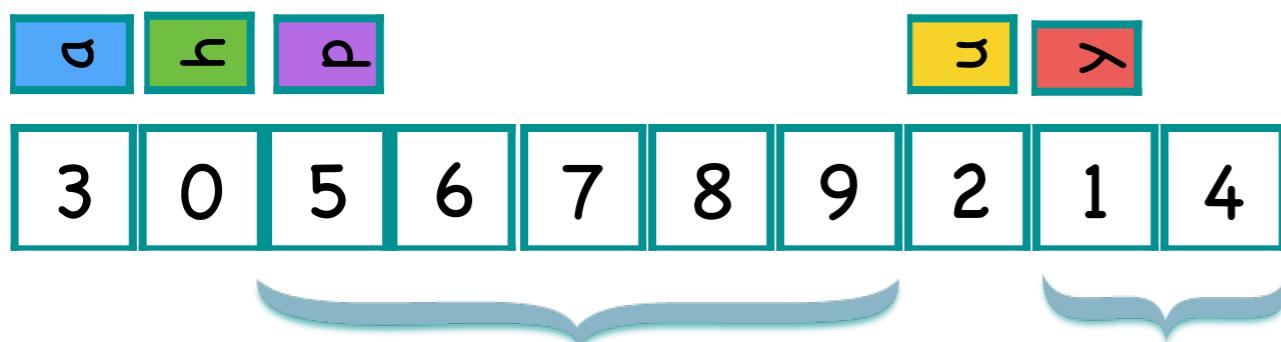
Solution: Follow  
NextCharIdx pointers until  
you hit a sampled value

# Succinct Data Representation

---

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--

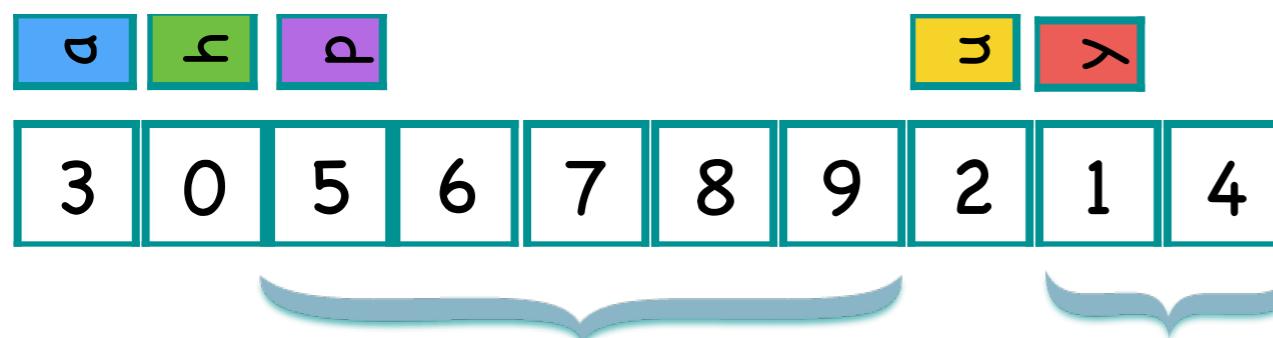


Increasing Integer  
Sequences

# Succinct Data Representation

0	1	2	3	4	5	6	7	8	9
h	a	p	p	y	p	u	p	p	y

1		7		5		3		9	
---	--	---	--	---	--	---	--	---	--



Increasing Integer  
Sequences

Can be compressed!

# Data Model and API

---

# Data Model and API

---

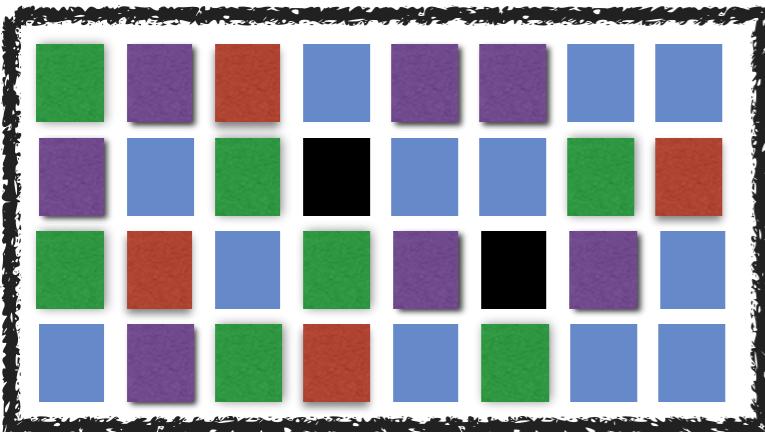
Input: flat (unstructured) files

# Data Model and API

---

Input: flat (unstructured) files

Original Input

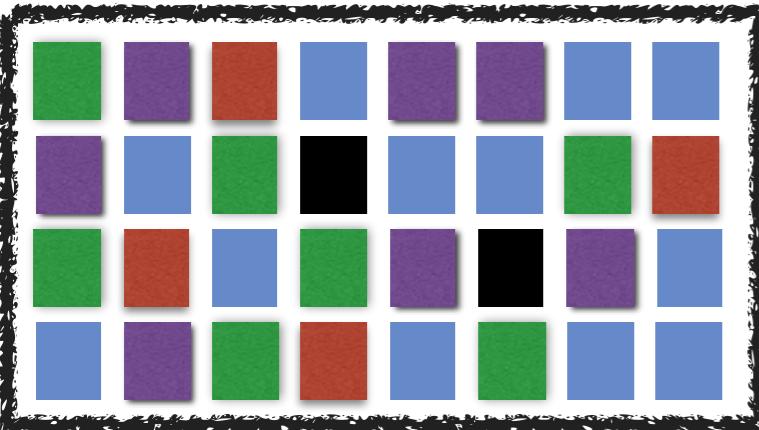


# Data Model and API

---

Input: flat (unstructured) files

Original Input



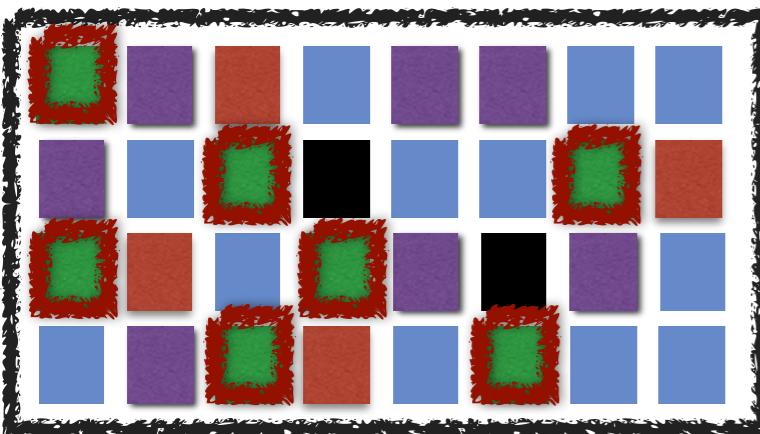
Succinct



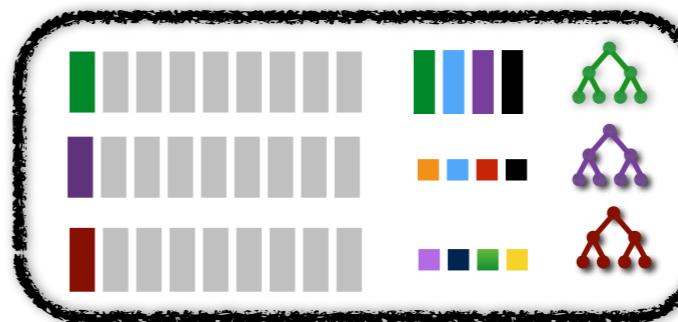
# Data Model and API

Input: flat (unstructured) files

Original Input



Succinct



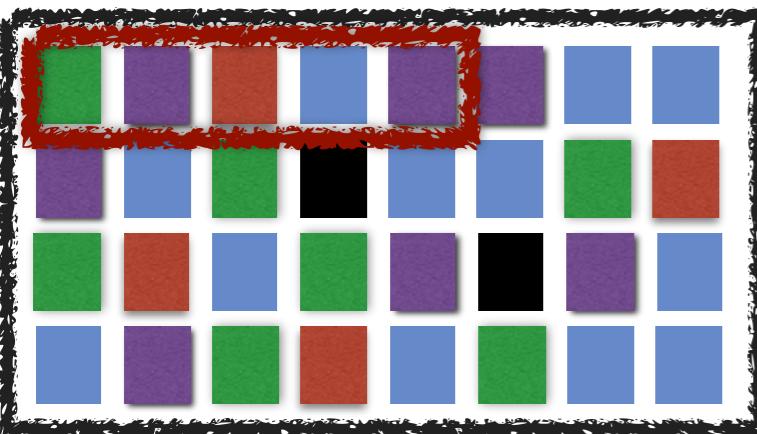
$\text{Search}(\blacksquare) = \{0, 10, 14, 16, 19, 26, 29\}$

Search: returns offsets of arbitrary strings in uncompressed file

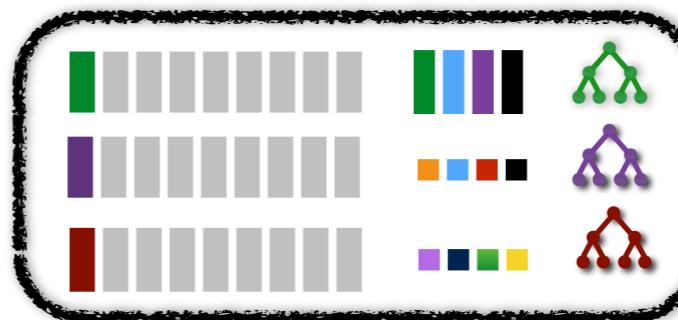
# Data Model and API

Input: flat (unstructured) files

Original Input



Succinct



$\text{Search}(\blacksquare) = \{0, 10, 14, 16, 19, 26, 29\}$

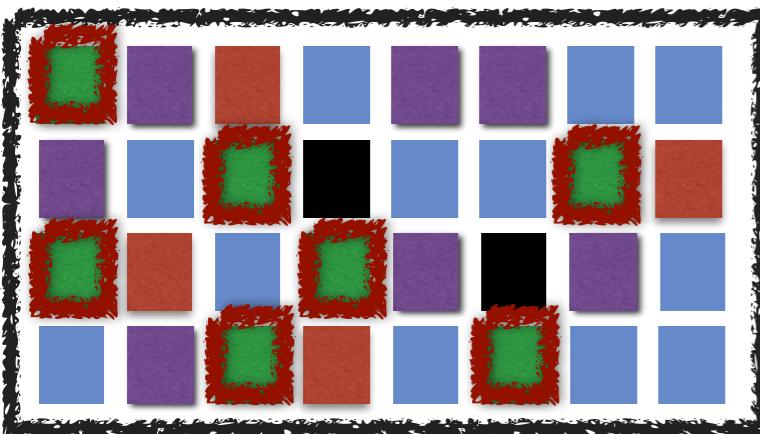
$\text{Extract}(0, 5) = \{\blacksquare, \blacksquare, \blacksquare, \blacksquare, \blacksquare\}$

Extract: returns data at arbitrary offsets in uncompressed file

# Data Model and API

Input: flat (unstructured) files

Original Input



Succinct



$\text{Search}(\blacksquare) = \{0, 10, 14, 16, 19, 26, 29\}$

$\text{Extract}(0, 5) = \{\blacksquare, \blacksquare, \blacksquare, \blacksquare, \blacksquare\}$

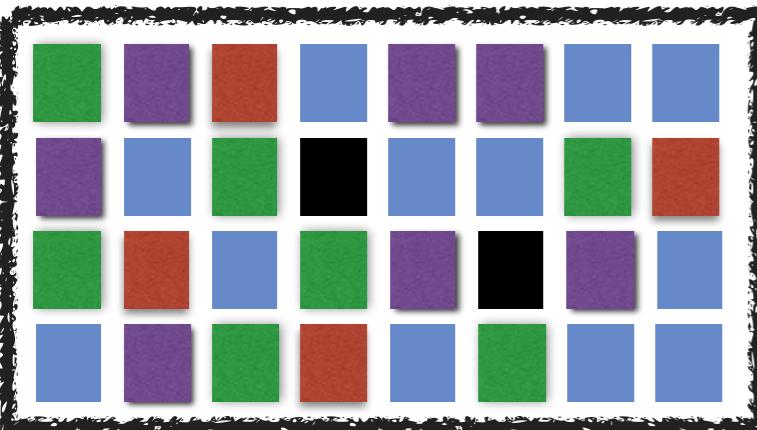
$\text{Count}(\blacksquare) = 7$

| Count: returns count of arbitrary strings in uncompressed file

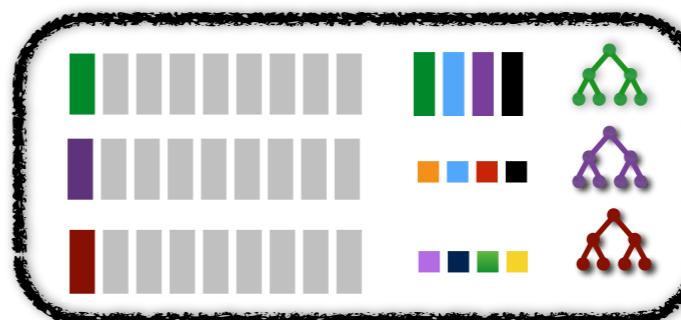
# Data Model and API

Input: flat (unstructured) files

Original Input



Succinct



$\text{Search}(\blacksquare) = \{0, 10, 14, 16, 19, 26, 29\}$

$\text{Extract}(0, 5) = \{\blacksquare, \blacksquare, \blacksquare, \blacksquare, \blacksquare\}$

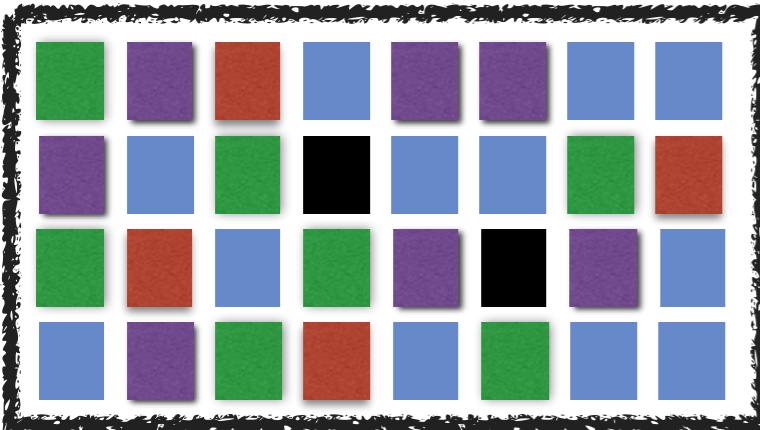
$\text{Count}(\blacksquare) = 7$

$\text{Append}(\blacksquare, \blacksquare, \blacksquare, \blacksquare, \blacksquare)$

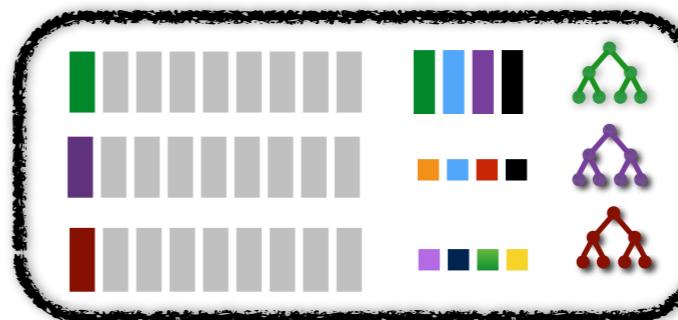
# Data Model and API

Input: flat (unstructured) files

Original Input



Succinct



$\text{Search}(\blacksquare) = \{0, 10, 14, 16, 19, 26, 29\}$

$\text{Extract}(0, 5) = \{\blacksquare, \blacksquare, \blacksquare, \blacksquare, \blacksquare\}$

$\text{Count}(\blacksquare) = 7$

$\text{Append}(\blacksquare, \blacksquare, \blacksquare, \blacksquare, \blacksquare)$

Range and Wildcard queries

# Power of queries on flat files

---

Why flat unstructured files?

# Power of queries on flat files

---

## Why flat unstructured files?

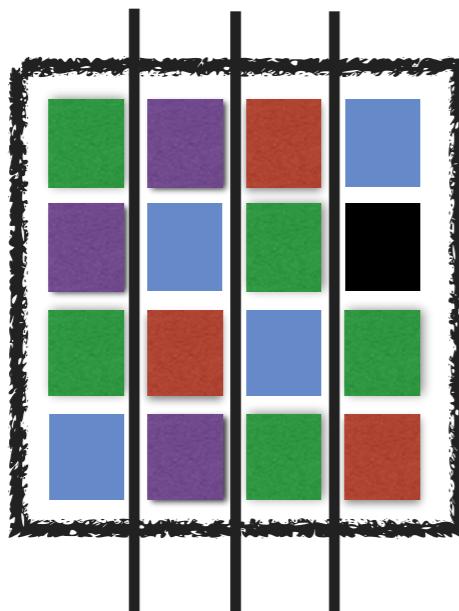
- Many powerful abstractions on top
  - Key-value store [Dynamo, MICA]
  - Tables [Cassandra, BigTable]
  - Documents [MongoDB]

# Power of queries on flat files

---

## Why flat unstructured files?

- Many powerful abstractions on top
  - Key-value store [Dynamo, MICA]
  - Tables [Cassandra, BigTable]
  - Documents [MongoDB]

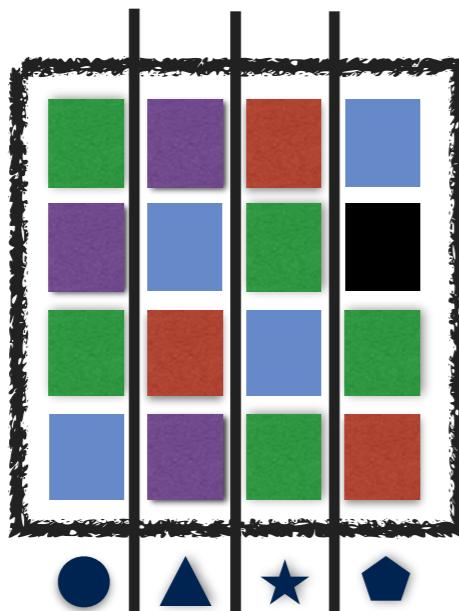


# Power of queries on flat files

---

## Why flat unstructured files?

- Many powerful abstractions on top
  - Key-value store [Dynamo, MICA]
  - Tables [Cassandra, BigTable]
  - Documents [MongoDB]

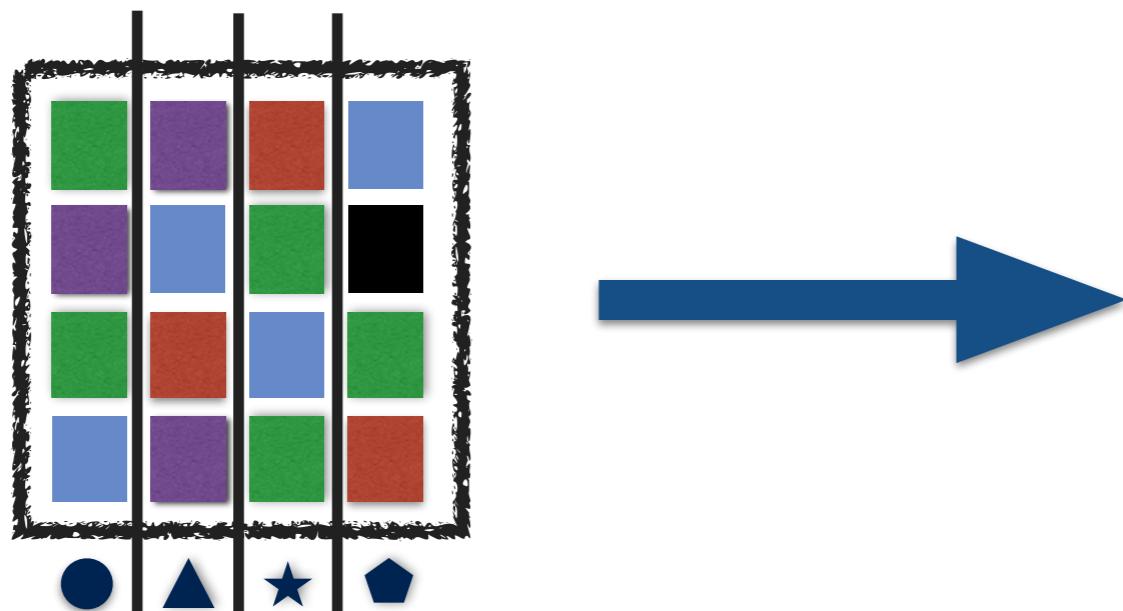


# Power of queries on flat files

---

## Why flat unstructured files?

- Many powerful abstractions on top
  - Key-value store [Dynamo, MICA]
  - Tables [Cassandra, BigTable]
  - Documents [MongoDB]

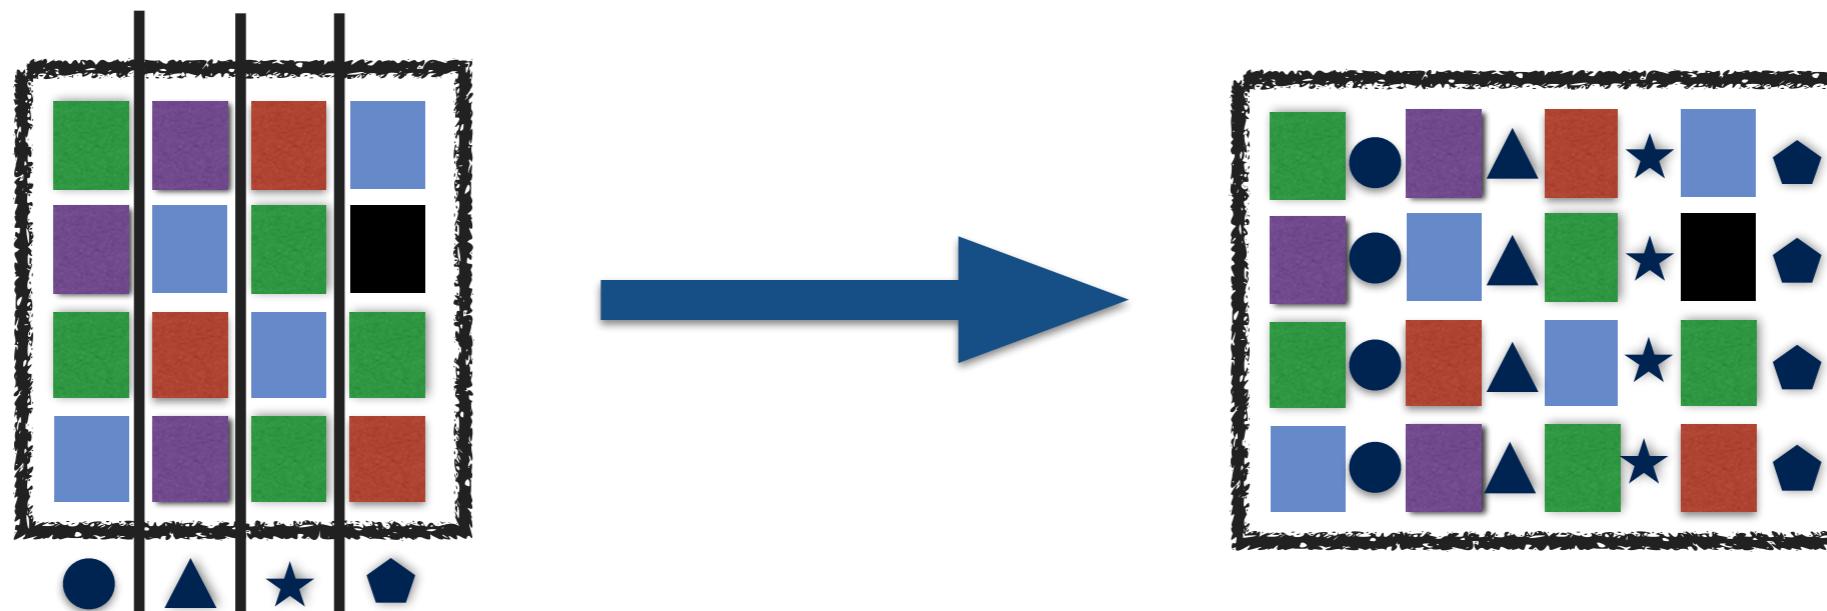


# Power of queries on flat files

---

## Why flat unstructured files?

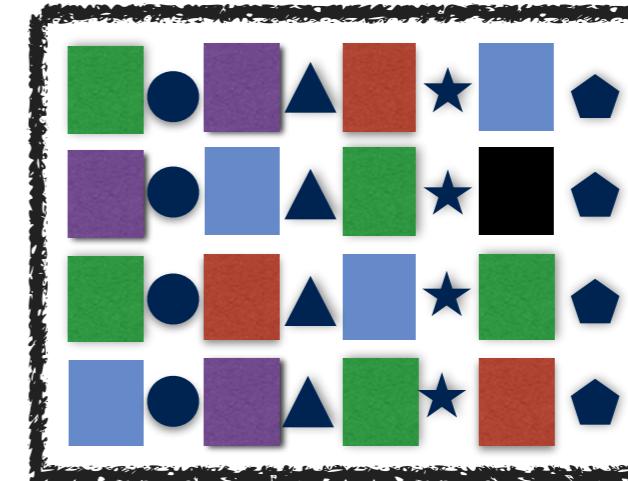
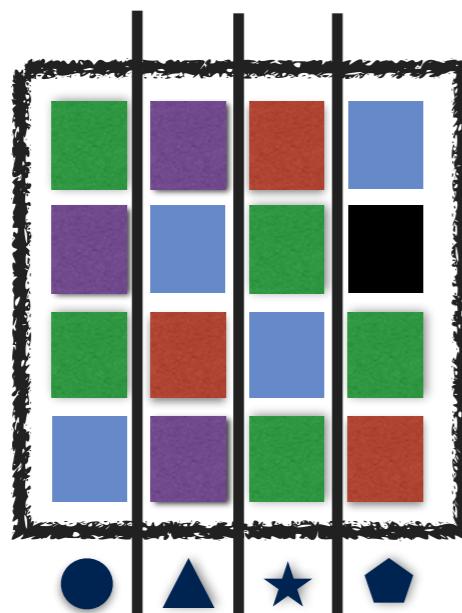
- Many powerful abstractions on top
  - Key-value store [Dynamo, MICA]
  - Tables [Cassandra, BigTable]
  - Documents [MongoDB]



# Power of queries on flat files

## Why flat unstructured files?

- Many powerful abstractions on top
  - Key-value store [Dynamo, MICA]
  - Tables [Cassandra, BigTable]
  - Documents [MongoDB]



Search(Column1, ■ )

Search( ■ ●)

# Handling Appends

---

Fine-grained updates of compressed data challenging

- Succinct does not support in-place updates
- Uses a multi-store architecture for appends
- Similar to SILT [SOSP'11]

# Handling Appends

---

Fine-grained updates of compressed data challenging

- Succinct does not support in-place updates
- Uses a multi-store architecture for appends
- Similar to SILT [SOSP'11]

Challenge

# Handling Appends

---

## Fine-grained updates of compressed data challenging

- Succinct does not support in-place updates
- Uses a multi-store architecture for appends
- Similar to SILT [SOSP'11]

## Challenge

- Multi-store architectures not new
- Challenge: supporting queries efficiently
- Details in paper

# End-to-end System

---

# End-to-end System

---



LogStore

# End-to-end System

---



SuffixStore



LogStore

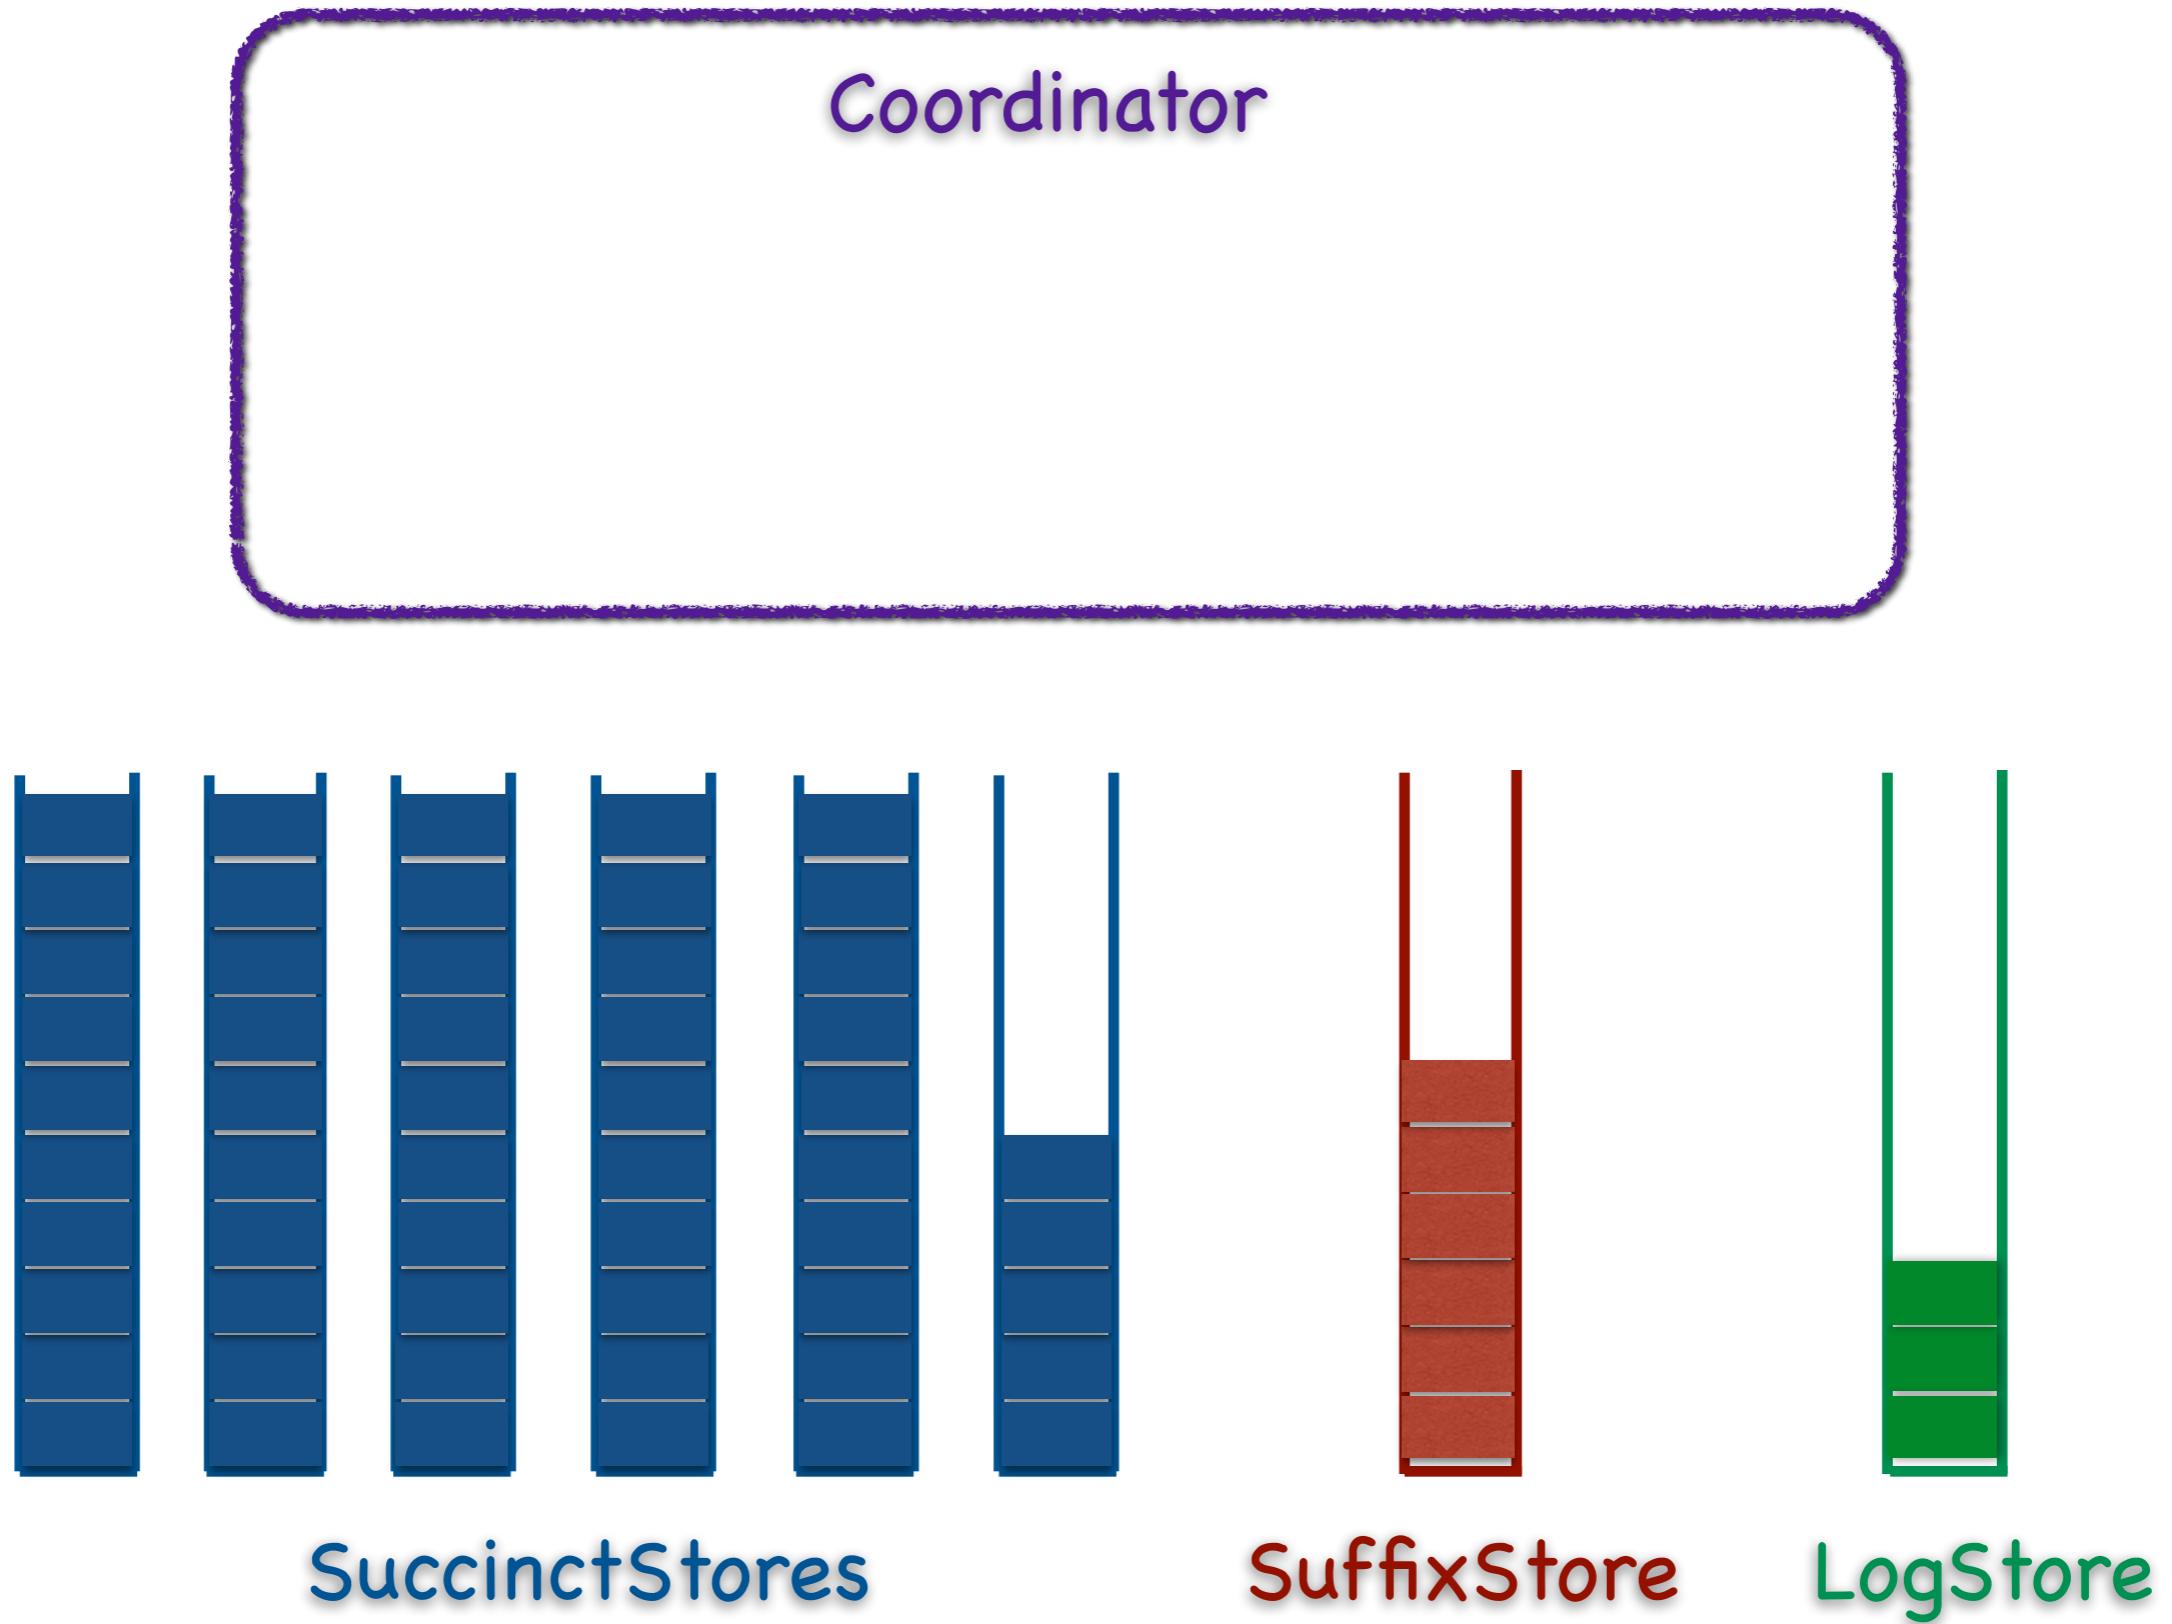
# End-to-end System

---



# End-to-end System

---



# End-to-end System

## Coordinator

### Cluster Management

- Maintains lists of active servers
- Periodic heartbeats



# End-to-end System

Coordinator

## Data Management

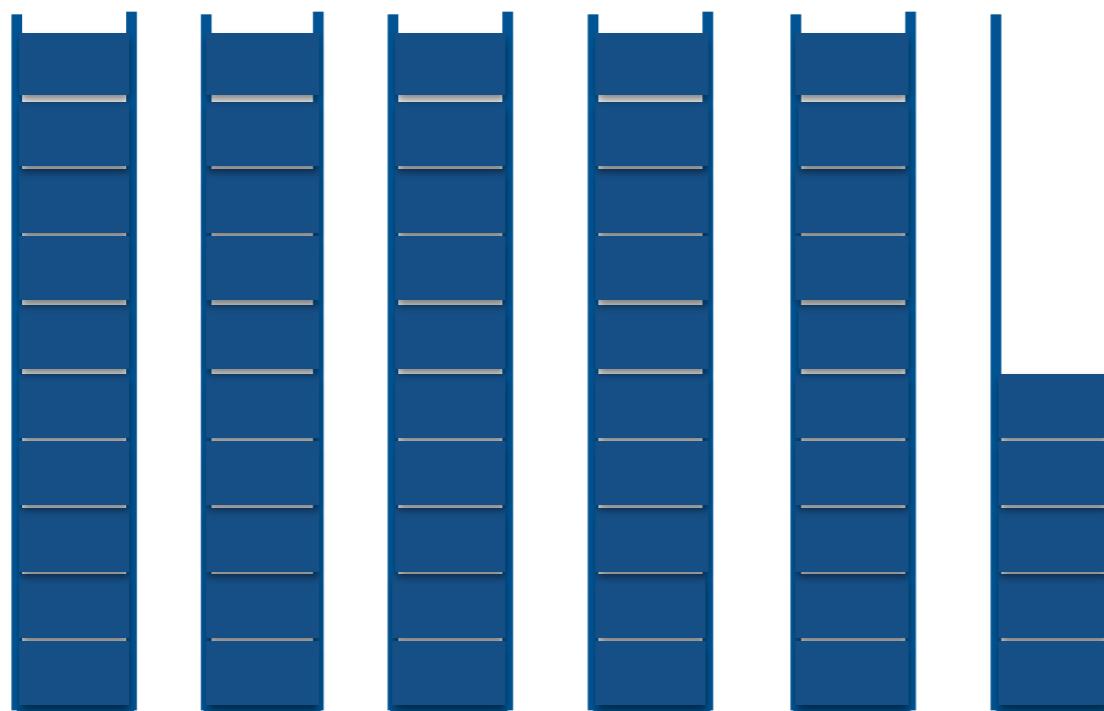
- Updated pointers to locate data
- Pushes these pointers to servers



# End-to-end System

---

Coordinator



SuccinctStores



SuffixStore

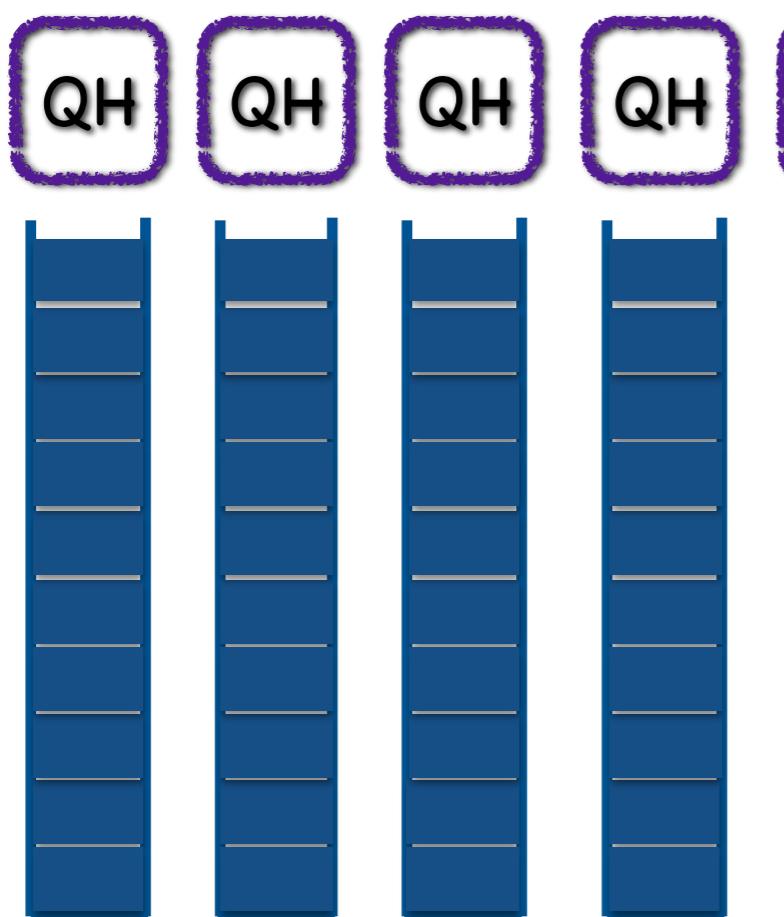


LogStore

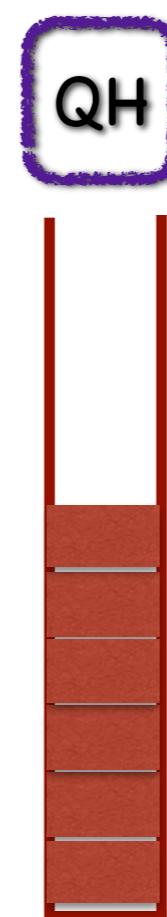
# End-to-end System

---

Coordinator



SuccinctStores

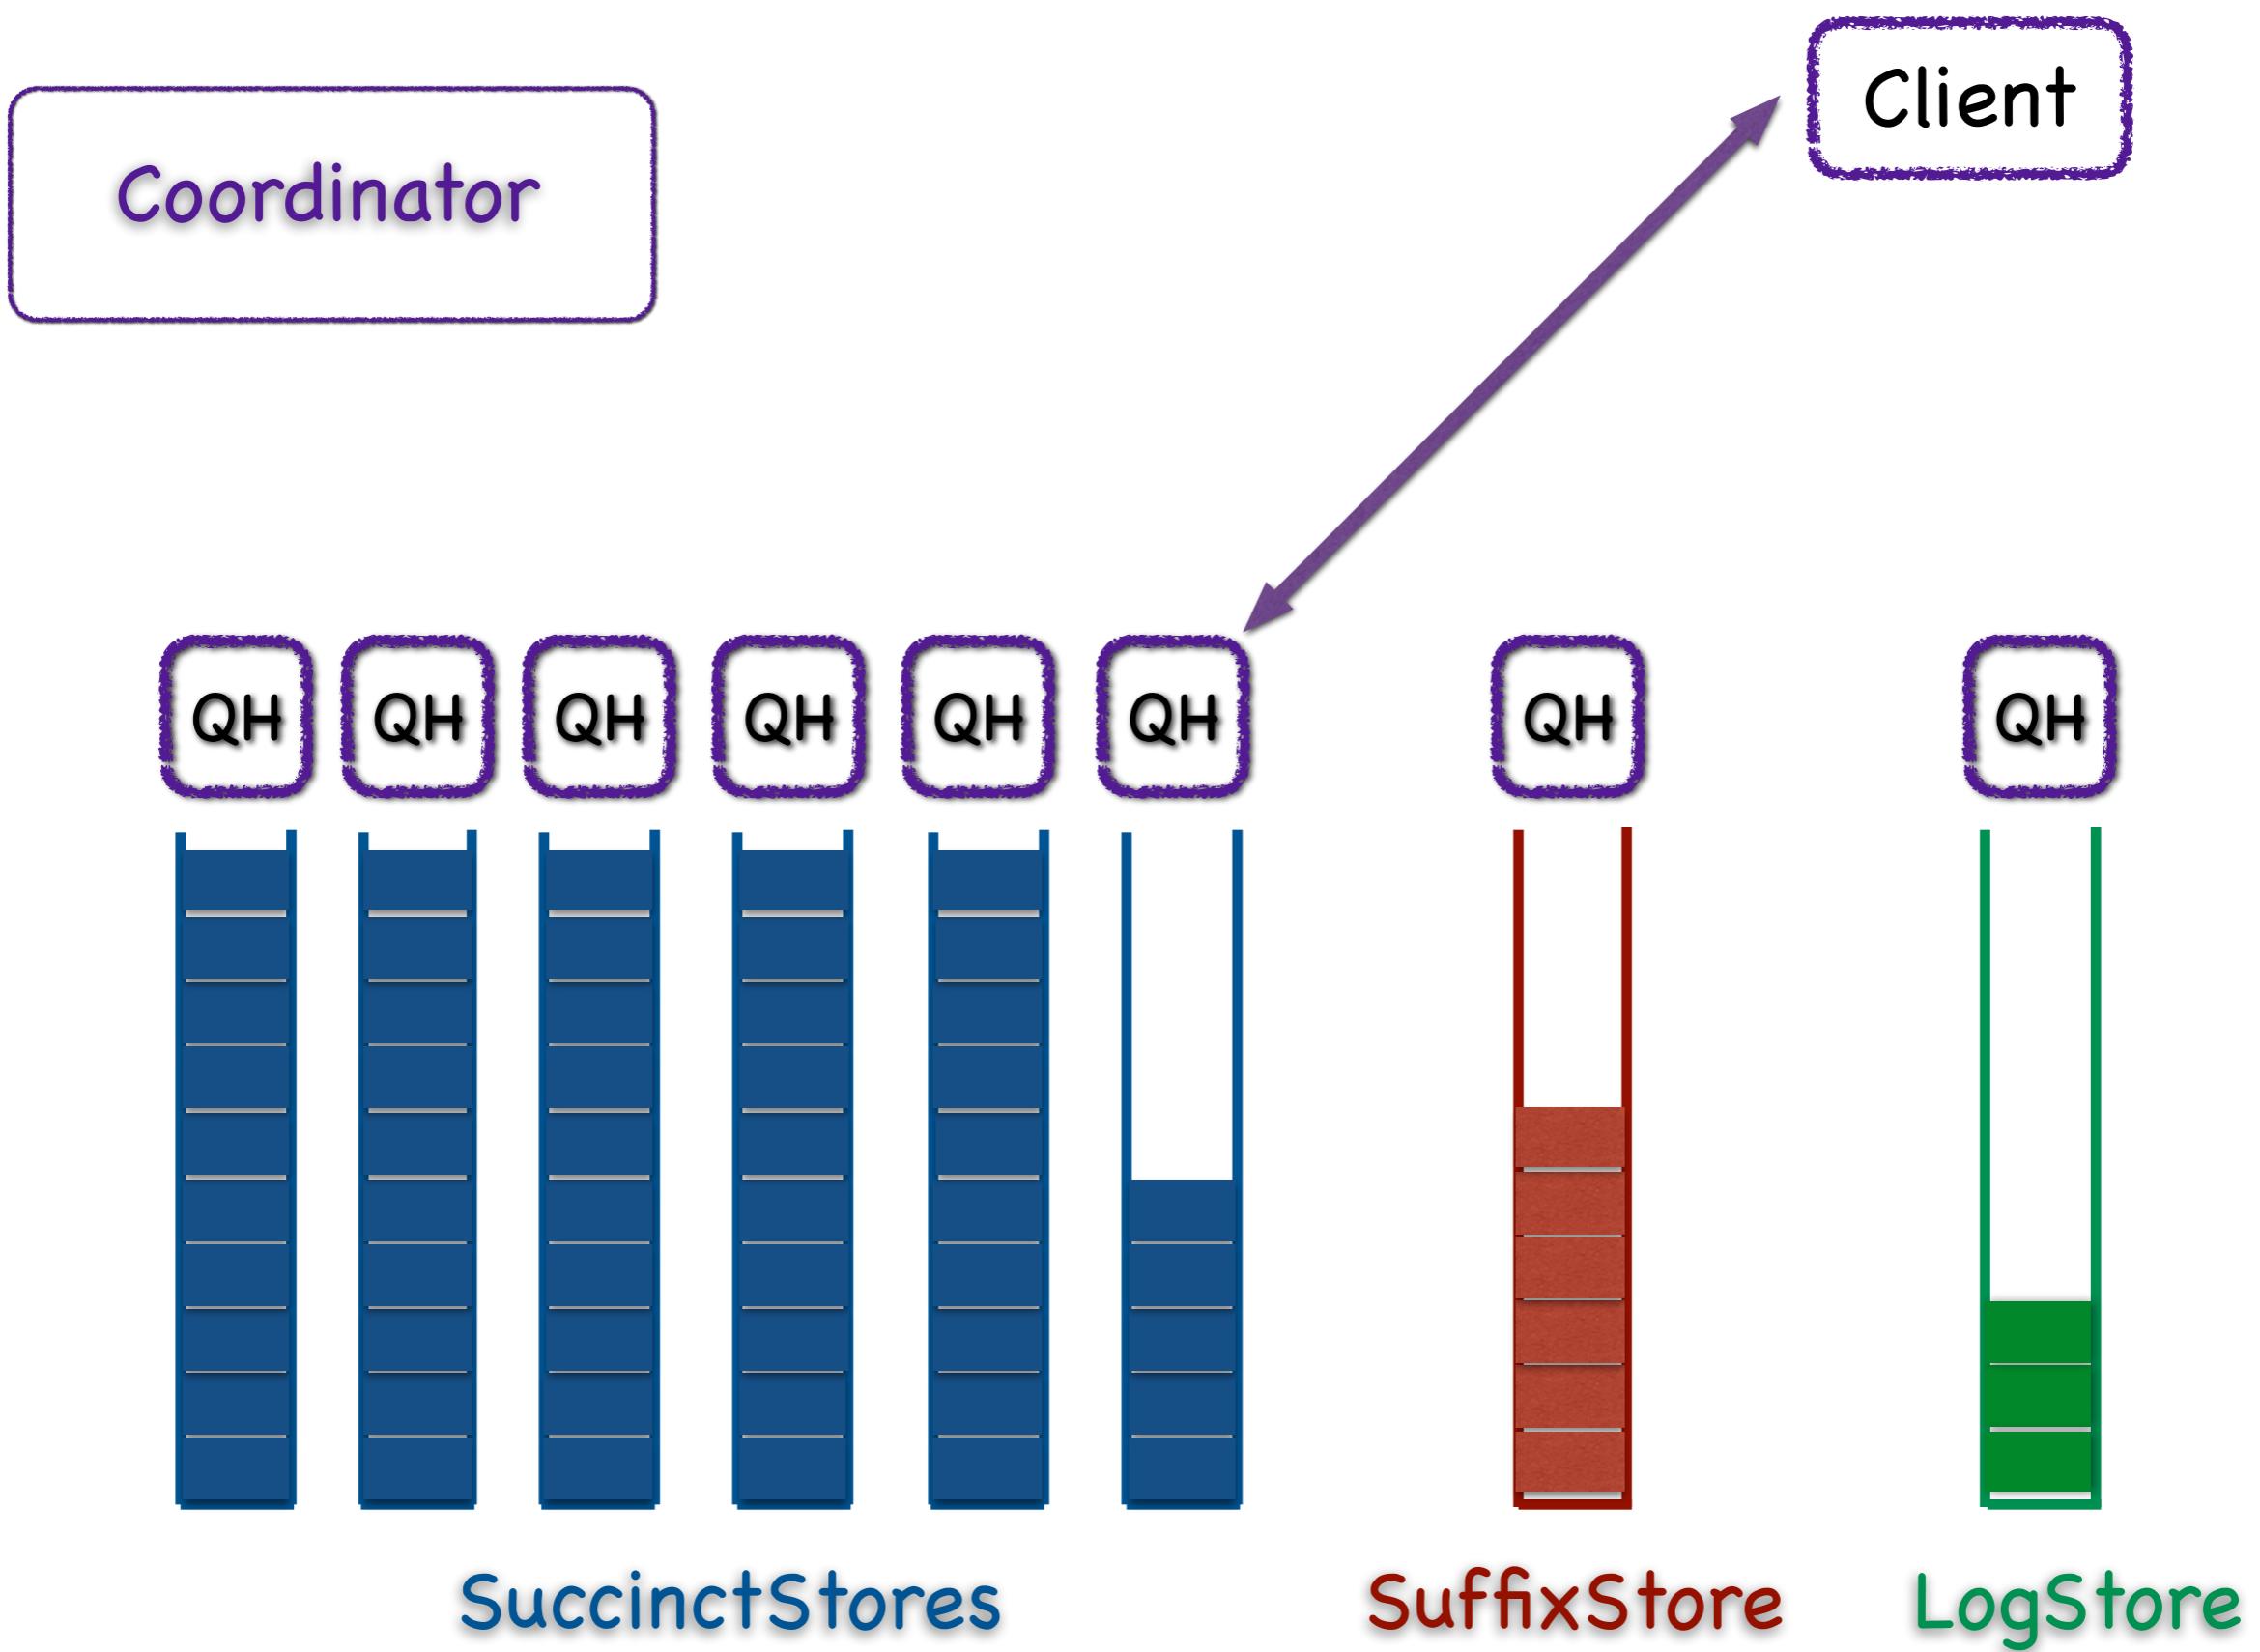


SuffixStore

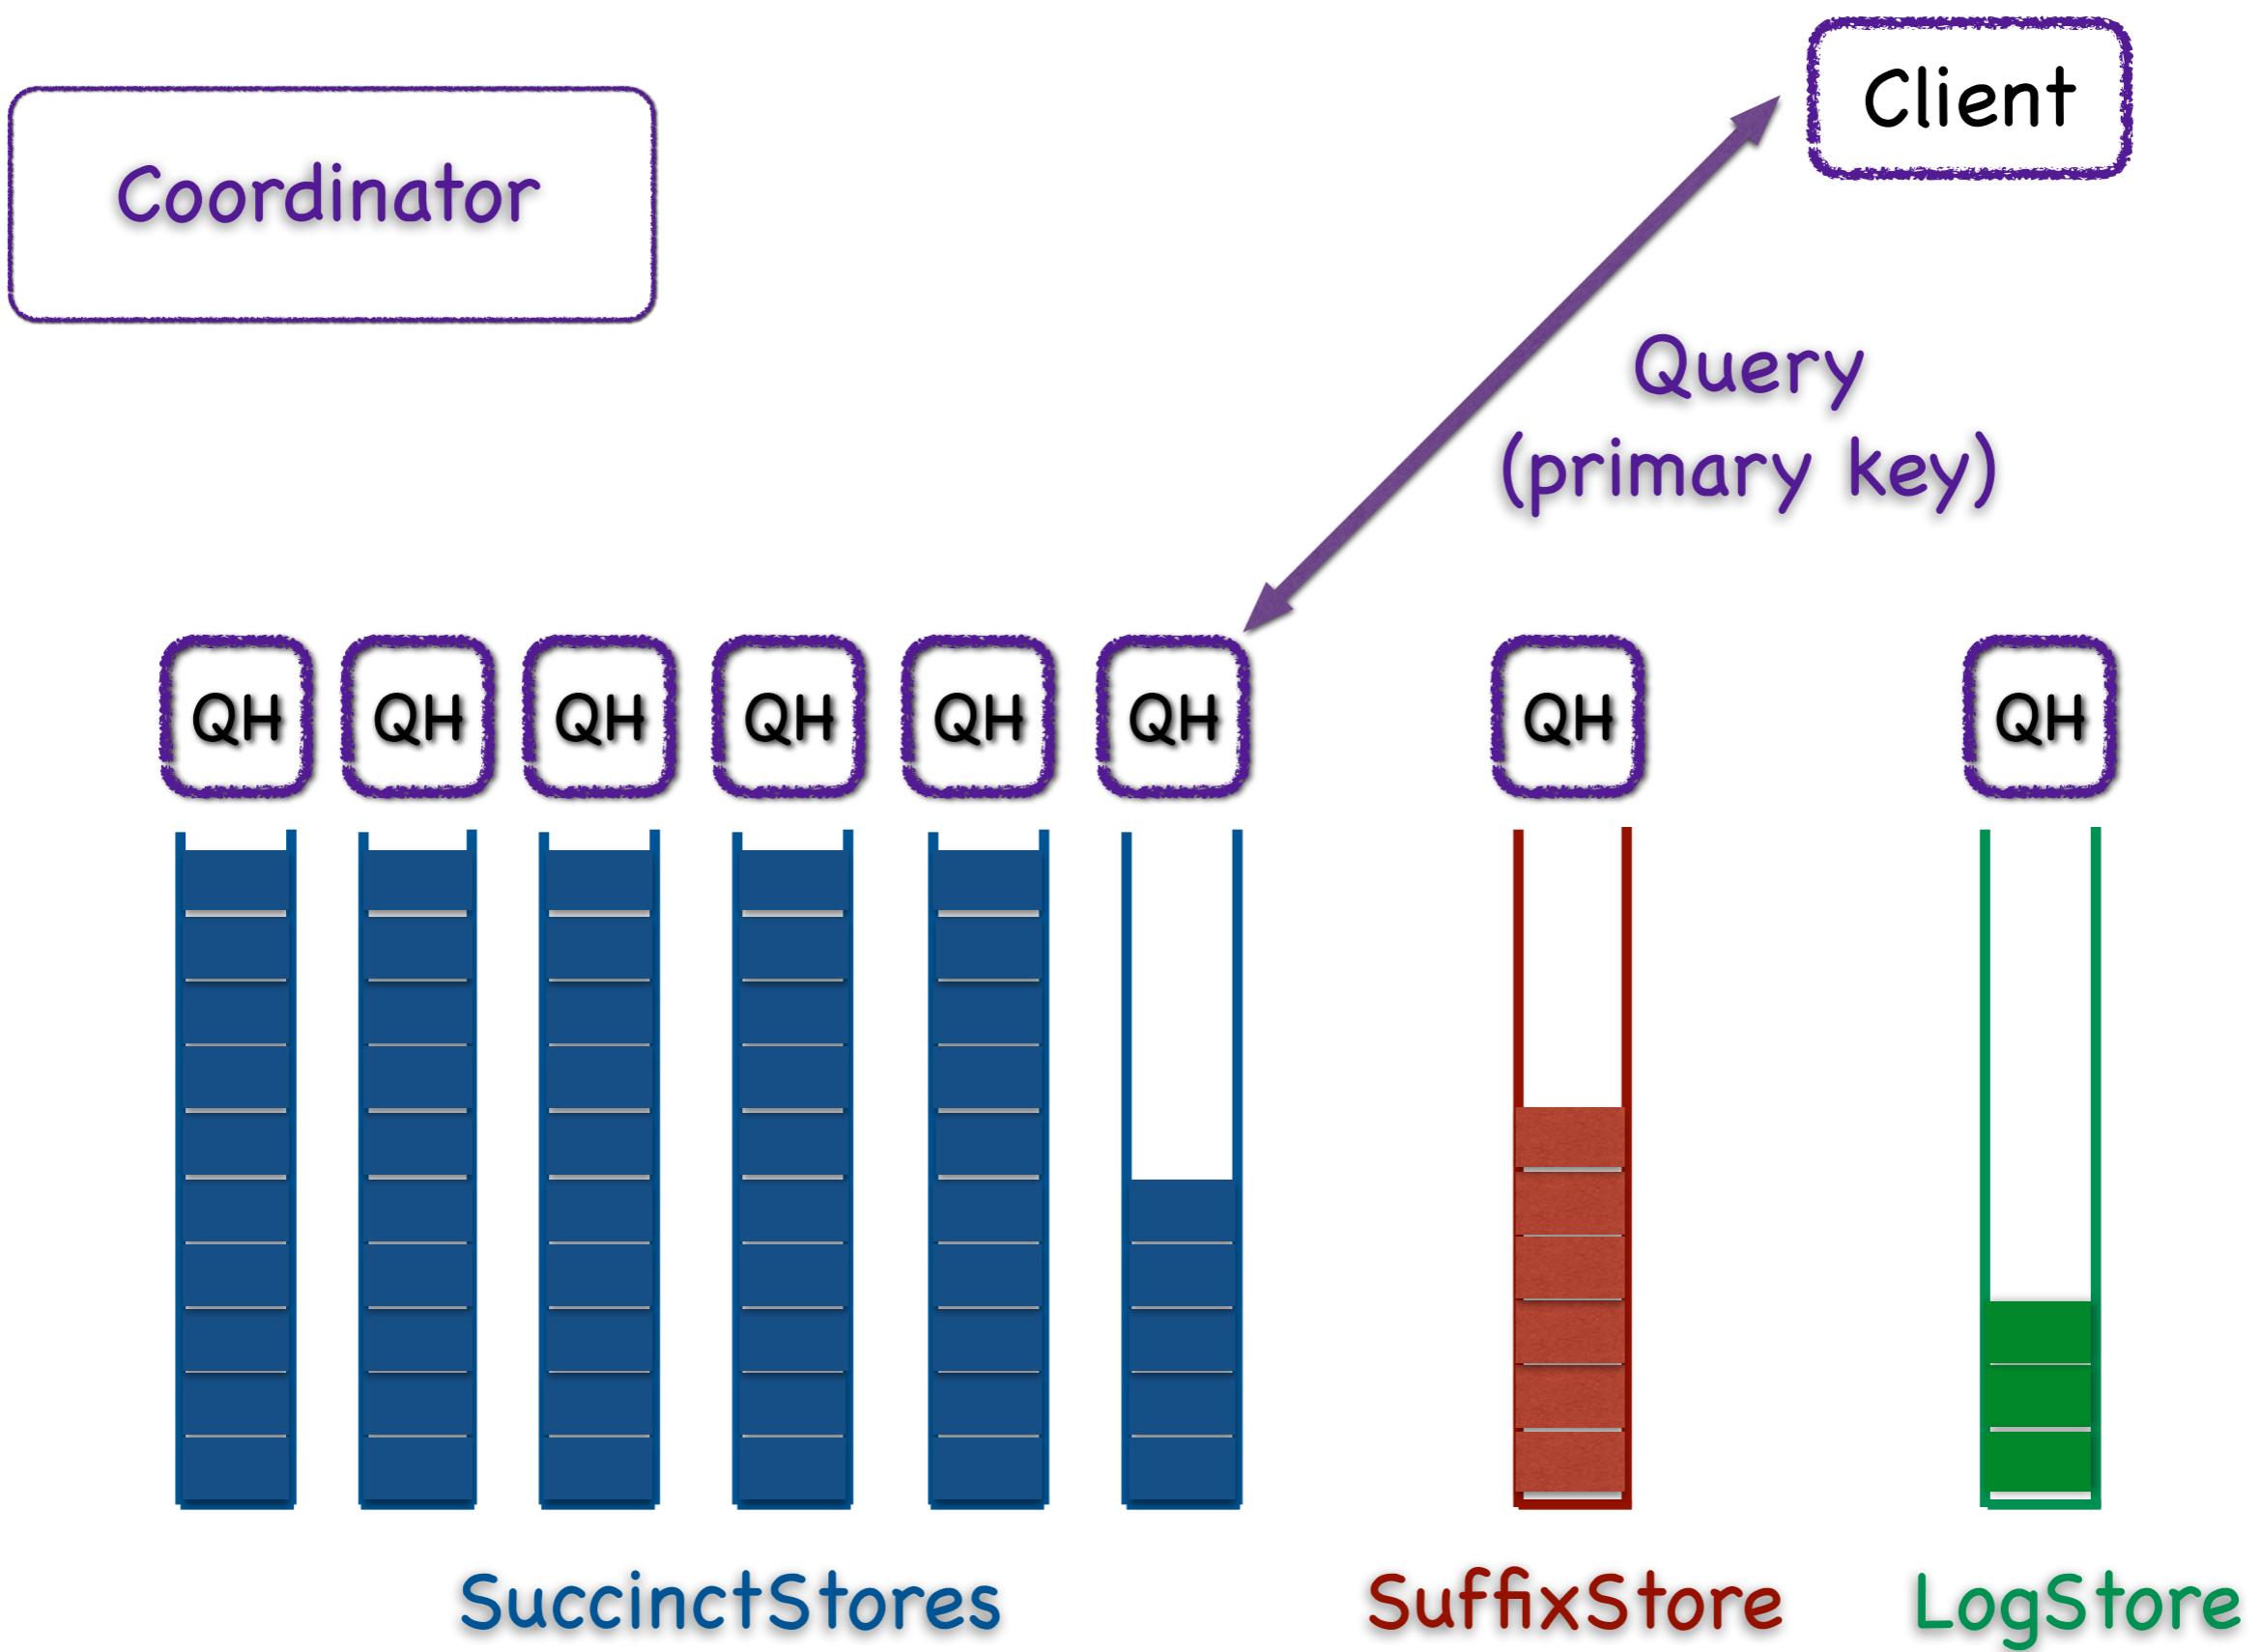


LogStore

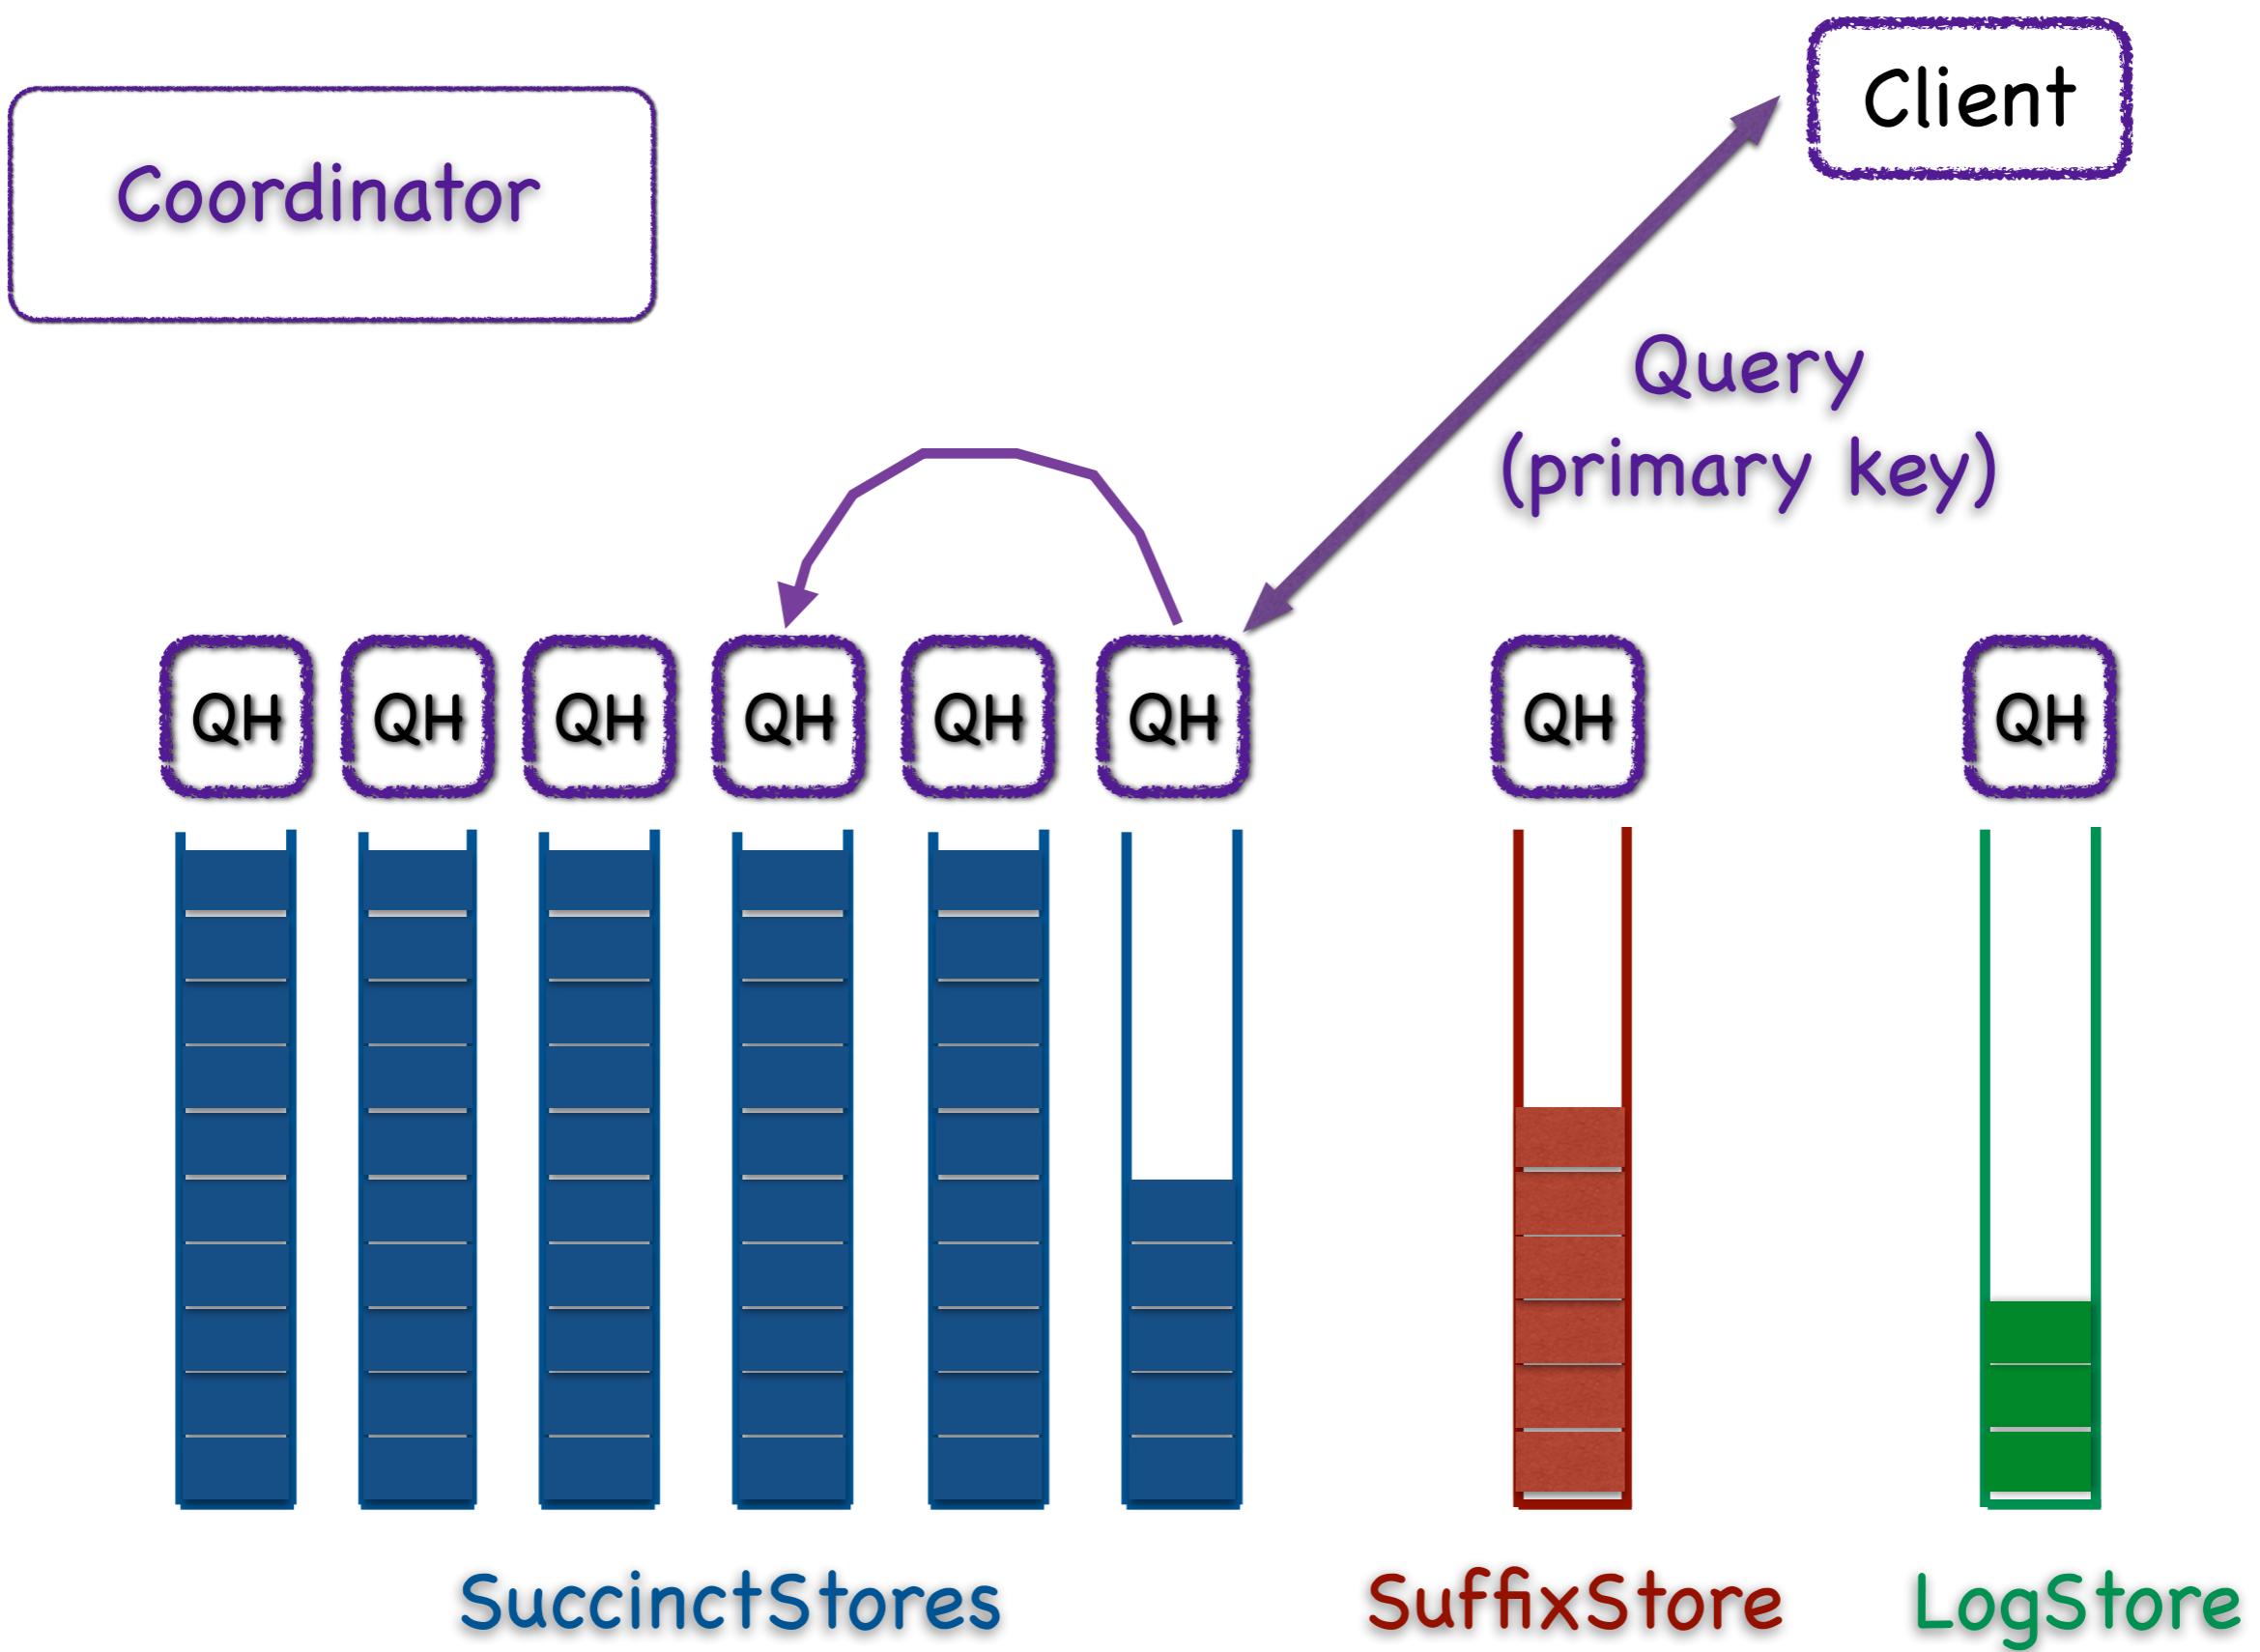
# End-to-end System



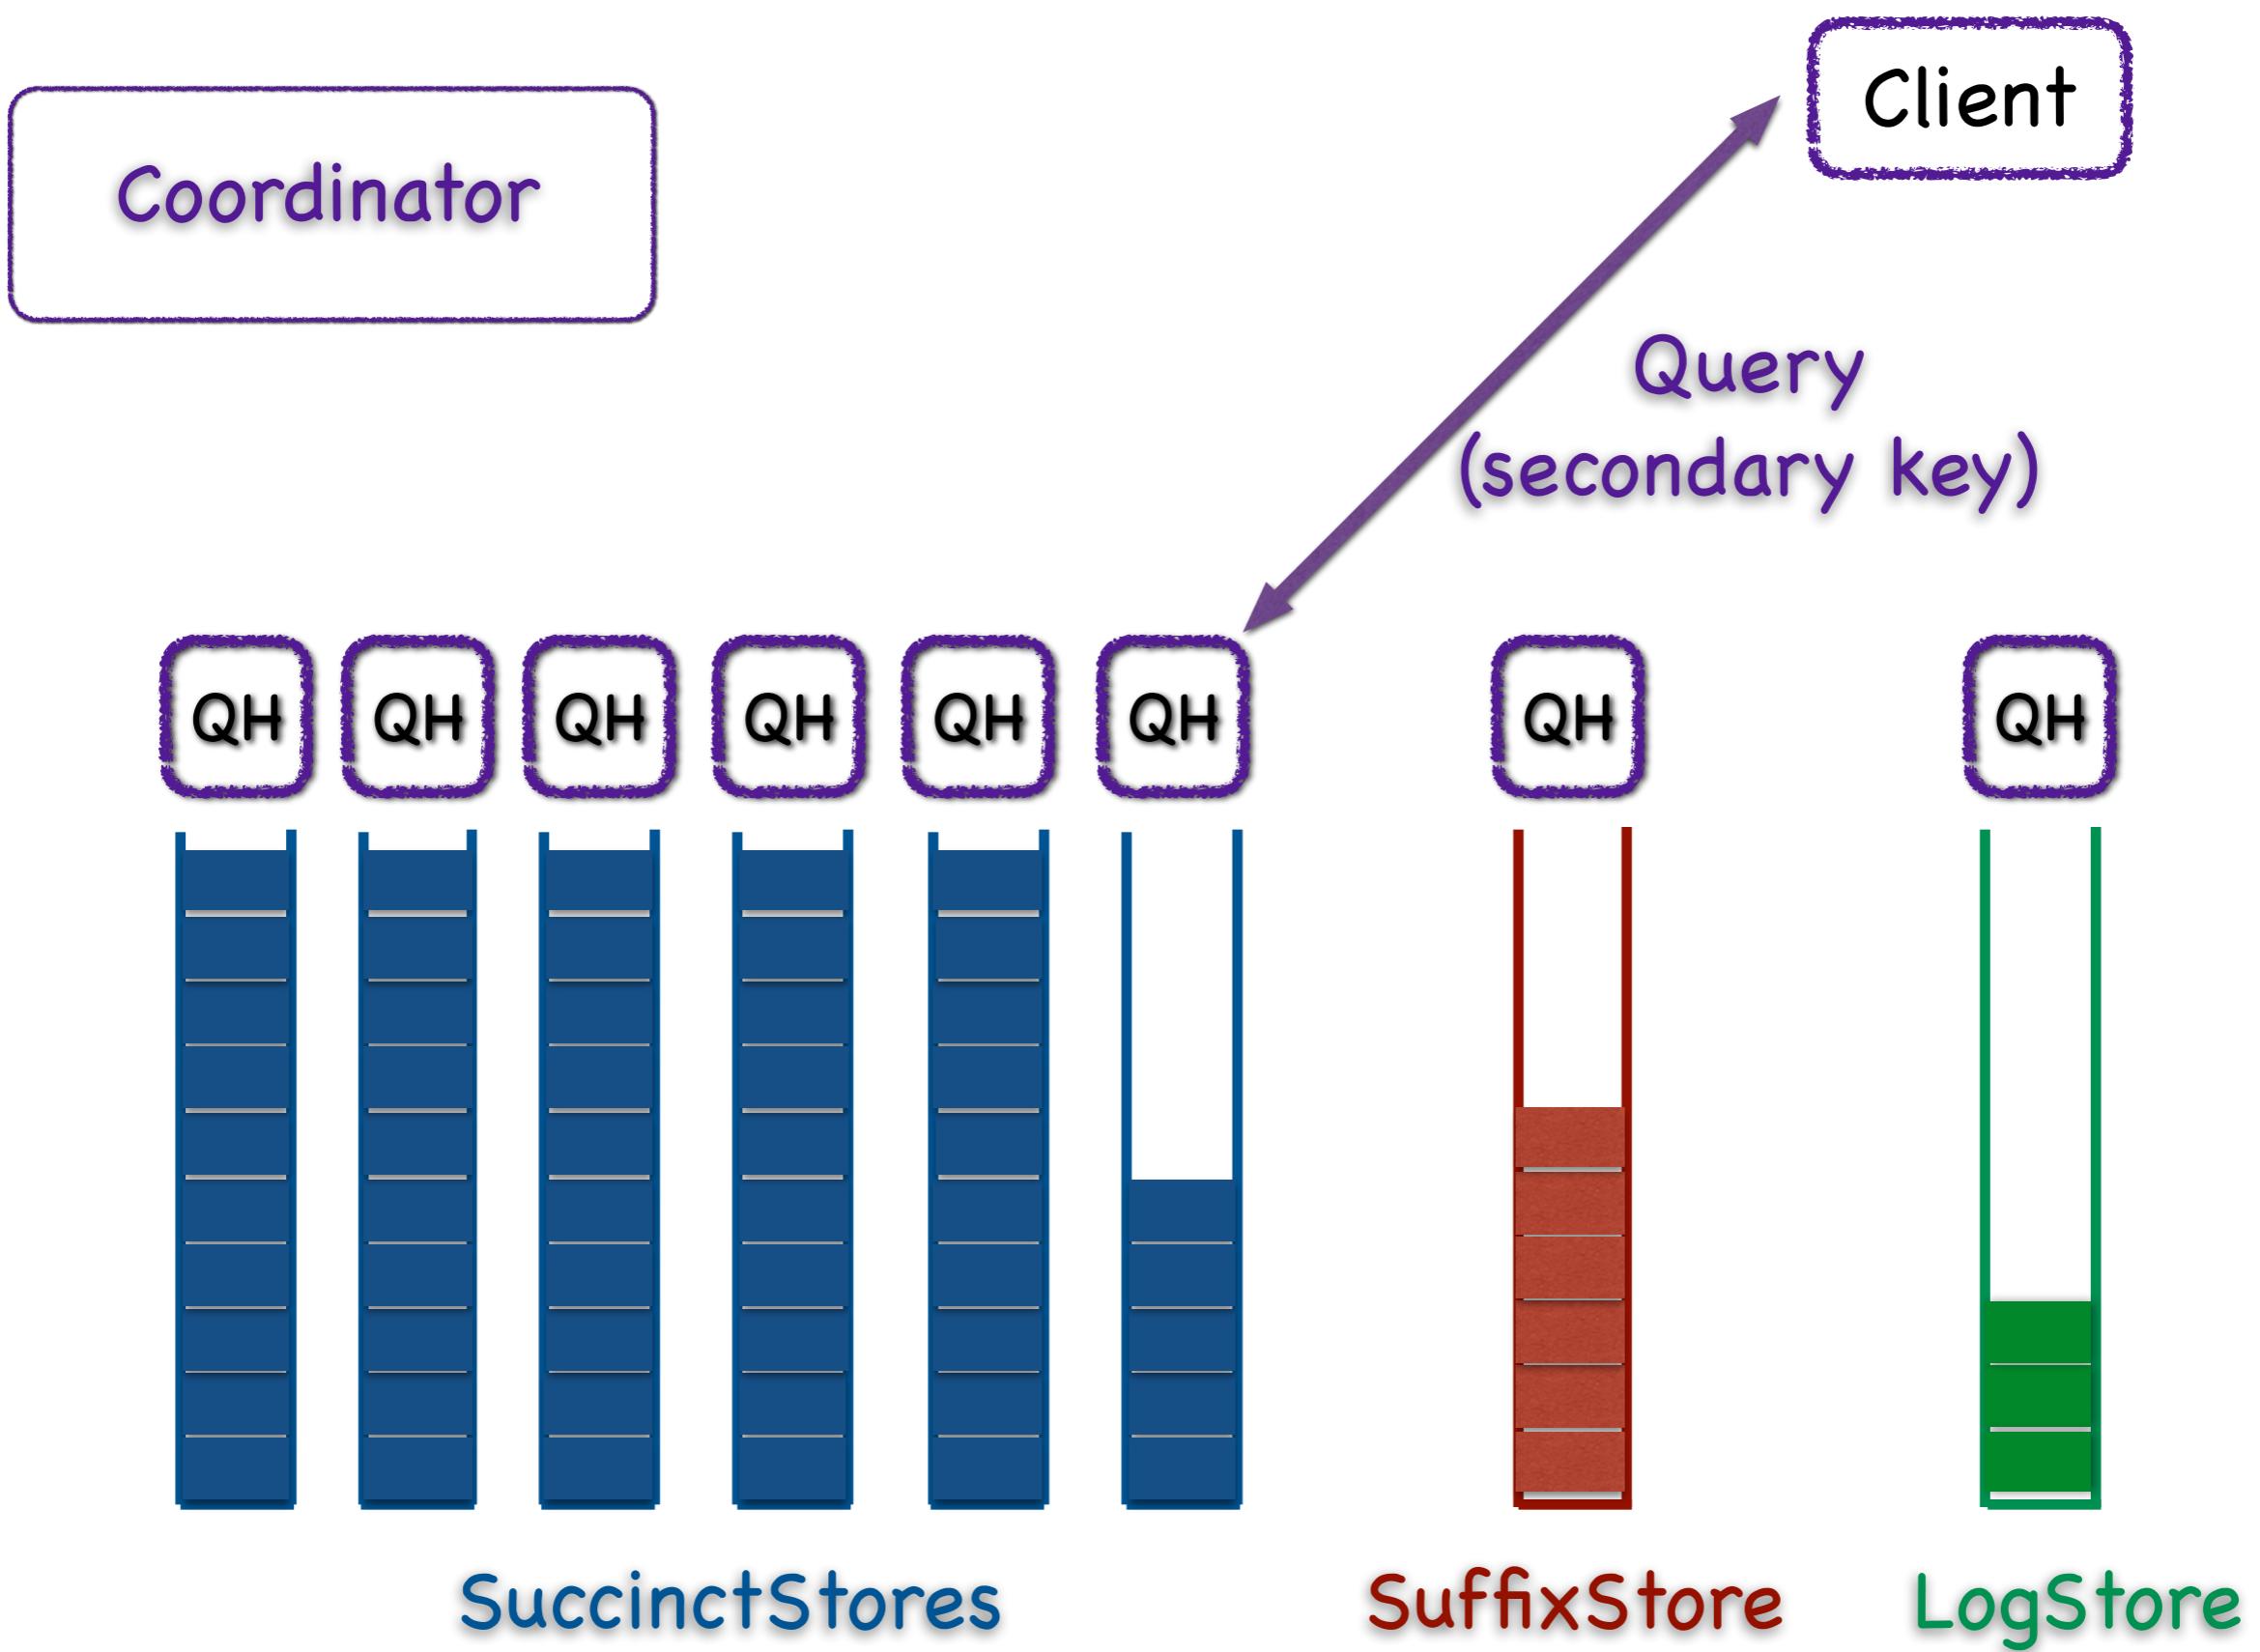
# End-to-end System



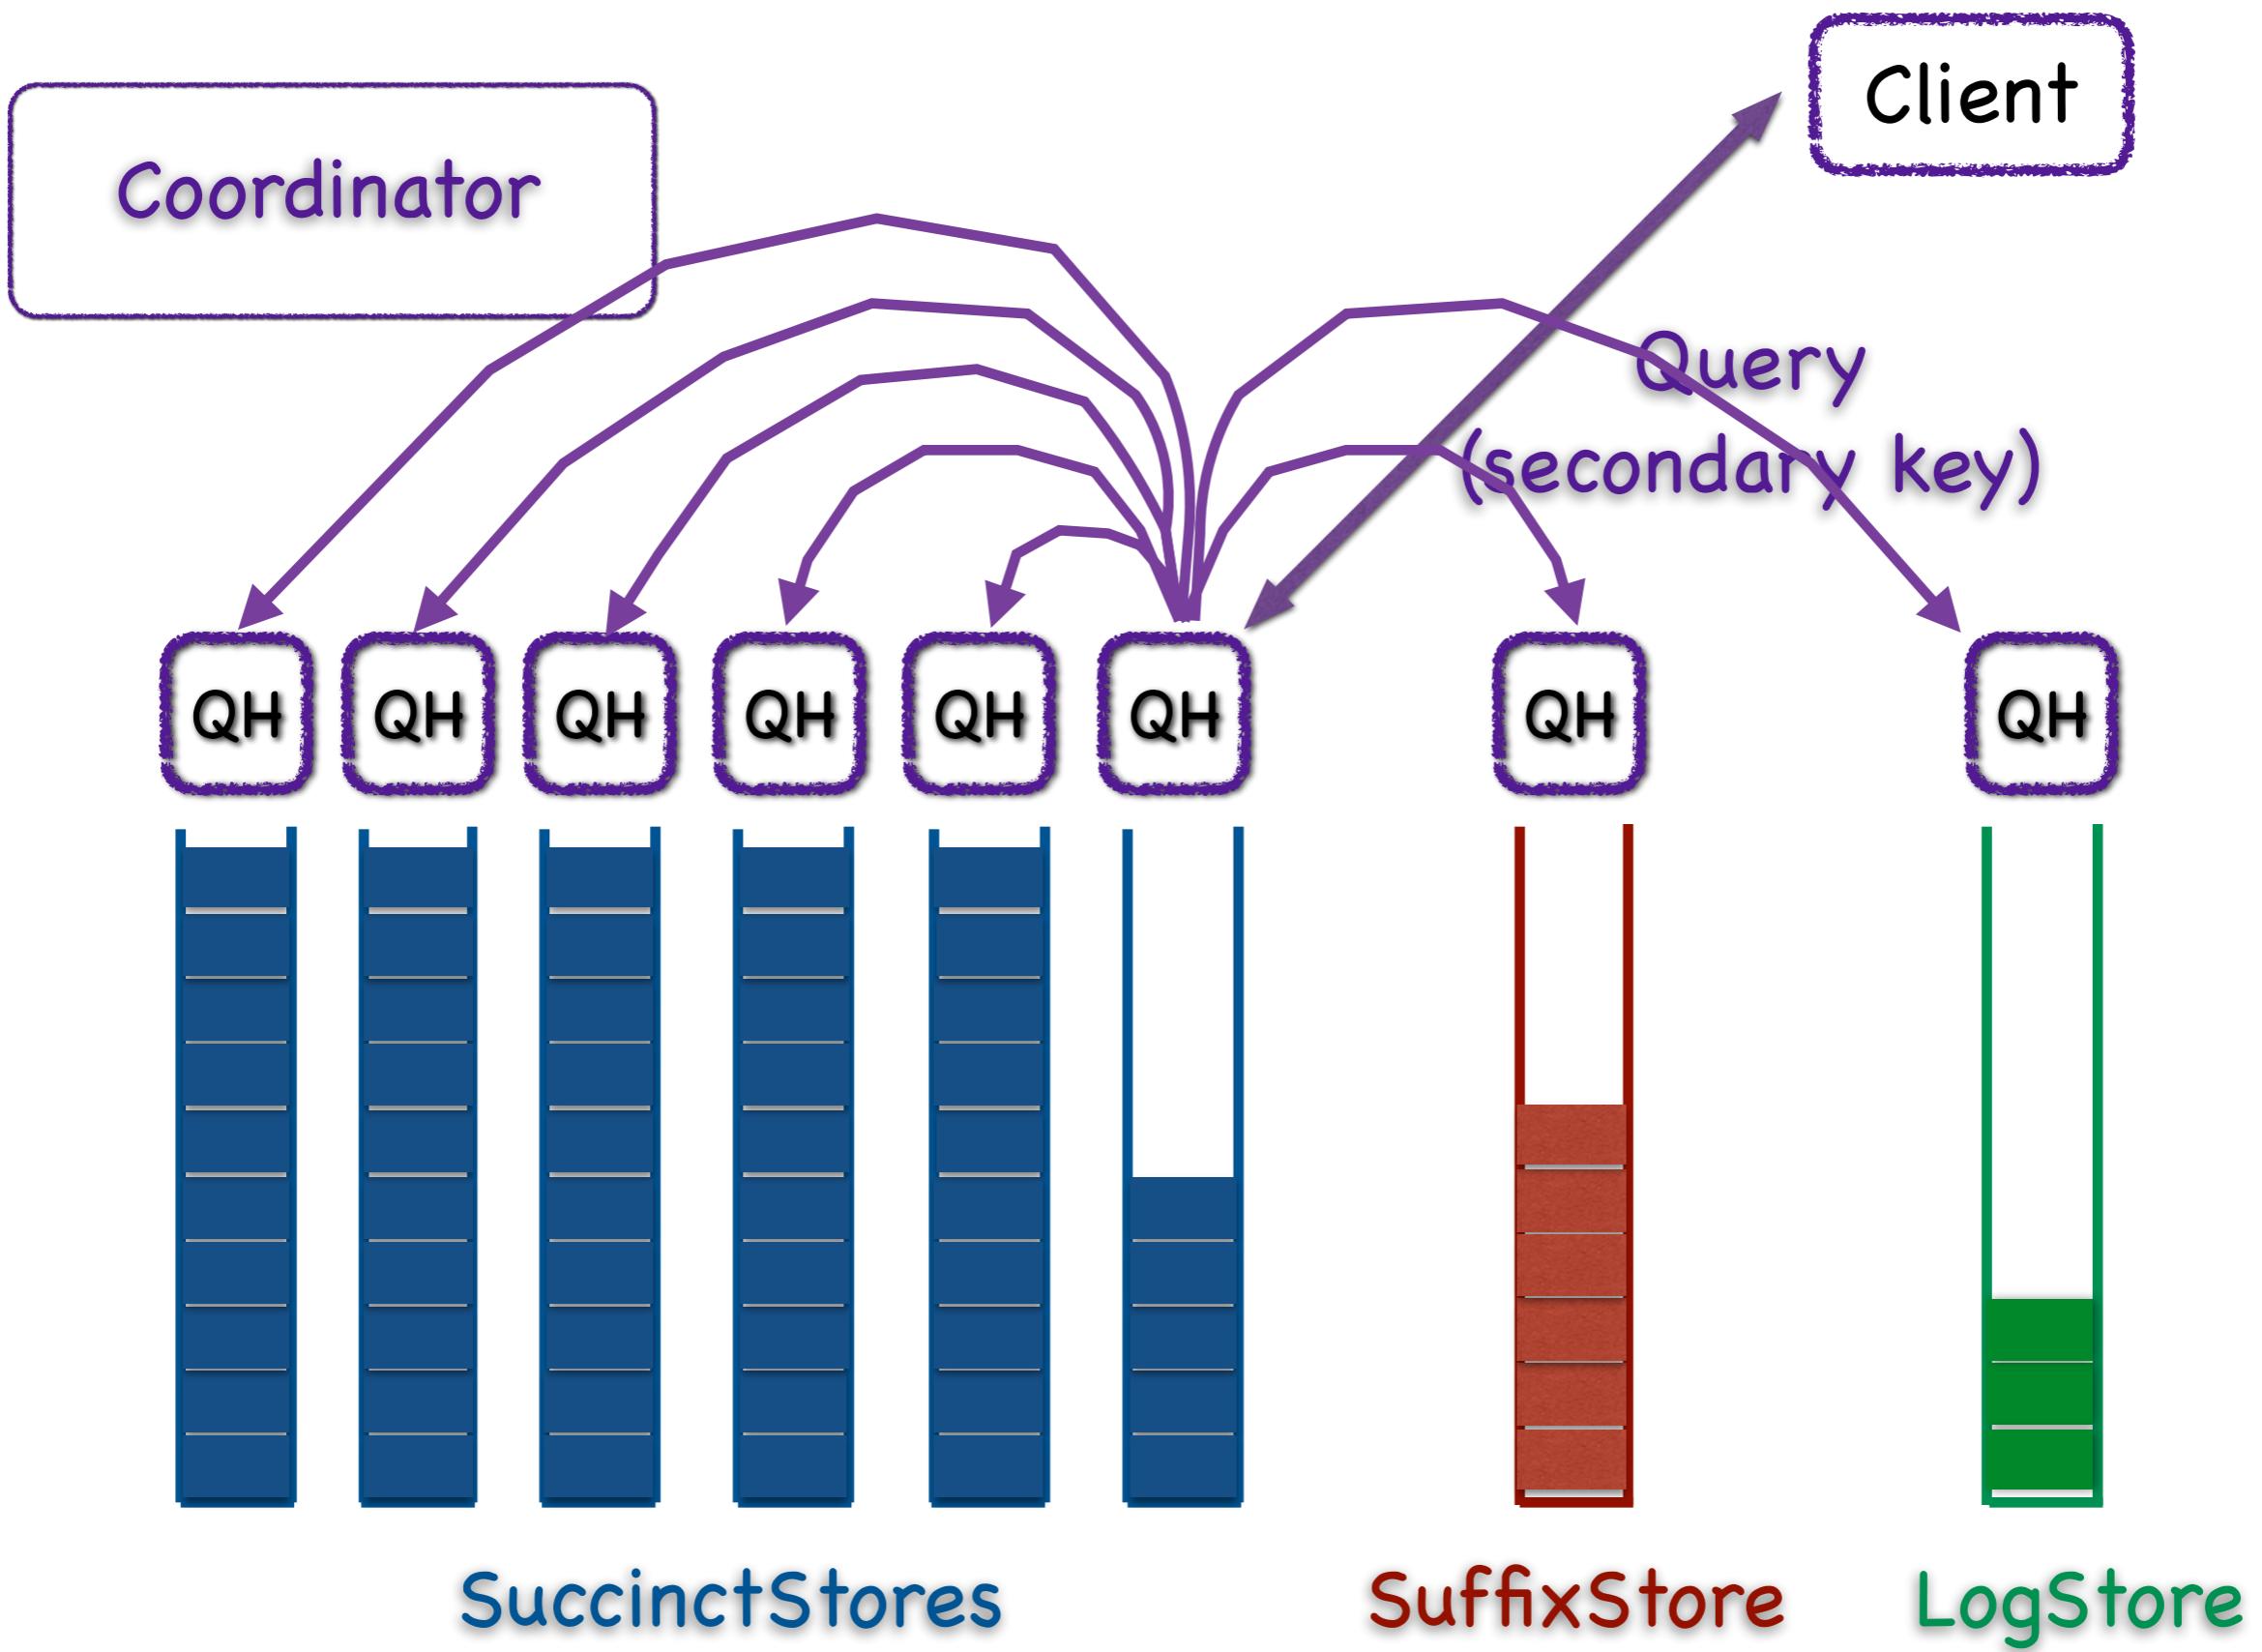
# End-to-end System



# End-to-end System



# End-to-end System



# Evaluation

---

## Datasets

SmallKV, LargeKV from Conviva customers

Record length: 140B (SmallKV), 1.3kB (LargeKV)

## Cluster

Amazon EC2, 10 machines, m1.xlarge machines

## Workload

YCSB for primary; YCSB mapped to secondary keys

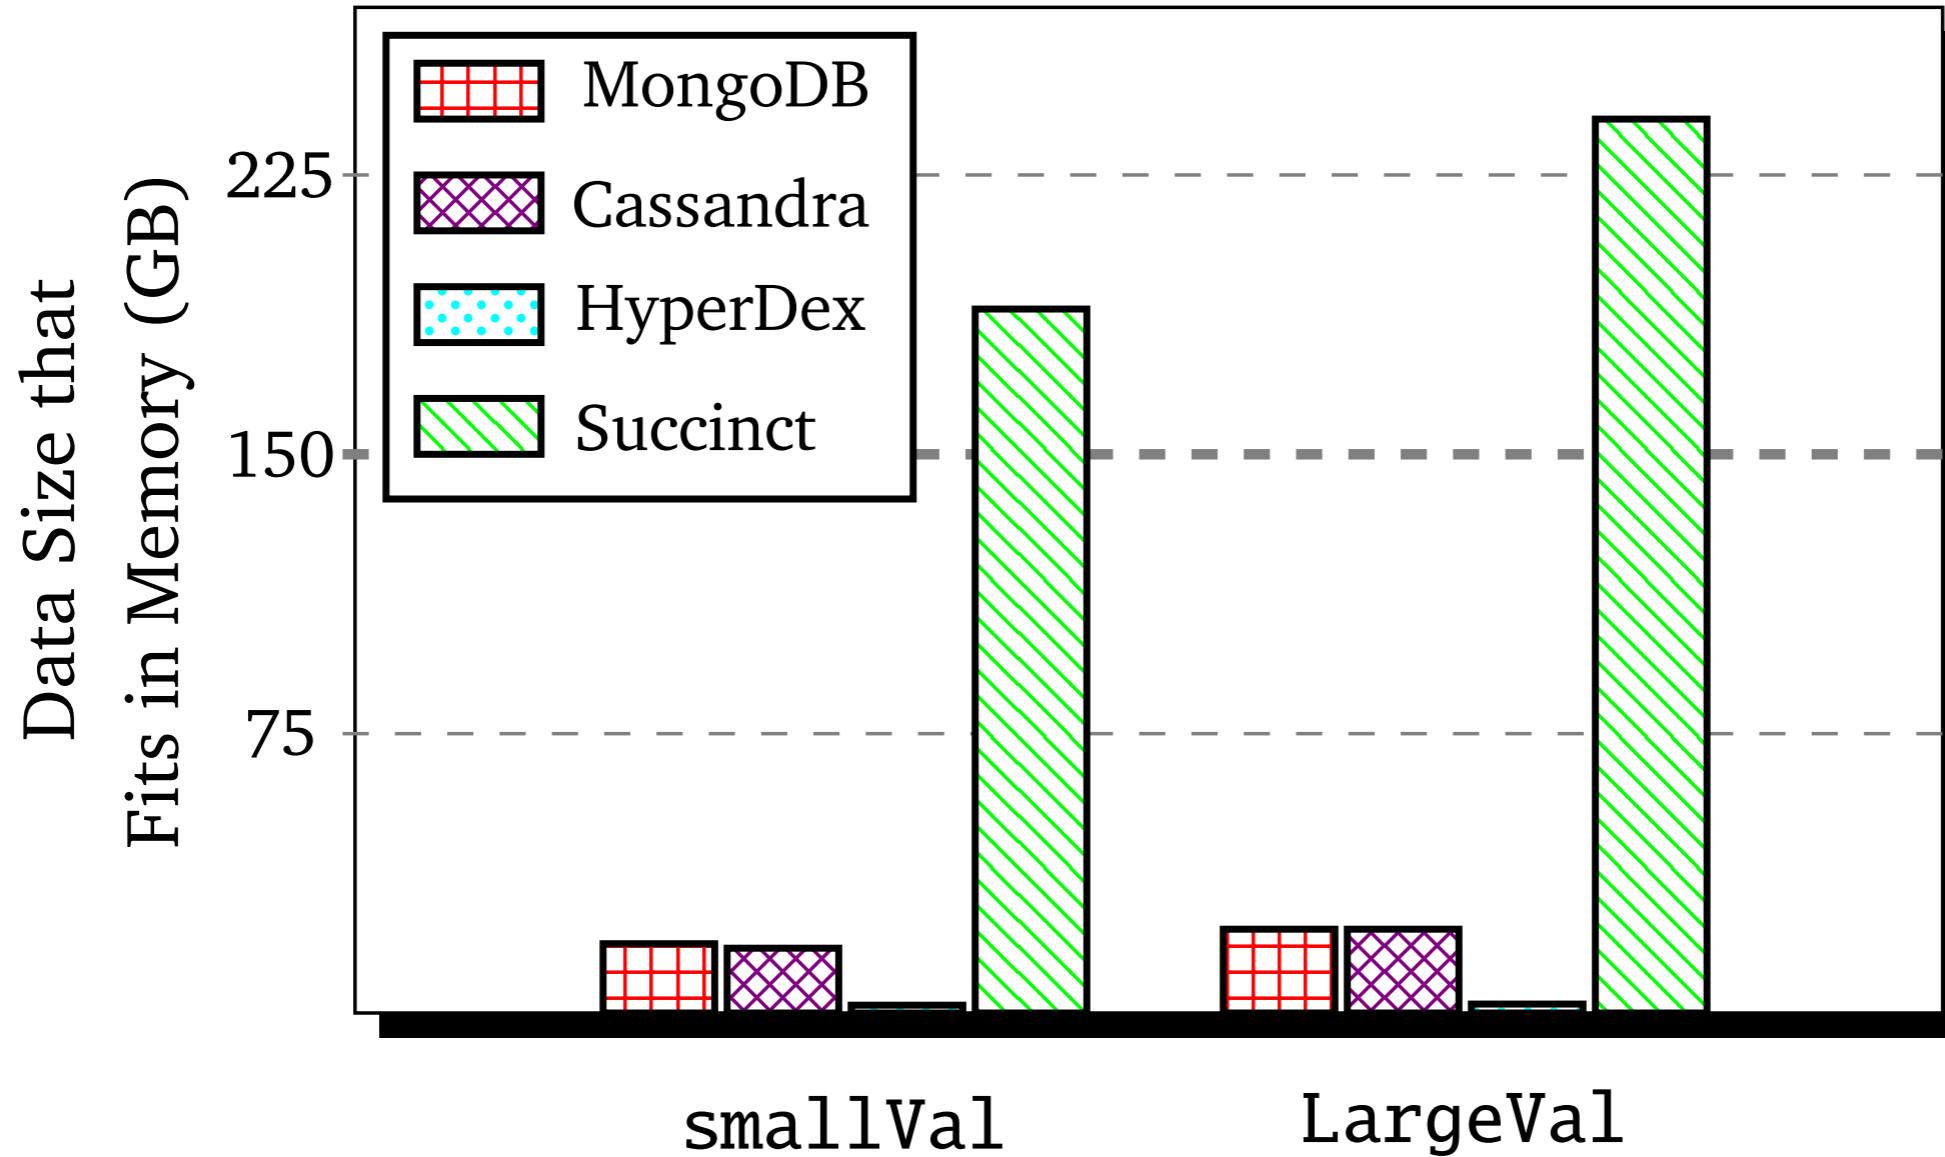
## Systems

MongoDB, Cassandra, HyperDex, DB-X

## Caveat

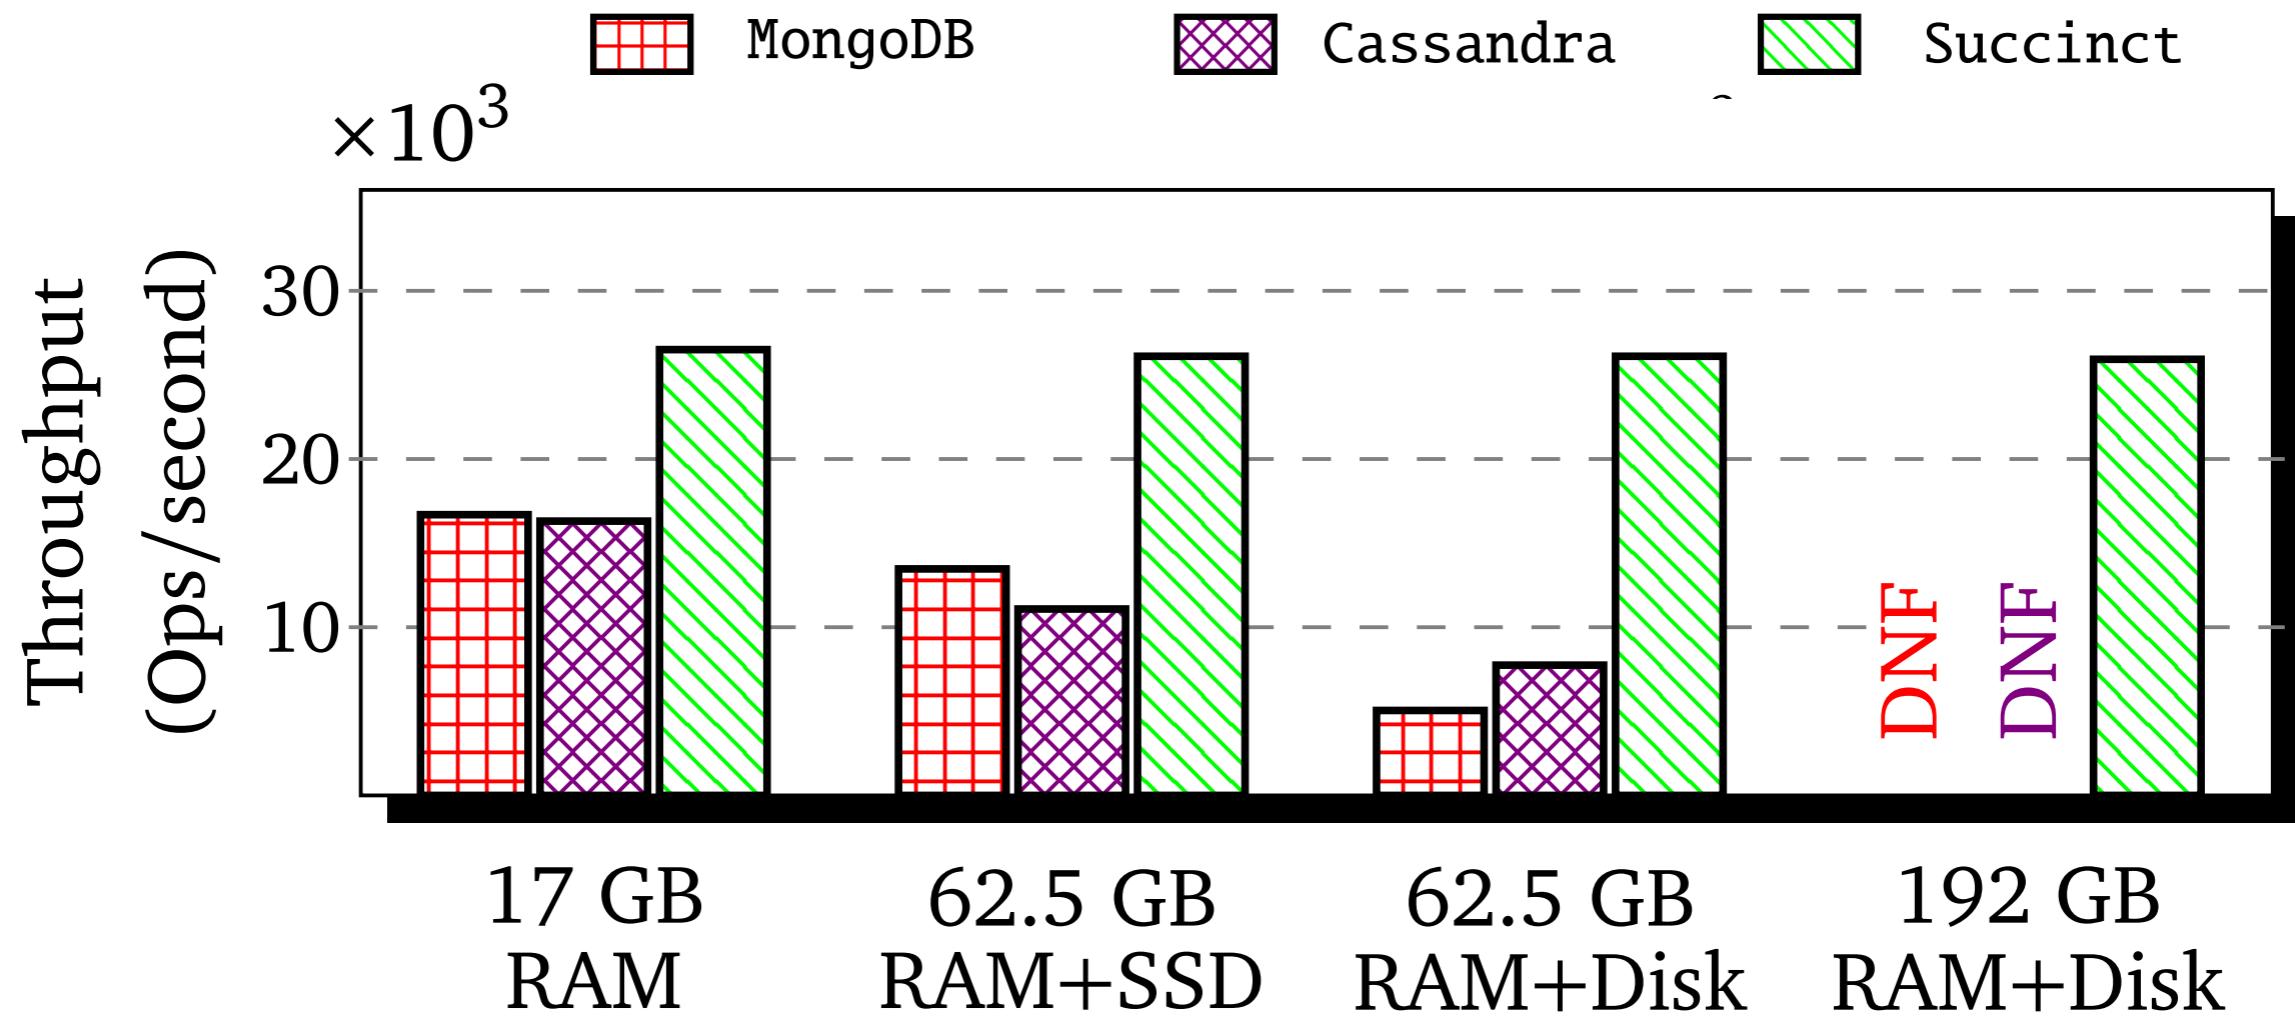
Absolute numbers are dataset dependent

# 10-11x lower storage



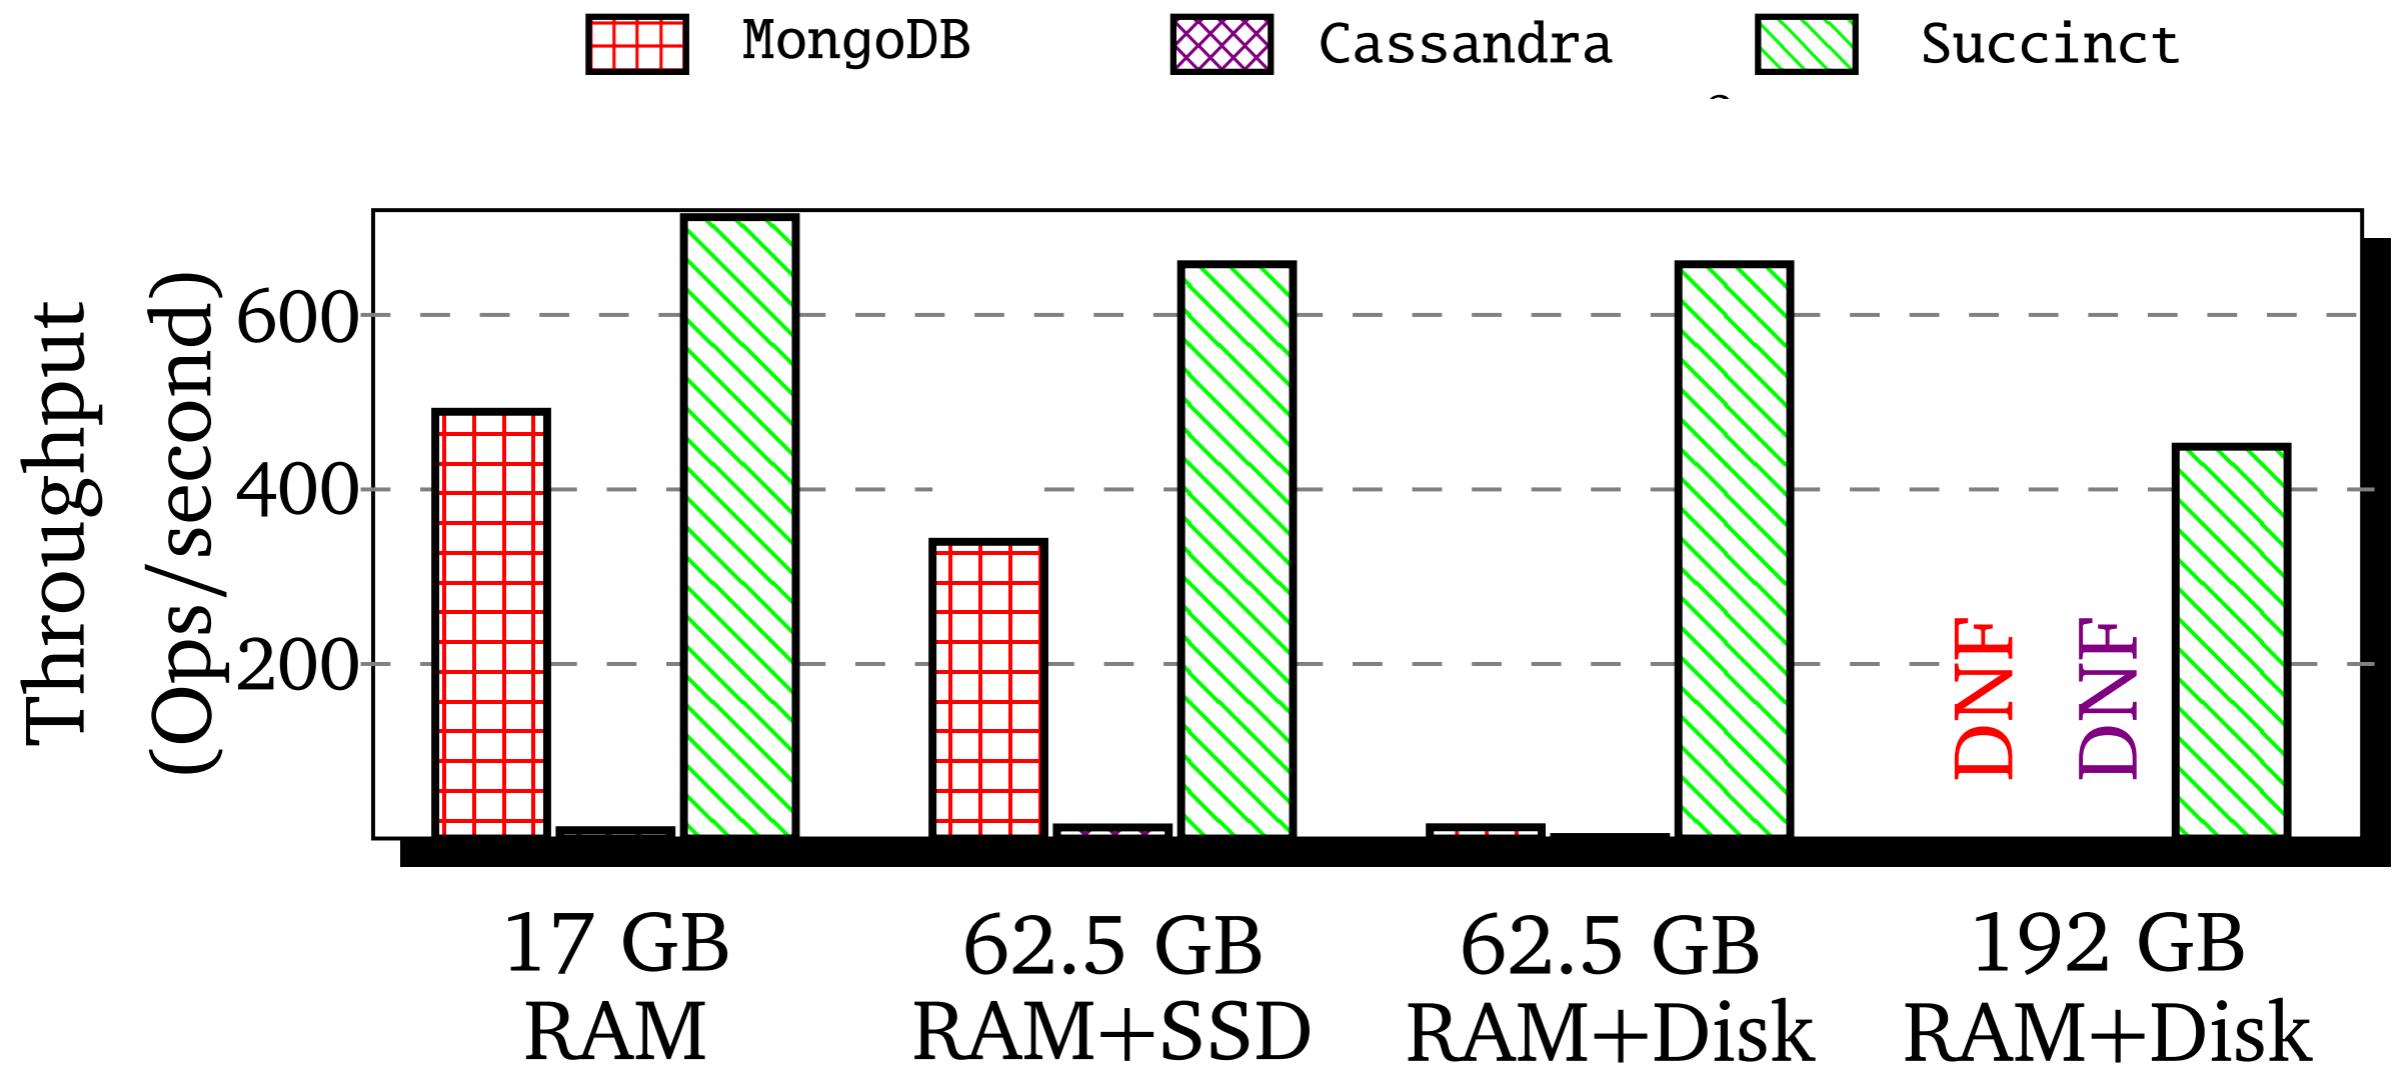
Take-away: Succinct pushes 10-11x more data in faster storage than systems with similar functionality (including every single bit stored)

# Throughput for 95% GET + 5% PUT (SmallKV)



Take-away: Succinct achieves performance comparable to existing open-source systems for queries on primary attributes

# Throughput for 95% SEARCH + 5% PUT (SmallKV)



Take-away: Succinct by pushing more data in faster storage provides performance similar to existing systems for 10-11x larger data sizes

# Summary

---

Succinct pushes 10-11x more data in faster storage  
(compared to popular open-source systems)

Succinct executes queries directly on compressed representation

Avoids data scans and data decompression costs

Enables interactive queries on secondary keys  
(for semi-structured data)