

# Note of Succinct: Enabling Queries on Compressed Data

Cheng Zhao

August 25, 2017

## 1 Motivation

Evaluation of popular open-source data stores (MongoDB, Cassandra) using real-world datasets shows that indexes can be much as 8x larger than the input data size. Existing data stores either **resort** to using complex memory management techniques for identifying and caching “hot” data or **simply executing queries off-disk or off-SSD**. In either case, latency and throughput advantages of indexes drop compared to in-memory query execution.

## 2 Succinct Overview

- a distributed data store
- operates at a new point in the design space
- memory efficiency close to data scans and latency close to indexes
- is able to store more data in memory, avoiding latency and throughput degradation due to off-disk or off-SSD query execution

Contributions:

- Enables efficient queries directly on a compressed representation of the input data
  - a new data structure
  - a new query algorithm
- Efficiently supports data appends by chaining multiple stores
  - a small log-structure store optimized for fine-grained appends
  - an intermediate store optimized for query efficiency while supporting bulk appends
  - an immutable store that stores most of the data
- Exposes a minimal, yet powerful, API that operates on flat unstructured files

## 2.1 Succinct Interface

For string `abcdeabczabgz`,

- `search(f, "ab")` returns `[0, 6, 10]`.
- `wildcardsearch(f, prefix="ab", suffix="z", dist=2)` returns `[6, 9]` for `abcz`, `[10, 13]` for `abgz`.

For semi-structured data (KV), like figure 1, transforming the input data into flat files.




Figure 1: Succinct supports queries on semi-structured data by transforming the input data into flat files

## 3 Querying on Compressed Data

### 3.1 Compressed Suffix array and Succinct data representation

Suffix array Wikipedia




Figure 2

Figure 2 (a) shows that AoS stores the first char of each suffix only. Since suffixes are sorted, only the first AoS index at which each character occurs (e.g.,  $\{(\$, 0), (a, 1), (b, 4), (n, 5)\}$ ) need be stored. Succinct uses *Skewed Wavelet Tree* to compress each row independently.

Succinct uses sampling by “value” strategy. For sampling rate  $\alpha$ , Succinct stores all AoS2Input values that are a multiple of  $\alpha$ . This allows storing each sampled value `val` as `val/α`, leading to a more space-efficient representation. Using  $\alpha = 2$  for example of Figure 3a, for instance, the sampled

AoS2Input values are  $\{6, 0, 4, 2\}$ , which can be stored as  $\{3, 0, 2, 1\}$ . Sampled Input2AoS then becomes  $\{1, 3, 2, 0\}$  with  $i$ -th value being the index into sampled AoS2Input where  $i$  is stored. Succinct stores a small amount of additional information to locate sampled AoS2Input indexes.




Figure 3: **3a** for reducing the space usage of AoS2Input. **3b** for two-dimensional NextCharIdx representation.

For instance, consider the query `search(anan)`; all occurrences of string "nan" are contained in the cell  $\langle n, an \rangle$ . To find all occurrences of string `anan`, our algorithm performs a binary search only in the cell  $\langle a, na \rangle$  in the next step. After this step, the algorithm has the indexes for which suffixes start with "a" and are followed by "nan", the desired string. For a string of length  $m$ , the above algorithm performs  $2(mt1)$  binary searches, two per NextCharIdx cell.




Figure 4: **4a** for The Input2AoS provides the inverse mapping of AoS2Input. **4b** for reducing the space usage of Input2AoS.

`extract` functionality is shown in Figure 4. For instance, to execute `extract(3, 3)`, we find the next smaller sampled index (`Input2AoS[2]`) and corresponding suffix (`AoS[2]="nana$"`). We then remove the first character since the difference between the desired index and the closest sampled index was 1; hence the result "ana\$".