

# ZipG

## A Memory-efficient Graph Store for Interactive Queries

---

Cheng Zhao

September 6, 2017

# Table of contents

Motivation

Preliminary: Succinct

ZipG Design

Evaluation

Reference

# Motivation

---

- 分布式图计算系统往往较难处理 “*Find friends of Alice who live in Ithaca*” 这类查询，因为典型的图查询几乎没有“**局部性**”（它可能会访问图的任意部分；查找所有包括 *Alice's friends* 的块）
- 传统的块压缩技术 (e.g., gzip) 也因为缺少“局部性”而在处理图查询时效率很低
- 一些系统在压缩后的图上进行查询，但他们往往忽略点和边的属性，并且只能进行有限的操作 (e.g., 查找与点相连的边，子图匹配) [1, 2, 3, 4, 5, 6, 7]

## Preliminary: Succinct

---

# ZipG Design

---

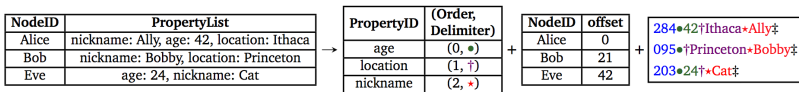
# Graph Representation

## • NodeFile

- 记录每个属性的字典序和使用的分隔符;
- 记录每个 NodeID 对于的offset;
- 在每条 PropertyList 的开始记录每列属性占用的空间。

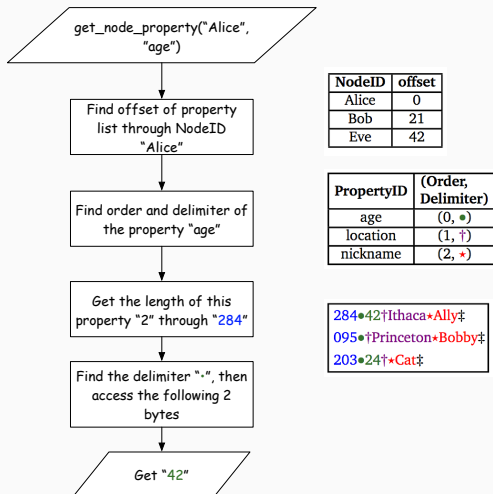
主要处理两类请求:

1. 给定(NodeID, List<propertyID>) 找出该节点的对应的属性值;
2. 给定一个属性列表, 找到所有与之匹配的节点。



# Query Execution

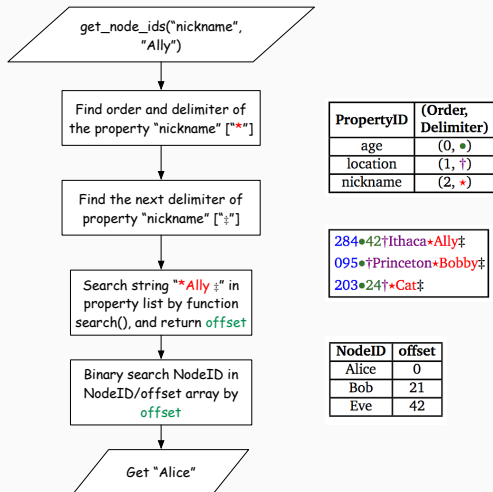
- `get_node_property(NodeID, propertyID)`





# Query Execution

- `get_node_ids(propertyID, propertyValue)`



# Graph Representation

- EdgeFile

- 一条 EdgeRecord 记录一组同一源点相同类型的边
- 时间戳存储需要支持二分查找，存储的两个极端
  1. **Space-efficient** 用最少的空间表示每个时间戳，但是需要用分隔符或者存储长度来标识每个时间戳；或者使用 *delta encoding* 存储
  2. **Latency-efficient** 固定一个全局时间戳长度，每条 EdgeRecord 存储相同长度的时间戳

• EdgeCount \* Tlength + Dlength

```
$S_1\#EdgeType_1,Metadata,T_0,...,T_M,D_0,...,D_M,PropertyList_0,...,PropertyList_M  
$S_1\#EdgeType_2,Metadata,T_0,...,T_M,D_0,...,D_M,PropertyList_0,...,PropertyList_M  
⋮  
$S_f\#EdgeType_k,Metadata,T_0,...,T_M,D_0,...,D_M,PropertyList_0,...,PropertyList_M
```

# Graph Representation

- EdgeFile

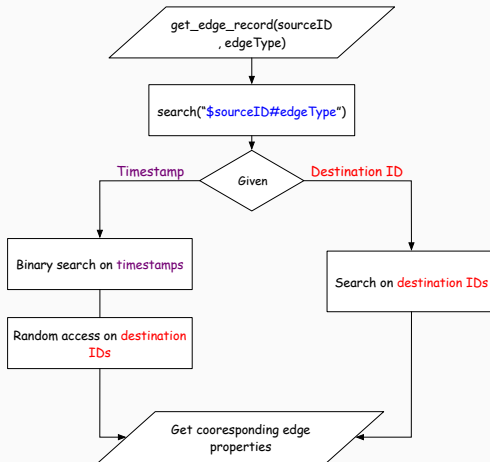
- 一条 EdgeRecord 记录一组同一源点相同类型的边
- Metadata 存储边的个数，时间戳的长度，目的节点的长度
- 以此条 EdgeRecord 内需要存储的时间戳最大长度作为每个时间戳的长度（每条 EdgeRecord 的时间戳长度可能不同），并在Metadata中记录
- 以相同方法存储对应的目的节点，并在Metadata中记录其长度
- 边属性的存储方法类似于节点属性，但是目前ZipG不支持在此上做查询

•EdgeCount★Tlength†DLength

$S_1 \# \text{EdgeType}_1, \text{Metadata}, T_0, \dots, T_M, D_0, \dots, D_M, \text{PropertyList}_0, \dots, \text{PropertyList}_M$   
 $S_1 \# \text{EdgeType}_2, \text{Metadata}, T_0, \dots, T_M, D_0, \dots, D_M, \text{PropertyList}_0, \dots, \text{PropertyList}_M$   
 $\vdots$   
 $S_f \# \text{EdgeType}_k, \text{Metadata}, T_0, \dots, T_M, D_0, \dots, D_M, \text{PropertyList}_0, \dots, \text{PropertyList}_M$

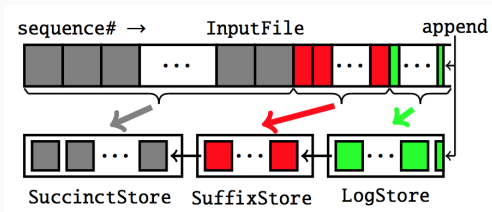
# Query Execution

- `get_edge_record(sourceID, edgeType)`



根据CPU的核心数划分数据（1 shard per core），根据 NodeID 将节点哈希到各个 shard 上，并存储所有与节点的出边信息。

# Succinct Store



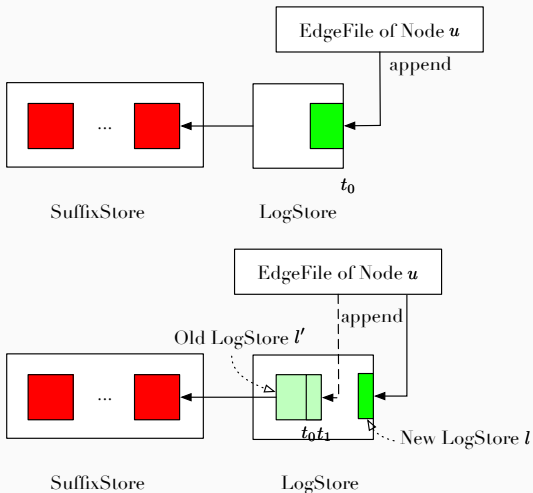
|             | Succinct Store      | Suffix Store      | Log Store         |
|-------------|---------------------|-------------------|-------------------|
| Stores      | Comp. Data (§3.1)   | Data + AoS2Input  | Data + Inv. Index |
| Appends     | -                   | Bulk              | Fine              |
| Queries     | §3.2                | Index             | Scans+ Inv. Index |
| #Machines   | $n-2$               | 1                 | 1                 |
| %Data(est.) | > 99.98%            | < 0.016%          | < 0.001%          |
| Memory      | $\approx 0.4\times$ | $\approx 5\times$ | $\approx 9\times$ |

ZipG 在整个系统中只用一个 LogStore (query-optimized)，一旦其大小超过阈值即压缩为 memory-optimized 存储方式并创建一个新的 LogStore. 其优势在于：

1. 不需要对数据进行解压和再压缩；
2. 单 LogStore 比之前的方法（每个Worker上一个LogStore）更省内存；
3. 避免了多 LogStore 时同步读写的麻烦。

# Fanned Updates

- 对于非结构化和半结构化的数据，单LogStore就比多机的LogStore更好；
- 对于图结构数据，使用单LogStore会导致 fragmented storage





# Fanned Updates

Fanned Updates 在节点/边第一次出现的shard上设置 update pointers, 从而避免了访问所有的 shard 。

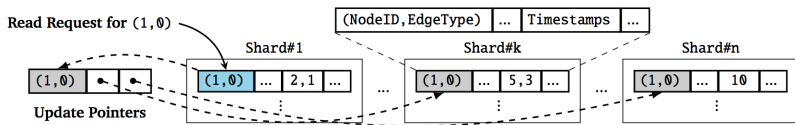


Figure 3: Update Pointers for the EdgeFile (§3.5).

# Evaluation

---

- `assoc_range(id, atype, idx, limit)`  
从位置 `idx` 开始, 获取不超过 `limit` 条源点为 `id`, 类型为 `atype` 的边
- `assoc_get(id1, atype, id2set, hi, lo)`  
获取所有的源点为 `id1`, 类型为 `atype`, 时间戳在 `[hi, lo)` 内, 汇点在 `id2set` 中的边
- `assoc_time_range(id, atype, hi, lo, limit)`  
获取不超过源点为 `id1`, 类型为 `atype`, 时间戳在 `[hi, lo)` 内的边

**Table 2:** Queries in TAO [29] and LinkBench [24] workloads.

| Query            | Execution in ZipG | TAO %  | LinkBench % |
|------------------|-------------------|--------|-------------|
| assoc_range      | Algorithm 1       | 40.8   | 50.6        |
| obj_get          | get_node_property | 28.8   | 12.9        |
| assoc_get        | Algorithm 2       | 15.7   | 0.52        |
| assoc_count      | get_edge_record   | 11.7   | 4.9         |
| assoc_time_range | Algorithm 3       | 2.8    | 0.15        |
| assoc_add        | append            | 0.1    | 9.0         |
| obj_update       | delete, append    | 0.04   | 7.4         |
| obj_add          | append            | 0.03   | 2.6         |
| assoc_del        | delete            | 0.02   | 3.0         |
| obj_del          | delete            | < 0.01 | 1.0         |
| assoc_update     | delete, append    | < 0.01 | 8.0         |

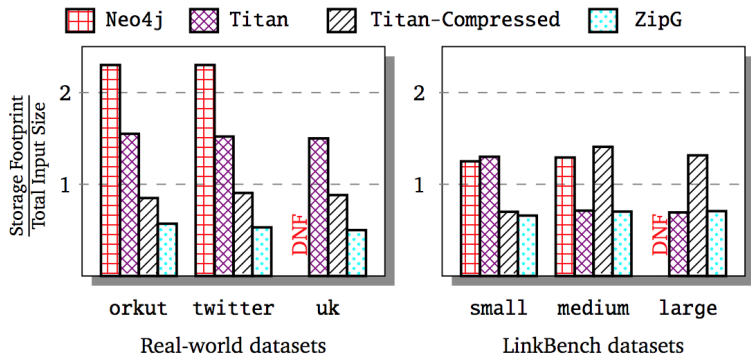
**Table 3:** The Graph Search Workload and implementation using ZipG API; p1 and p2 are node properties, id and type are NodeID and EdgeType. All queries occur in equal proportion in the workload.

| QID | Example                     | Execution in ZipG           |
|-----|-----------------------------|-----------------------------|
| GS1 | All friends of Alice        | get_neighbor_ids(id,*,*)    |
| GS2 | Alice's friends in Ithaca   | get_neighbor_ids(id,*,{p1}) |
| GS3 | Musicians in Ithaca         | get_node_ids({p1,p2})       |
| GS4 | Close friends of Alice      | get_neighbor_ids(id,type,*) |
| GS5 | All data on Alice's friends | assoc_range(id,type,0,*)    |

**Table 4:** Datasets used in our evaluation.

|            | Dataset      | #nodes & #edges    | Type   | On-disk Size |
|------------|--------------|--------------------|--------|--------------|
| Real-world | orkut [41]   | ~ 3M & ~ 117M      | social | 20 GB        |
|            | twitter [28] | ~ 41M & ~ 1.5B     | social | 250 GB       |
|            | uk [28]      | ~ 105M & ~ 3.7B    | web    | 636 GB       |
| LinkBench  | small        | ~ 32.3M & ~ 141.7M | social | 20 GB        |
|            | medium       | ~ 403.6M & ~ 1.76B | social | 250 GB       |
|            | large        | ~ 1.02B & ~ 4.48B  | social | 636 GB       |

- For real-world datasets, Each node has an average propertyList of 640 bytes distributed across 40 propertyIDs. Each edge is randomly assigned one of 5 distinct EdgeTypes, a POSIX timestamp drawn from a span of 50 days, and a 128- byte long edge property.
- LinkBench models Facebook's database workload for social graph queries. For LinkBench datasets, these datasets mimic the Orkut, Twitter and UK graphs in terms of their total on-disk size. LinkBench assigns a single property to each node and edge in the graph, with the properties having a median size of 128 bytes.



**Figure 5: ZipG's storage footprint (§5.1) is 1.8–4× lower than Neo4j and 1.8–2× lower than Titan. DNF denotes that the experiment did not finish after 48 hours of data loading.**

# Single Server

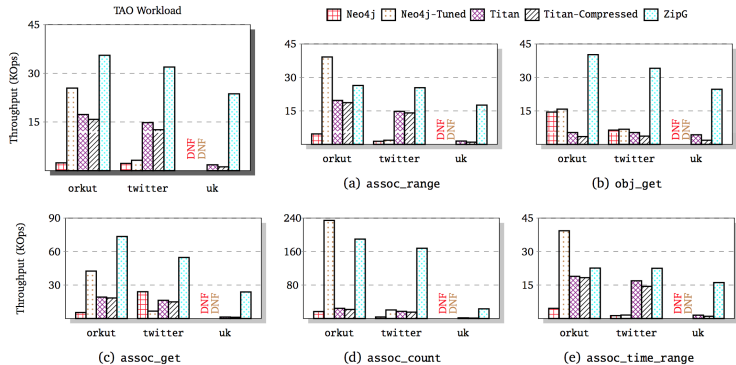


Figure 6: Single server throughput for the TAO workload, and its top 5 component queries in isolation. DNF indicates that the experiment did not finish after 48 hours of data loading. Note that the figures have different y-scales.

# Distributed Servers

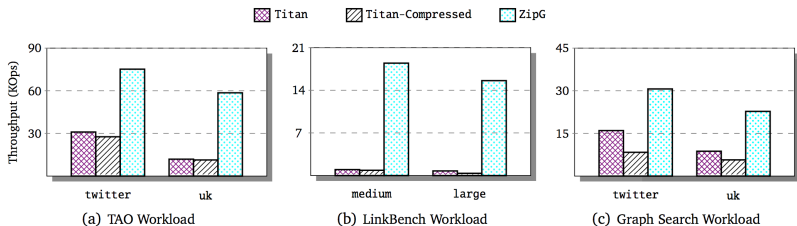









Figure 9: Throughput for TAO, LinkBench and Graph Search workloads for the distributed cluster.



## Reference

---

# Reference

-  Boldi, Paolo and Vigna, Sebastiano *The webgraph framework I: compression techniques*. WWW, 2004
-  Flavio Schierichetti, et al. *On compressing social networks*. SIGMOD, 2009
-  Wenfei Fan, et al. *Query preserving graph compression*. SIGMOD, 2012
-  Hernández C, et al. *Compression of web and social graphs supporting neighbor and community queries* SIGKDD, 2011
-  Hernández C, et al. *Compressed representations for web and social graphs*. KAIS, 2014
-  Maccioni A, Abadi D J. *Scalable pattern matching over compressed graphs via dedensification*. SIGKDD, 2016
-  Maserrat H, Pei J. *Neighbor query friendly compression of social networks*. SIGKDD, 2010